

CSE 31

Computer Organization

Lecture 8 – C Memory Management

Announcement

▶ Labs

- Lab 3 due this week (with 7 days grace period after due date)
 - Demo is REQUIRED to receive full credit
- Lab 4 out this week
 - Due at 11:59pm on the same day of your next lab
 - You must demo your submission to your TA within 14 days

▶ Reading assignment

- Reading 02 (zyBooks 2.1 – 2.9) due **tonight**, 27-SEP and Reading 03 (zyBooks 3.1 - 3.7, 3.9) due 11-OCT
 - Complete Participation Activities in each section to receive grade towards Participation
 - IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses

Announcement

▶ Homework assignment

- Homework 01 (zyBooks 1.1 – 1.5) due **tonight**, 27-SEP and Homework 02 (zyBooks 2.1 - 2.9) due 04-OCT
 - Complete *Challenge Activities* in each section to receive grade towards Homework
 - IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses


C Memory Management

- ▶ C has 3 pools of memory (based on the nature of usage)
 - Static storage: global variable storage, basically permanent, entire program run
 - The Stack: local variable storage, parameters, return address (location of “activation records” in Java or “stack frame” in C)
 - The Heap (dynamic malloc storage): data lives until deallocated by programmer
- ▶ C requires knowing where things are in memory, otherwise things don't work as expected
 - Java hides location of objects

Normal C Memory Management

$\sim FFFF\ FFFF_{hex}$

- ▶ A program's **address space** contains 4 regions:
 - **stack**: local variables, grows downward
 - **heap**: space requested for pointers via `malloc()` ; resizes dynamically, grows upward
 - **static data**: variables declared outside main, does not grow or shrink
 - **code**: loaded when program starts, does not change



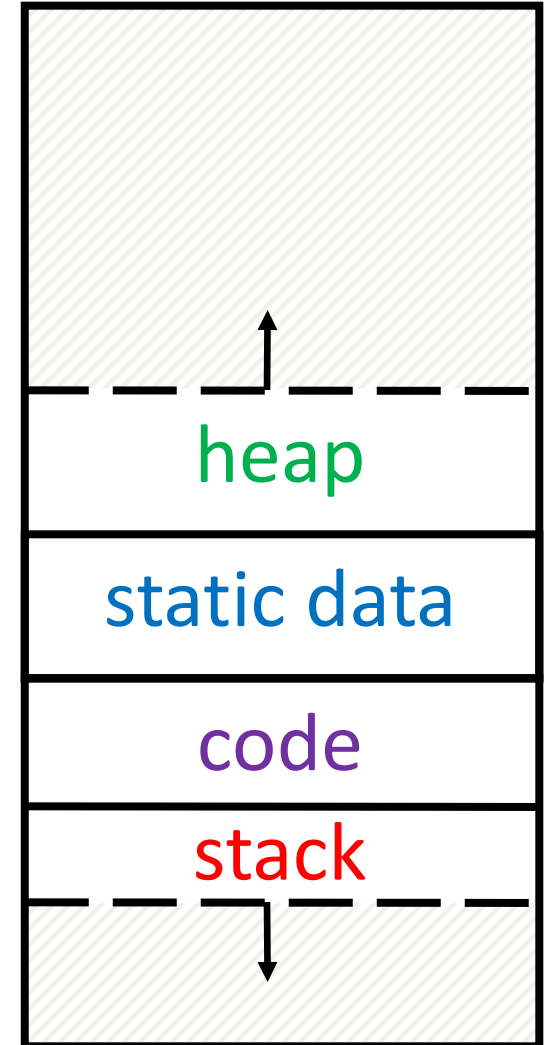
$\sim 0_{hex}$

For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory

Intel 80x86 C Memory Management

- ▶ A C program's 80x86 address space :
 - **heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward
 - **static data**: variables declared outside main, does not grow or shrink
 - **code**: loaded when program starts, does not change
 - **stack**: local variables, grows downward

~ 08000000_{hex}




Where are variables allocated?

- ▶ If declared outside of any function
 - allocated in “static” storage
- ▶ If declared inside of a function
 - allocated in the “stack”
 - freed when a function returns.
 - That’s why the scope is within the function
- ▶ Note: `main()` is a function!

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

Stack frames

- ▶ Stack frame includes storage for:
 - Return “instruction” address
 - Parameters (input arguments)
 - Space for other local variables
- ▶ Stack frames:
 - contiguous blocks of memory for a function
 - stack pointer tells where top stack frame is
- ▶ When a function ends, stack frame is “**popped off**” the stack; frees memory for future stack frames



Stack

- ▶ Last In, First Out (LIFO) data structure

stack

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack
grows
down



Stack


- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack Pointer



stack



Stack
grows
down



Stack


- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack Pointer



stack



Stack
grows
down

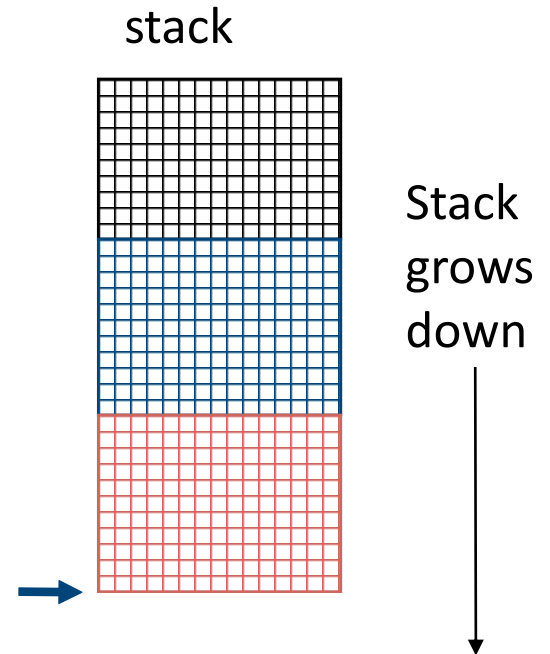


Stack

- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack Pointer




Stack

- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```


Stack Pointer



Stack

- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```




Stack

- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

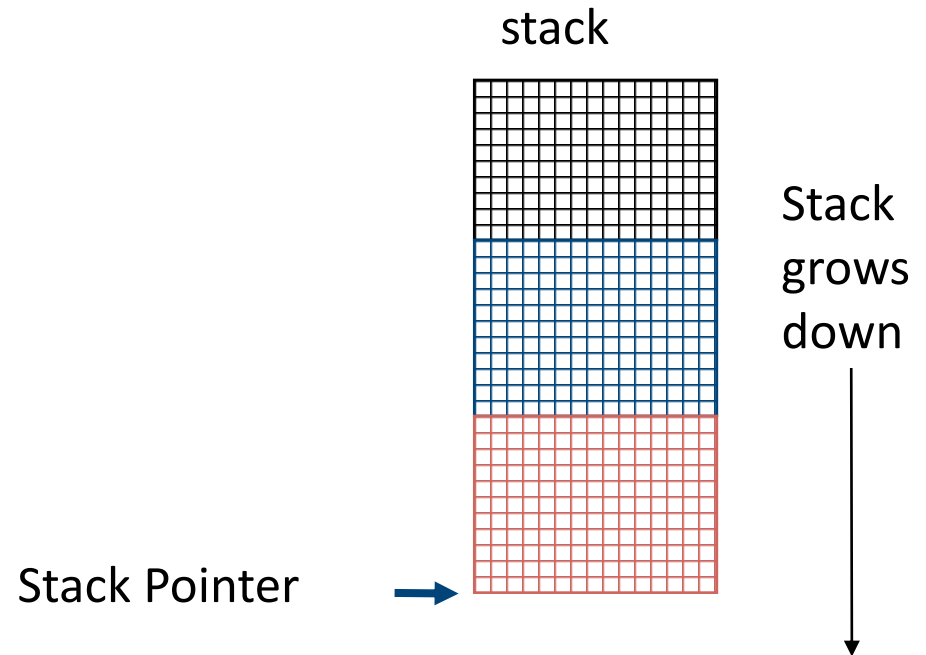
Stack Pointer



Stack

- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```



Stack


- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack Pointer



stack



Stack
grows
down



Stack


- ▶ Last In, First Out (LIFO) data structure

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```

Stack Pointer



stack



Stack
grows
down



Stack

- ▶ Last In, First Out (LIFO) data structure

stack

```
main () {  
    a(0);  
}  
void a (int m) {  
    b(1);  
}  
void b (int n) {  
    c(2);  
}  
void c (int o) {  
    d(3);  
}  
void d (int p) {  
}
```


Stack
grows
down



Who cares about stack management?

- ▶ Pointers in C allow access to deallocated memory, leading to hard-to-find bugs !

```
int *ptr () {  
    int y;  
    y = 3;  
    return &y;  
}
```



```
int main () {  
    int *stackAddr, content;  
    stackAddr = ptr();  
    content = *stackAddr;  
    printf("%d", content); /* 3 */  
    content = *stackAddr;  
    printf("%d", content); /* -2 */  
    return 0;  
}
```

The Heap (Dynamic memory)

- ▶ Large pool of memory, not allocated in contiguous order
 - back-to-back requests for heap memory could result in blocks very far apart
 - where Java/C++ **new** command allocates memory
- ▶ In C, specify number of bytes of memory explicitly to allocate item

```
int *ptr;  
ptr = (int *) malloc(sizeof(int));  
/* malloc returns type (void *),  
so need to cast to right type */
```

- `malloc()`: Allocates raw, uninitialized memory from heap

Memory Management

- ▶ How do we manage memory?
 - Code, Static
 - Simple
 - They never grow or shrink
 - Stack
 - Simple
 - Stack frames are created and destroyed in last-in, first-out (LIFO) order
 - Heap
 - Tricky
 - Memory can be allocated / deallocated at any time


Heap Management Requirements

- ▶ Want `malloc()` and `free()` to run quickly.
- ▶ Want minimal memory overhead
- ▶ Want to avoid **fragmentation***
 - When most of our free memory is in many small chunks
 - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.

* This is technically called *external fragmentation*


Heap Management

- ▶ An example
 - Request R1 for 100 bytes
 - Request R2 for 1 byte
 - Memory from R1 is freed



Heap Management

- ▶ An example
 - Request R1 for 100 bytes
 - Request R2 for 1 byte
 - Memory from R1 is freed
 - Request R3 for 50 bytes



K&R Malloc/Free Implementation

- ▶ From Section 8.7 of K&R
 - Code in the book uses some C language features we haven't discussed and is written in a very terse style, don't worry if you can't decipher the code
- ▶ Each block of memory is preceded by a header that has two fields:
 - **size** of the block
 - a **pointer to the next** block
- ▶ All **free blocks** are kept in a circular linked list, the pointer field is unused in an allocated block

K&R Implementation

- ▶ `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- ▶ `free()` checks if the blocks adjacent to the freed block are also free
 - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block
 - Otherwise, the freed block is just added to the free list