

CSE 31

Computer Organization

**Lecture 9 – C Memory Management (cont.),
Integer Representation**



Announcement

▶ Labs

- Lab 3 due this week (with 7 days grace period after due date)
 - Demo is REQUIRED to receive full credit
- Lab 4 out this week
 - Due at 11:59pm on the same day of your next lab
 - You must demo your submission to your TA within 14 days

▶ Reading assignment

- Reading 03 (zyBooks 3.1 - 3.7, 3.9) due 11-OCT
 - Complete Participation Activities in each section to receive grade towards Participation
 - IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses

Announcement


- ▶ Homework assignment
 - Homework 02 (zyBooks 2.1 - 2.9) due 04-OCT
 - Complete *Challenge Activities* in each section to receive grade towards Homework
 - IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses

Normal C Memory Management (review)

- ▶ A program's **address space** contains 4 regions:
 - **stack**: local variables, grows downward
 - **heap**: space requested for pointers via `malloc()` ; resizes dynamically, grows upward
 - **static data**: Initialized/uninitialized static and global variables
 - **code/text**: loaded when program starts, does not change

For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory

$\sim FFFF\ FFFF_{hex}$



$\sim 0_{hex}$

K&R Malloc/Free Implementation (review)

- ▶ Each block of memory is preceded by a header that has two fields:
 - **size** of the block
 - a **pointer to the next** block
- ▶ All **free blocks** are kept in a circular linked list, the pointer field is unused in an allocated block
- ▶ `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- ▶ `free()` checks if the blocks adjacent to the freed block are also free
 - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block
 - Otherwise, the freed block is just added to the free list

Choosing a block in `malloc()`

- ▶ If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
 - **best-fit**: choose the smallest block that is big enough for the request
 - **first-fit**: choose the first block we see that is big enough
 - **next-fit**: like first-fit but remember where we finished searching and resume searching from there

Tradeoffs of allocation policies

- ▶ **Best-fit:** Tries to limit fragmentation but at the cost of time (must examine all free blocks for each malloc).
 - Leaves lots of small blocks (why?)
- ▶ **First-fit:** Quicker than best-fit (why?) but potentially more fragmentation.
 - Tends to concentrate small blocks at the beginning of the free list (why?)
- ▶ **Next-fit:** Does not concentrate small blocks at front like first-fit, should be faster as a result.

Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

	1	2	3
a)	F	F	T
b)	F	T	T
c)	T	F	F
d)	T	F	T
e)	T	T	T

Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

	1	2	3
a)	F	F	T
b)	F	T	T
c)	T	F	F
d)	T	F	T
e)	T	T	T

Summary

- ▶ C has 3 pools of memory
 - Static storage: global and static variable storage, basically permanent, entire program run
 - The Stack: local variable storage, parameters, return address
 - The Heap (dynamic storage): `malloc()` grabs space from here, `free()` returns it.
- ▶ `malloc()` handles free space with freelist. Three different ways to find free space when given a request:
 - **First fit** (find first one that's free)
 - **Next fit** (same as first, but remembers where left off)
 - **Best fit** (finds most “snug” free space)

Slab Allocator

- ▶ A different approach to memory management (used in GNU `libc`)
- ▶ Divide blocks in to “large” and “small” by picking an arbitrary threshold size (say 128kB). Blocks larger than this threshold are managed with a **freelist** (as before).
- ▶ For small blocks, allocate blocks in sizes that are powers of 2
 - e.g., if program wants to allocate 20 bytes, actually give it 32 bytes


Slab Allocator

- ▶ Bookkeeping for small blocks is relatively easy
 - Use a **bitmap** for each range of blocks of the same size
- ▶ Allocating is easy and fast
 - Compute the size of the block to allocate and find a free bit in the corresponding bitmap.
- ▶ Freeing is also easy and fast
 - Figure out which slab the address belongs to and clear the corresponding bit.

Slab Allocator

16 byte blocks: 

32 byte blocks: 

64 byte blocks: 

16 byte block bitmap: 11011000

32 byte block bitmap: 0111

64 byte block bitmap: 00

Slab Allocator Tradeoffs

- ▶ Extremely fast for small blocks.
- ▶ Slower for large blocks
 - But presumably the program will take more time to do something with a large block, so the overhead is not as critical.
- ▶ Minimal space overhead
- ▶ No external fragmentation (as we defined it before)
 - For small blocks, but still have wasted space!

Internal vs. External Fragmentation

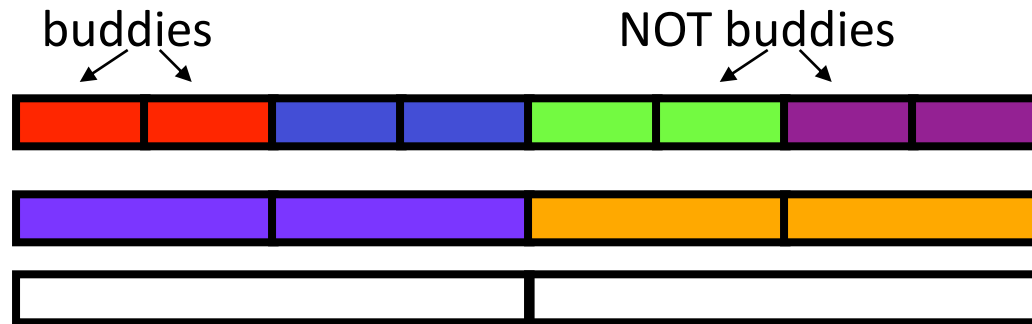
- ▶ With the slab allocator, difference between requested size and next power of 2 is wasted
 - e.g., if program wants to allocate 20 bytes and we give it a 32 byte block, 12 bytes are unused.
- ▶ We also refer to this as fragmentation, but call it **internal fragmentation** since the wasted space is actually within an allocated block.
- ▶ **External fragmentation**: wasted space between allocated blocks.

Buddy System

- ▶ Yet another memory management technique (used in Linux kernel)
- ▶ Like GNU's "slab allocator", but only allocate blocks in sizes that are powers of 2 (internal fragmentation is possible)
- ▶ Keep separate free lists for each size
 - e.g., separate free lists for 16 byte, 32 byte, 64 byte blocks, etc.

Buddy System

- ▶ If no free block of size n is available, find a block of size $2n$ and split it in to two blocks of size n
- ▶ When a block of size n is freed, if its neighbor of size n is also free, combine the blocks into a single block of size $2n$
 - **Buddy** is a block in other half larger block



- ▶ Same speed advantages as slab allocator

Allocation Schemes

- ▶ So which memory management scheme (K&R, slab, buddy) is best?
 - There is no single best approach for every application.
 - Different applications have different allocation / deallocation patterns.
 - A scheme that works well for one application may work poorly for another application.

Automatic Memory Management

- ▶ Dynamically allocated memory is difficult to track
 - Why not track it **automatically**?
- ▶ If we can keep track of what memory is in use, we can reclaim everything else.
 - Unreachable memory is called **garbage**, the process of reclaiming it is called **garbage collection**.
- ▶ So how do we track what is in use?

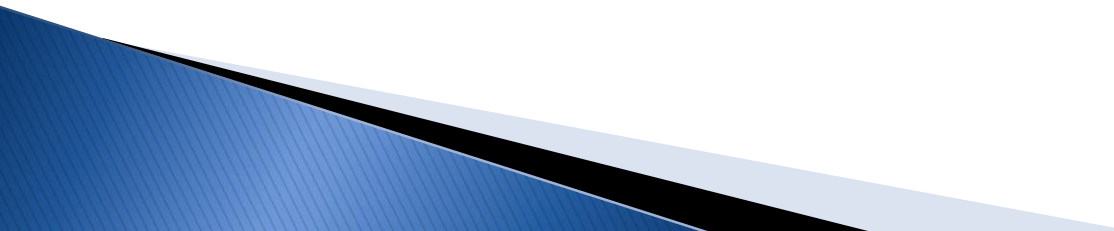
Tracking Memory Usage

- ▶ Techniques depend heavily on the programming language and rely on help from the compiler.
- ▶ Start with all pointers in global variables and local variables ([root set](#)).
- ▶ Recursively examine dynamically allocated objects we see a pointer to.
 - We can do this in **constant space** by reversing the pointers on the way down
- ▶ How do we recursively find pointers in dynamically allocated memory?

Tracking Memory Usage

- ▶ Again, it depends heavily on the programming language and compiler.
- ▶ Could have only a single type of dynamically allocated object in memory
 - E.g., simple Lisp/Scheme system with only `cons` cells
- ▶ Could use a **strongly typed** language (e.g., Java)
 - Don't allow conversion (casting) between arbitrary types.
 - C/C++ are not strongly typed.
- ▶ We will cover 3 schemes to collect garbage

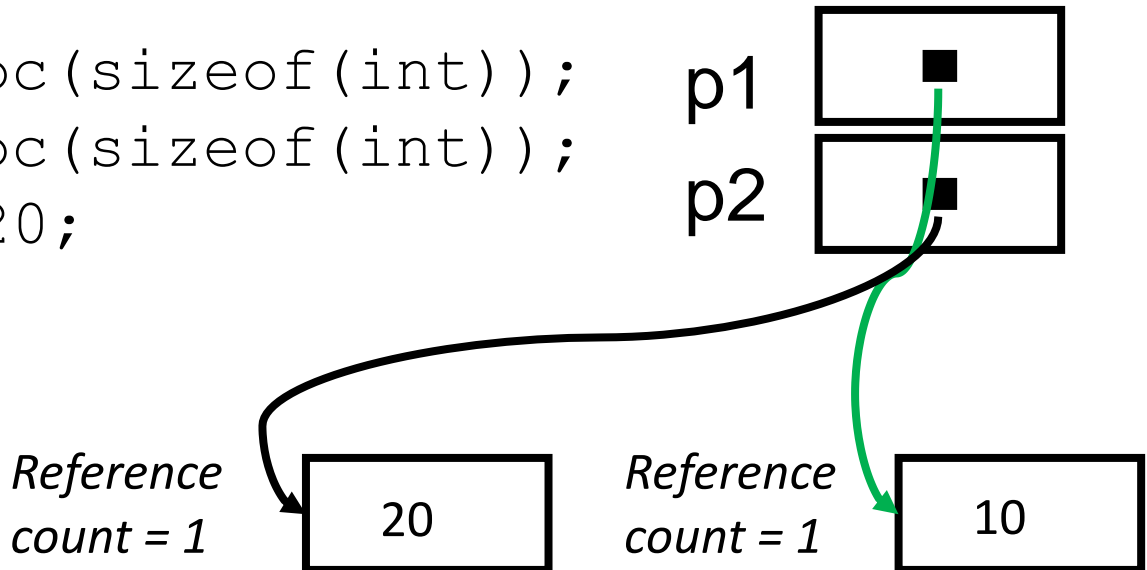
Scheme 1: Reference Counting

- ▶ For every chunk of dynamically allocated memory, keep a count of number of pointers that point to it.
 - ▶ When the count reaches 0, reclaim the memory.
 - ▶ Simple assignment statements can result in a lot of work, since may update reference counts of many items
- 

Reference Counting Example

- ▶ For every chunk of dynamically allocated memory, keep a count of number of pointers that point to it.
 - When the count reaches 0, reclaim.

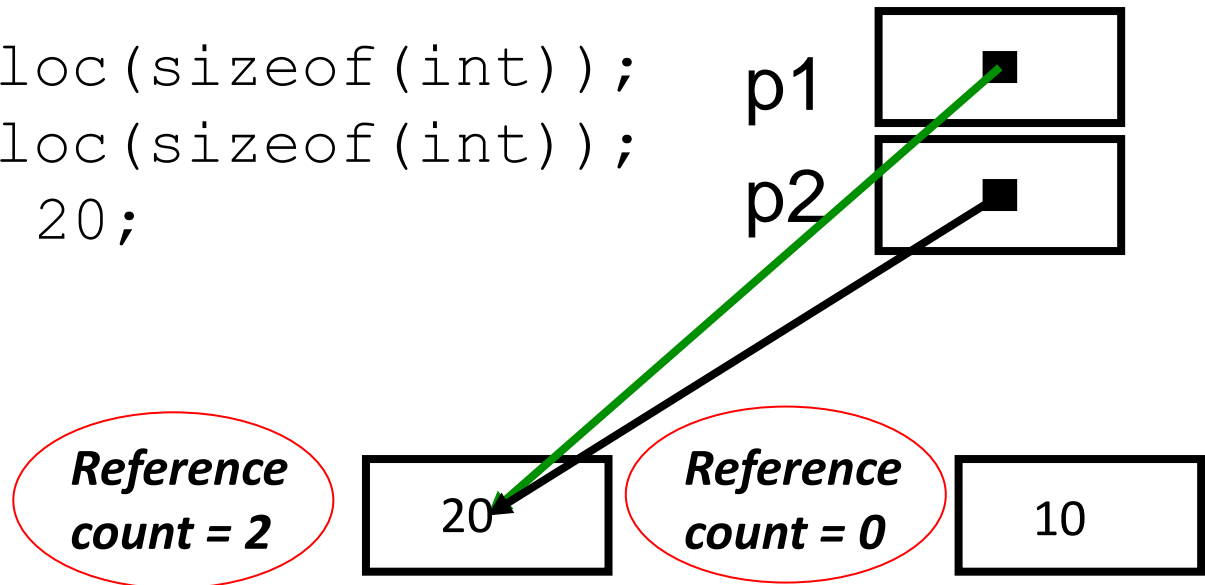
```
int *p1, *p2;  
p1 = (int *)malloc(sizeof(int));  
p2 = (int *)malloc(sizeof(int));  
*p1 = 10; *p2 = 20;
```



Reference Counting Example

- ▶ For every chunk of dynamically allocated memory, keep a count of number of pointers that point to it.
 - When the count reaches 0, reclaim.

```
int *p1, *p2;  
p1 = (int *)malloc(sizeof(int));  
p2 = (int *)malloc(sizeof(int));  
*p1 = 10; *p2 = 20;  
p1 = p2;
```




Reference Counting ($p1$, $p2$ are pointers)

$p1 = p2;$

- ▶ Increment reference count for $p2$
- ▶ If $p1$ held a valid value, decrement its reference count
- ▶ If the reference count for $p1$ is now 0, reclaim the storage it points to.
 - If the storage pointed to by $p1$ was pointing to other pointers, decrement all their reference counts, and so on...
- ▶ Must also decrement reference count when local variables cease to exist.

Reference Counting Flaws

- ▶ Extra overhead added to assignments, as well as ending a block of code.
- ▶ Does not work for circular structures!
 - E.g., doubly linked list:



Scheme 2: Mark and Sweep Garbage Collection

- ▶ Keep allocating new memory until memory is exhausted, then try to find unused memory.
- ▶ Consider objects in a graph, chunks of memory (objects) are graph nodes, pointers to memory are graph edges.
 - Edge from A to B → A stores pointer to B
- ▶ Can start with the root set, perform a graph traversal, find all usable memory!
- ▶ 2 Phases:
 1. Mark used nodes
 2. Sweep free ones, returning list of free nodes

Mark and Sweep

- ▶ Graph traversal is relatively easy to implement recursively


```
void traverse(struct graph_node *node) {  
    /* visit this node */  
    foreach child in node->children {  
        traverse(child);  
    }  
}
```

- ▶ But with recursion, state is stored on the execution stack.
 - Garbage collection is invoked when not much memory left
 - ...Oops!
- ▶ As before, we could traverse in constant space (by reversing pointers)


Scheme 3: Copying Garbage Collection

- ▶ Divide memory into two spaces, only one in use at any time.
- ▶ When active space is exhausted, traverse the active space, copying all objects to the other space, then make the new space active and continue.
 - Only reachable objects are copied!
- ▶ Use “forwarding pointers” to keep consistency
 - Simple solution to avoiding having to have a table of old and new addresses, and to mark objects already copied

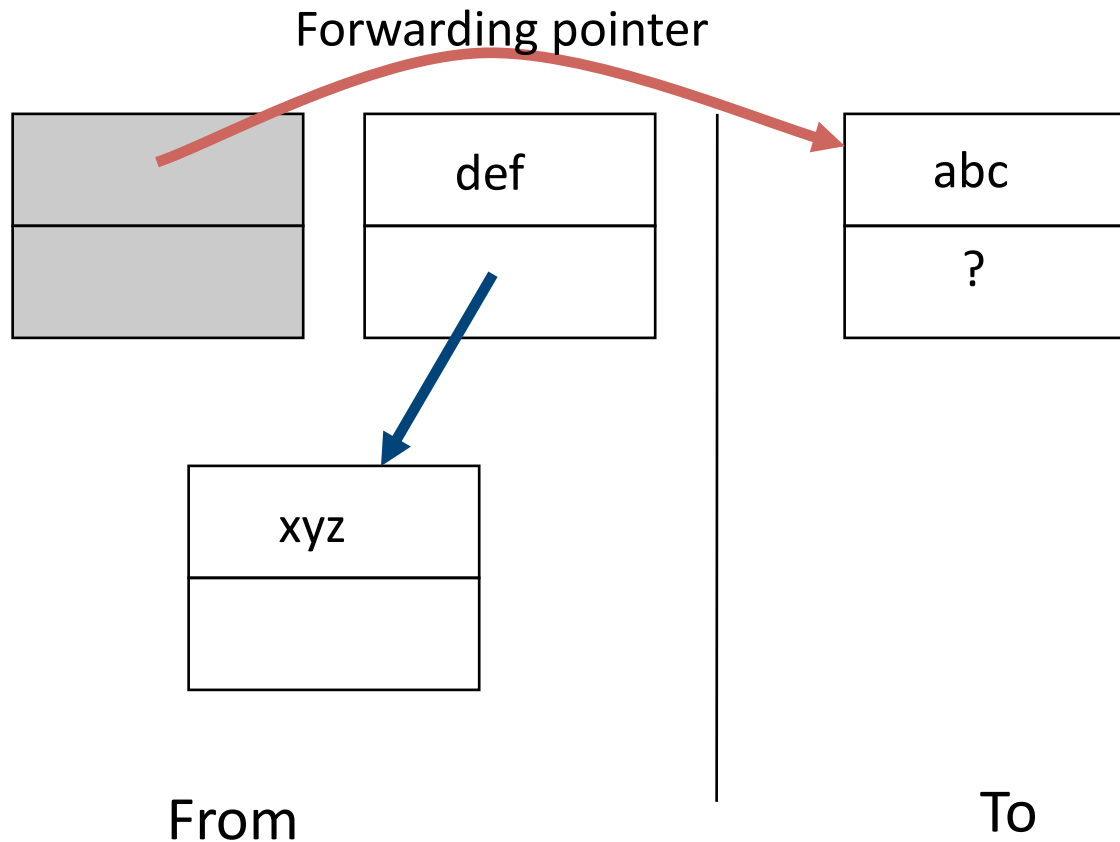
Forwarding Pointers: 1st copy “abc”



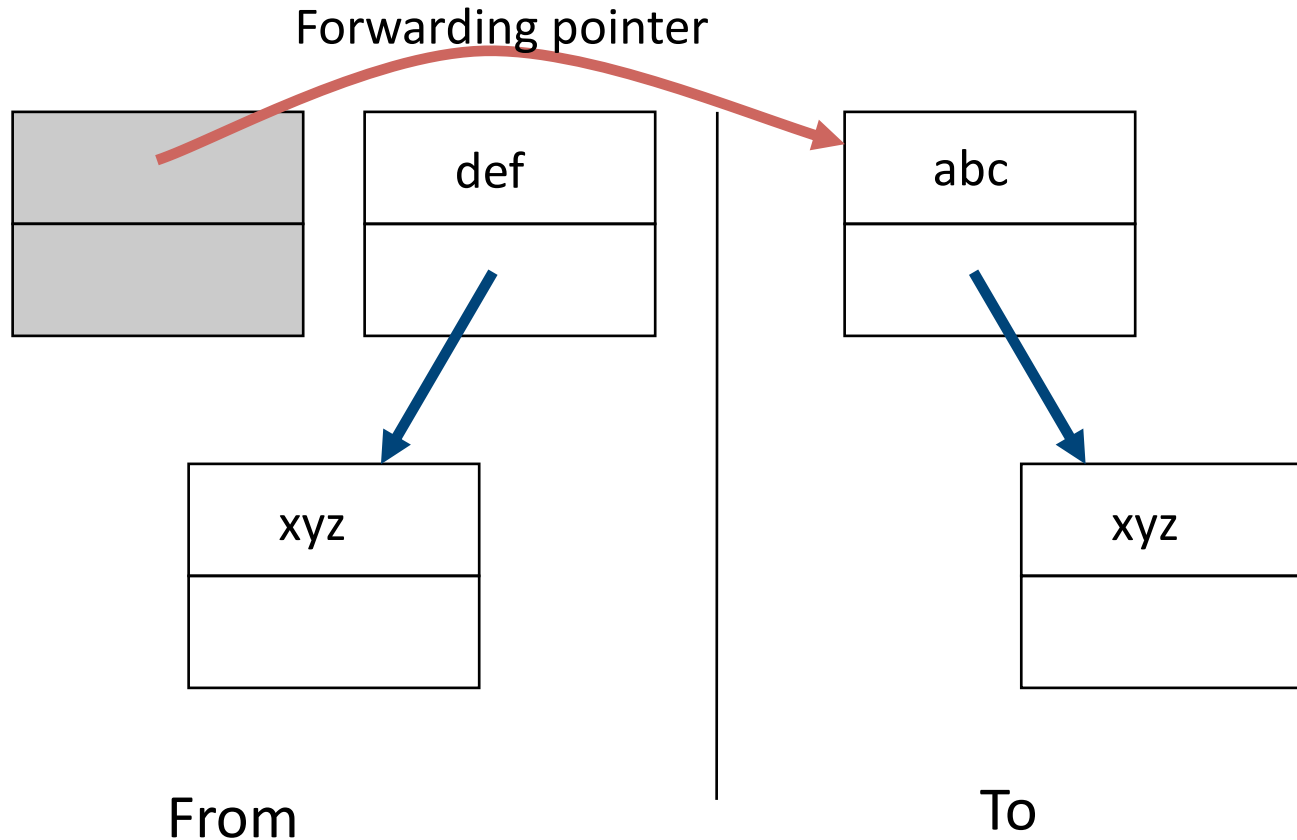
Forwarding Pointers: leave ptr to new abc



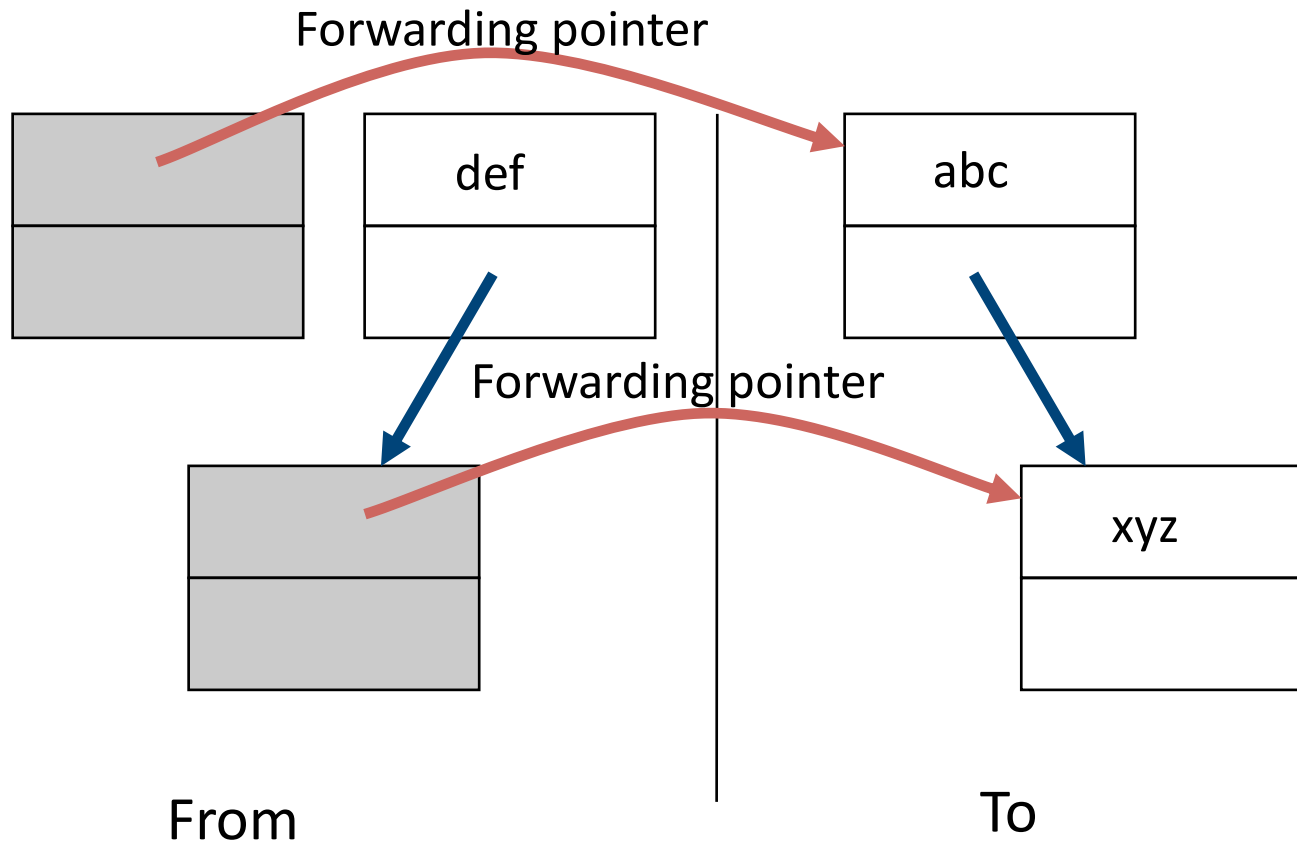
Forwarding Pointers : now copy “xyz”



Forwarding Pointers: leave ptr to new xyz

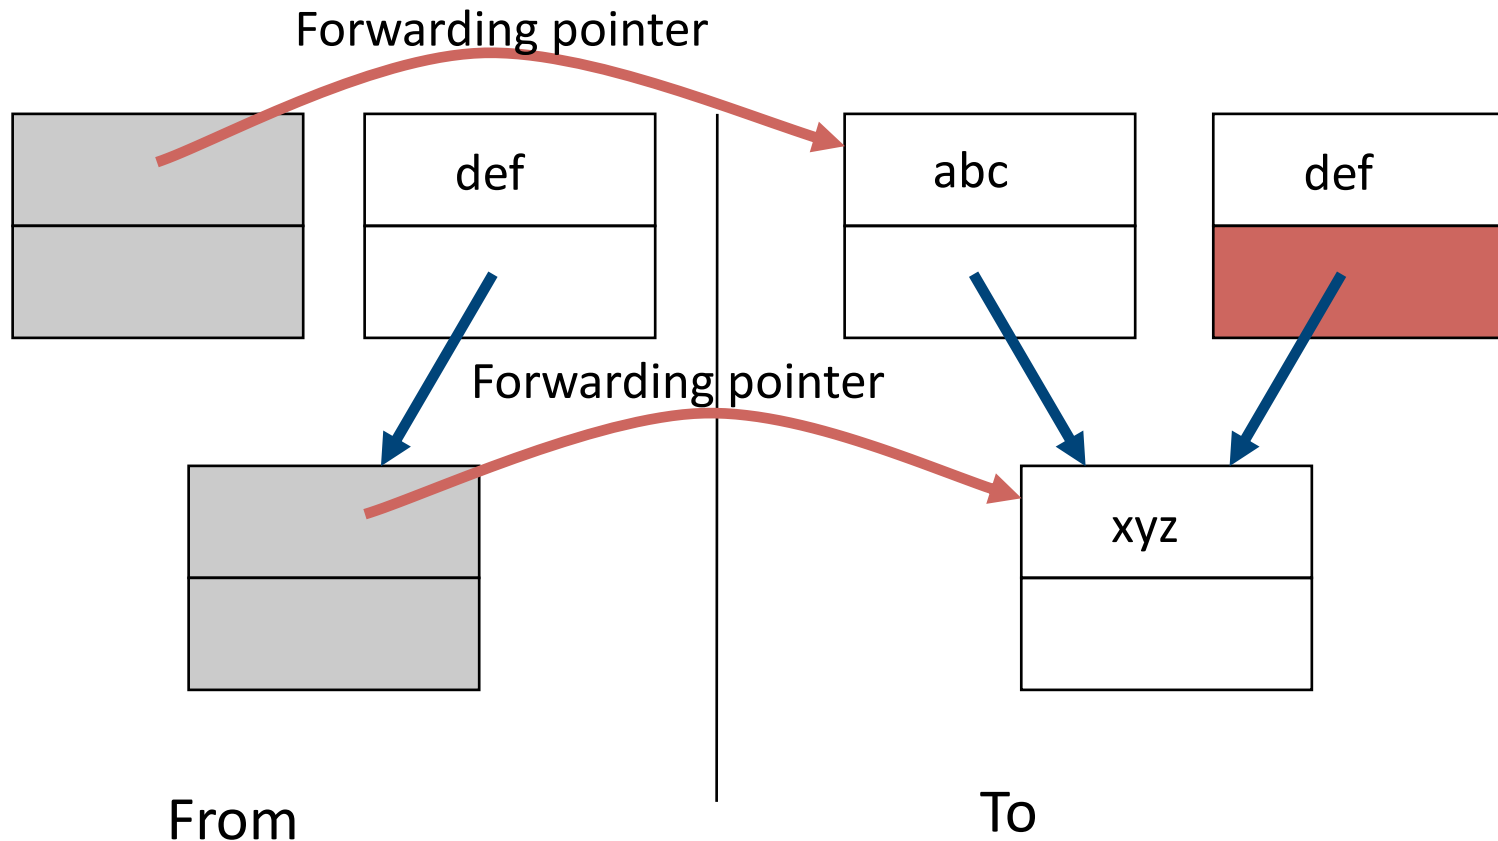


Forwarding Pointers: now copy “def”



*Since xyz was already copied,
def uses xyz's forwarding pointer
to find its new location*

Forwarding Pointers



*Since xyz was already copied,
def uses xyz's forwarding pointer
to find its new location*

Summary


- ▶ Several techniques for managing heap via malloc and free: best-, first-, next-fit
 - 2 types of memory fragmentation: internal & external; all suffer from some kind of frag.
 - Each technique has strengths and weaknesses, none is definitively best
- ▶ Automatic memory management relieves programmer from managing memory.
 - All require help from language and compiler
 - **Reference Count**: not for circular structures
 - **Mark and Sweep**: complicated and slow, works
 - **Copying**: Divides memory to copy good stuff

Number Representations

- ▶ What do these numbers mean?
 - 101
 - 0101
- ▶ Depends on what representation!

Representation and Meaning

- ▶ Objects are represented as collections of symbols (bits, digits)
- ▶ Their meaning is derived from what you do with them.



Representation (how many bits?)

▶ Characters?

- 26 letters → 5 bits ($2^5 = 32$)
- upper/lower case + punctuation → 7 bits (in 8 bits) ("ASCII")
- standard code to cover all the world's languages → 8,16,32 bits ("Unicode") www.unicode.com



▶ Logical values?

- 0 → False, 1 → True

▶ Colors?

Ex: *Red (00)*

Green (01)

Blue (11)

▶ Remember: N bits → at most 2^N things

How many bits to represent π ?

- a) 1
- b) 9 ($\pi = 3.14$, so that's 011 . 001 100)
- c) 64 (Since modern computers are 64-bit machines)
- d) Every bit the machine has!
- e) ∞

We are going to learn how to represent floating point numbers later!

What to do with representations of numbers?

► Just what we do with numbers!

- Add them
- Subtract them
- Multiply them
- Divide them
- Compare them


► Example: $10 + 7 = 17$

- ...so simple to add in binary that we can build circuits to do it!
- subtraction just as you would in decimal
- Comparison: How do you tell if $X > Y$?

$$\begin{array}{rcccc} & 1 & 1 & & \\ & 1 & 0 & 1 & 0 \\ + & 0 & 1 & 1 & 1 \\ \hline 1 & 0 & 0 & 0 & 1 \end{array}$$

What if too big?

- ▶ Binary bit patterns are simply representatives of numbers. Strictly speaking they are called “numerals”
- ▶ Numbers really have an ∞ number of digits
 - with almost all being same (00...0 or 11...1) except for a few of the rightmost digits
 - Just don't normally show leading digits
- ▶ If result of add (or -, *, /) cannot be represented by these rightmost HW bits, **overflow** is said to have occurred.



Negative Numbers


- ▶ So far, *unsigned numbers*



- ▶ Obvious solution: define leftmost bit to be sign!

- $0 \rightarrow +$, $1 \rightarrow -$
- Rest of bits can be numerical value of number

- ▶ Representation called *sign and magnitude*



Shortcomings of Sign Magnitude?

- ▶ Arithmetic circuit complicated
 - Special steps depending whether signs are the same or not
- ▶ Also, **two zeros**
 - $0x00000000 = +0_{\text{ten}}$
 - $0x80000000 = -0_{\text{ten}}$
 - What would two 0s mean for programming?
- ▶ Also, incrementing “binary odometer”, sometimes increases values, and sometimes decreases!
- ▶ Therefore sign and magnitude abandoned

Another try

► Complement the bits

- Example: $7_{10} = 00111_2$ $-7_{10} = 11000_2$
- Called **One's Complement**
- Note: positive numbers have leading 0s, negative numbers have leading 1s.
- What is -00000?
 - Answer: 11111



- How many positive numbers in N bits? 2^{N-1}
- How many negative numbers? 2^{N-1}

Shortcomings of One's complement?

- ▶ Arithmetic is less complicate than sign & magnitude.
- ▶ Still two zeros
 - $0x00000000 = +0_{\text{ten}}$
 - $0xFFFFFFFF = -0_{\text{ten}}$
- ▶ Although used for a while on some computer products, one's complement was eventually abandoned because another solution was better.

Standard Negative # Representation

- ▶ Problem is the negative mappings “overlap” with the positive ones (the two 0s). Want to shift the negative mappings left by one.
 - **Solution! For negative numbers, complement, then add 1 to the result**
- ▶ As with sign and magnitude, & one’s complement, leading 0s → positive, leading 1s → negative
 - 000000...xxx is ≥ 0 , 111111...xxx is < 0
 - except 1...1111 is -1, not -0
- ▶ This representation is **Two’s Complement**
- ▶ This makes the hardware simple!

In C: short, int, long, intN_t (C99) are all signed integers.

Two's Complement Formula

- ▶ Can represent positive and negative numbers in terms of the bit value times a power of 2:


$$d_{31} * -(2^{31}) + d_{30} * 2^{30} + \dots + d_2 * 2^2 + d_1 * 2^1 + d_0 * 2^0$$

- ▶ Example: 1101_{two}
 $= 1x - (2^3) + 1 * 2^2 + 0 * 2^1 + 1 * 2^0$
 $= -2^3 + 2^2 + 0 + 2^0$
 $= -8 + 4 + 0 + 1$
 $= -8 + 5$
 $= -3_{\text{ten}}$


Example: -3 to +3 to -3:

x :	1101 _{two}	(-3)
x' :	0010 _{two}	
+1 :	0011 _{two}	(3)
()' :	1100 _{two}	
+1 :	1101 _{two}	(-3)

2's Complement Number "line": N = 5




- ▶ 2^{N-1} non-negatives
- ▶ 2^{N-1} negatives
- ▶ **one zero**
- ▶ how many positives?
 - $2^{N-1} - 1$



Note: 1's and 2's complements are used to represent negative numbers only!

Bias Encoding: $N = 5$ (bias = 15)



- ▶ Want 00... to represent the smallest number
- ▶ value = unsigned - bias
- ▶ Bias for N bits = $2^{N-1} - 1$
- ▶ one zero
- ▶ how many positives?
 - 2^{N-1}
 - (more than 2's complement)

00000 00001 01110

01111 10000 ... 11110 11111


Binary odometer

Summary

- ▶ We represent “things” in computers as particular bit patterns:
 - N bits $\rightarrow 2^N$ things
- ▶ Different integer encodings have different benefits; 1s complement and sign/mag have most problems.
- ▶ **unsigned** (C99's `uintN_t`):



- ▶ **2's complement** (C99's `intN_t`): universal, learn it!



- ▶ Overflow: numbers ∞ ; computers finite \rightarrow errors!