

CSE100: Design and Analysis of Algorithms

Lecture 14 – Binary Search Trees (cont.)

Mar 8th 2022

Binary Search Trees and Red-Black Trees



Binary tree terminology (review)

Each node has two **children**.

The **left child** of **3** is **2**

The **right child** of **3** is **4**

The **parent** of **3** is **5**

2 is a **descendant** of **5**

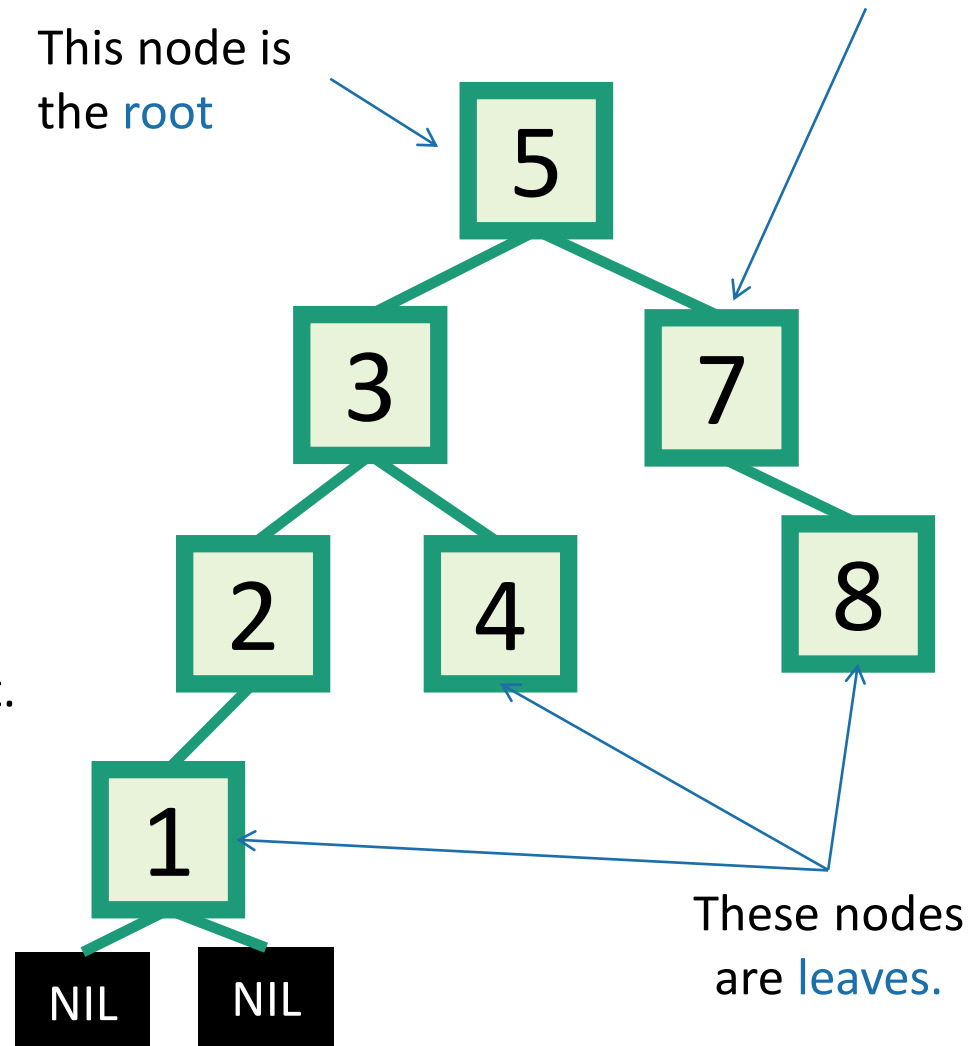
Each node has a pointer to its left child, right child, and parent.

Both **children** of **1** are NIL.
(We won't usually draw them).

The **height** of this tree is 3.
(Max number of edges from the root to a leaf).

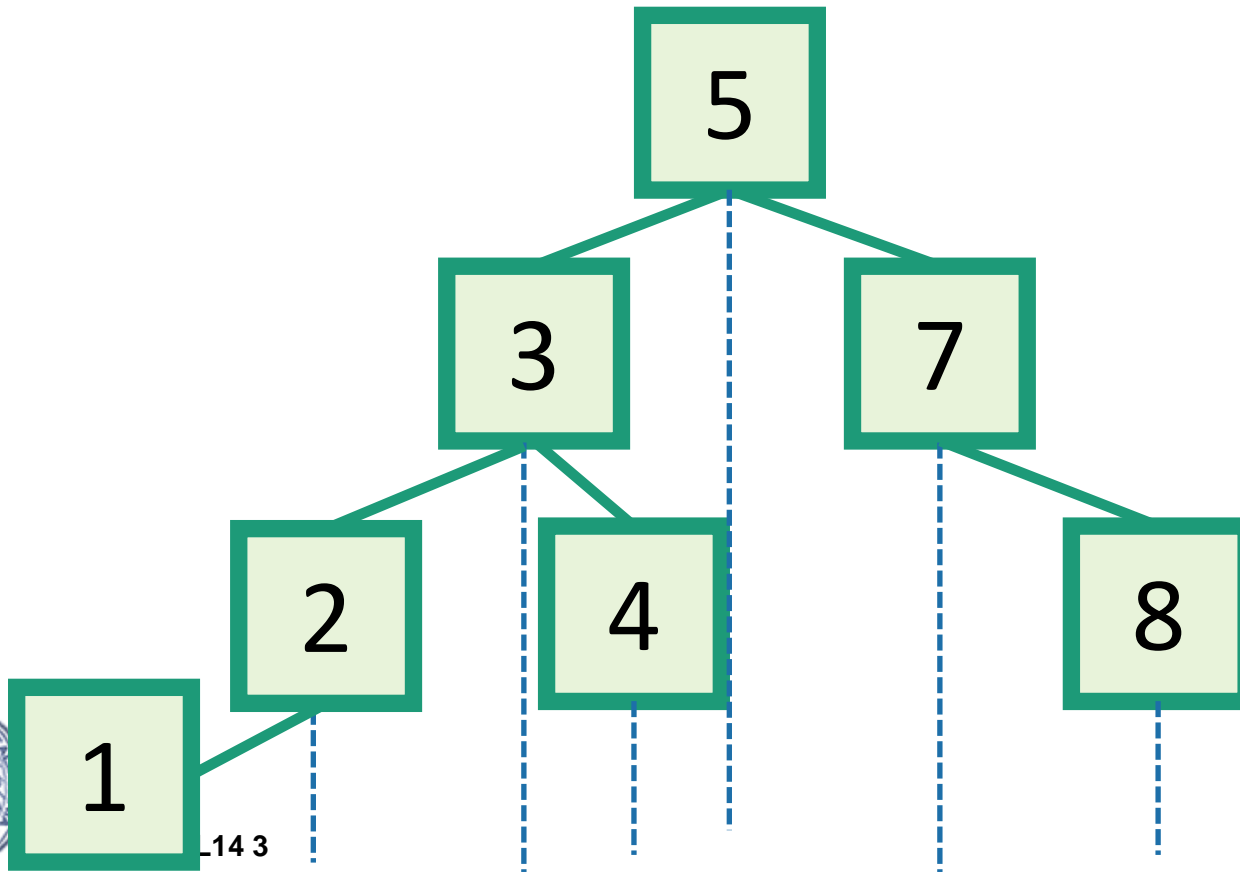
This node is the **root**

This is a **node**.
It has a **key** (7).



Binary Search Trees (review)

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



Q: Is this the only binary search tree I could possibly build with these values?

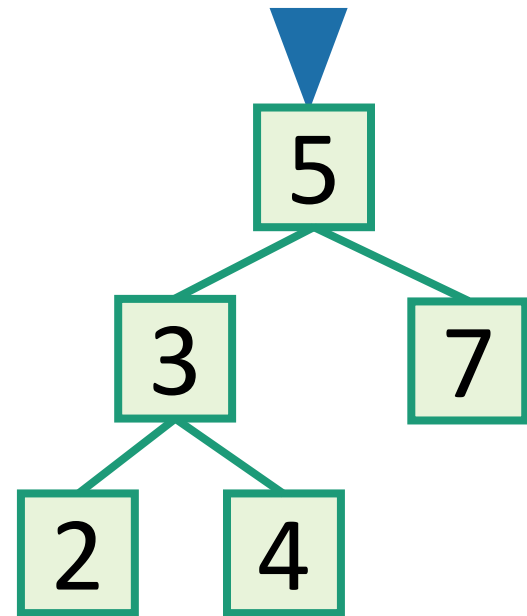
A: **No.** I made choices about which nodes to choose when. Any choices would have been fine.



Aside: In-Order Traversal of BSTs (review)

- Output all the elements in sorted order!

- `inOrderTraversal(x)`:
 - if $x \neq \text{NIL}$:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



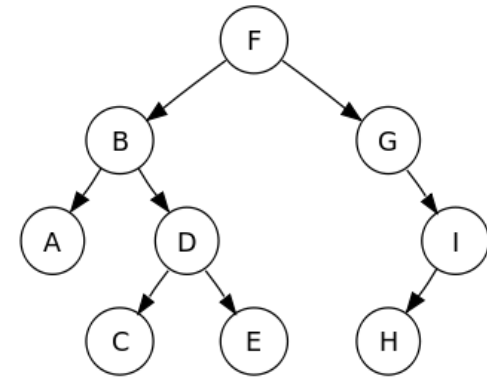
- Runs in time $O(n)$.

2 3 4 5 7 Sorted!



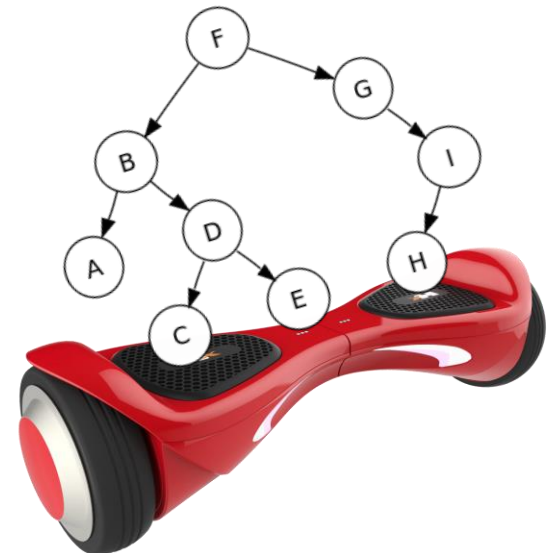
Today

- Begin a brief foray into data structures!
- Binary search trees (cont.)
 - You may remember these from CSE 30
 - They are better when they're balanced.



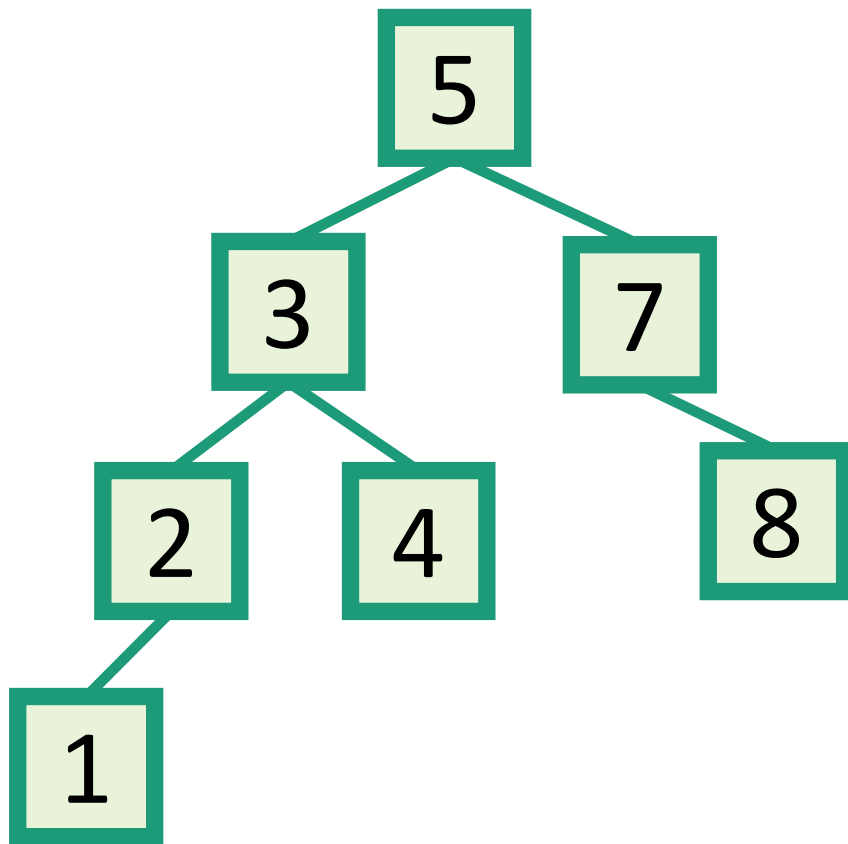
this will lead us to...

- Self-Balancing Binary Search Trees
 - **Red-Black** trees.



SEARCH in a Binary Search Tree

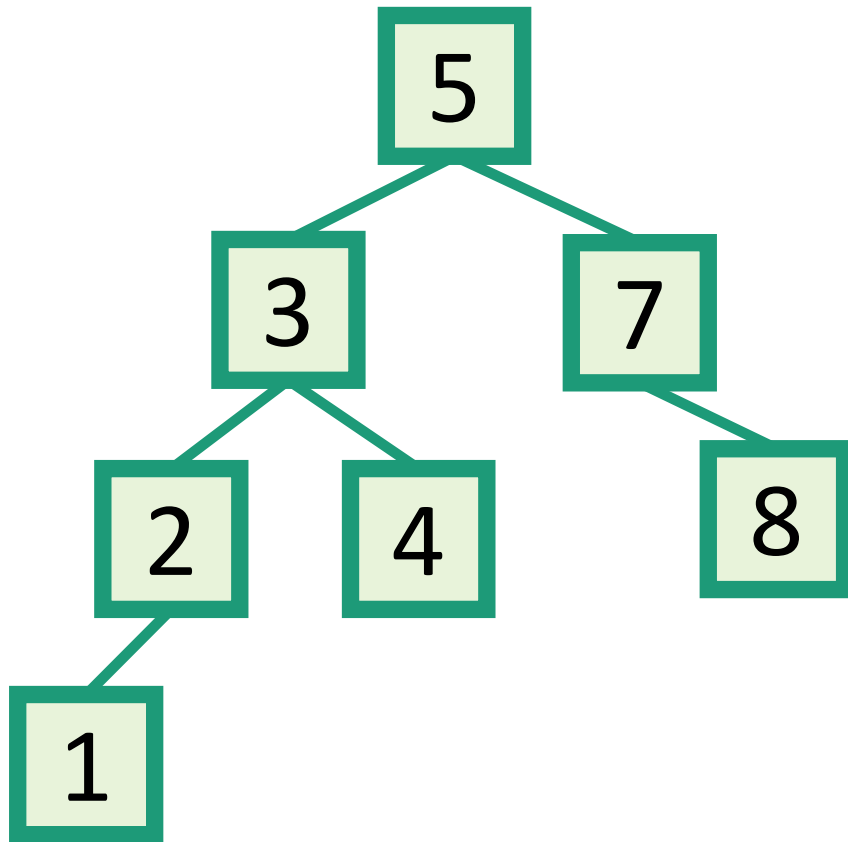
definition by example



SEARCH in a Binary Search Tree

definition by example

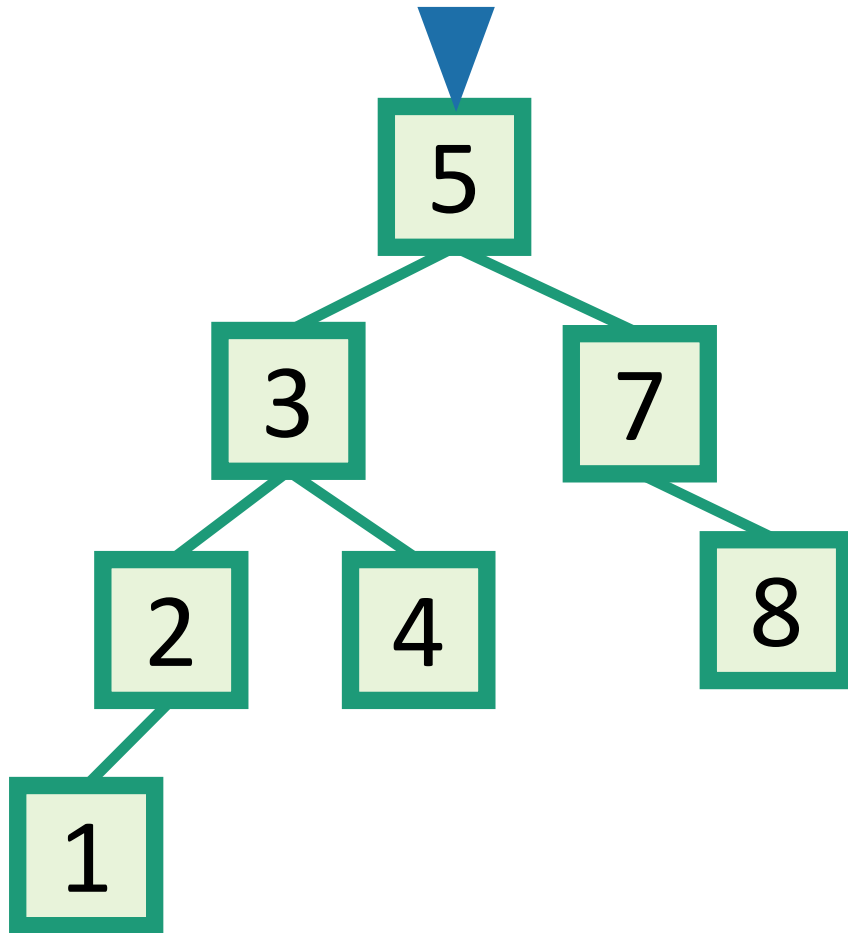
EXAMPLE: Search for 4.



SEARCH in a Binary Search Tree

definition by example

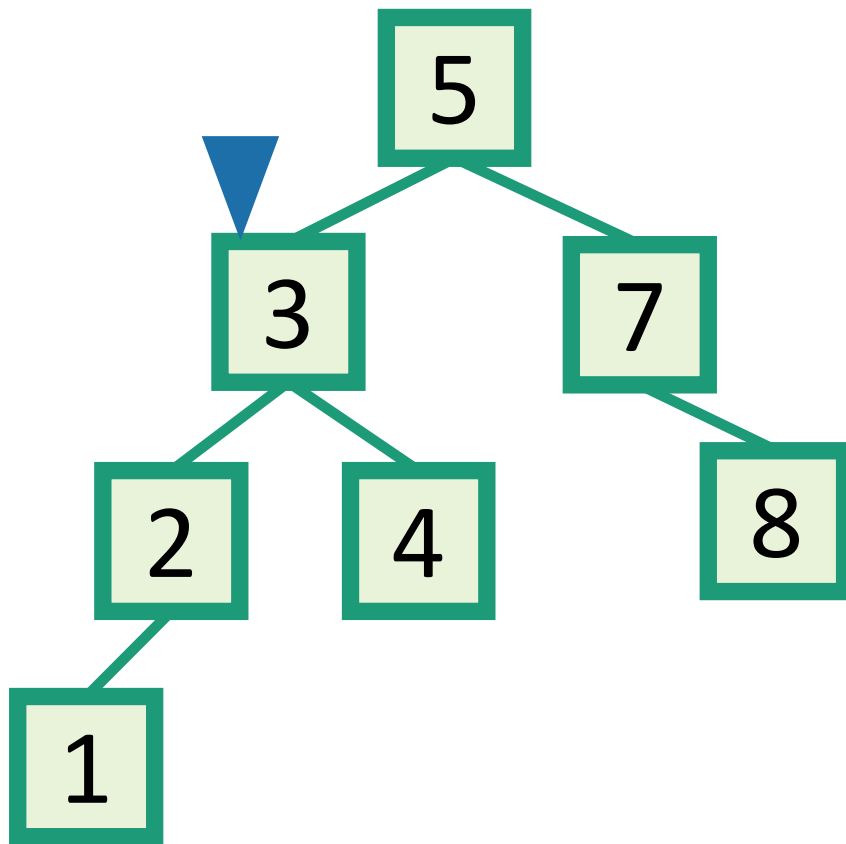
EXAMPLE: Search for 4.



SEARCH in a Binary Search Tree

definition by example

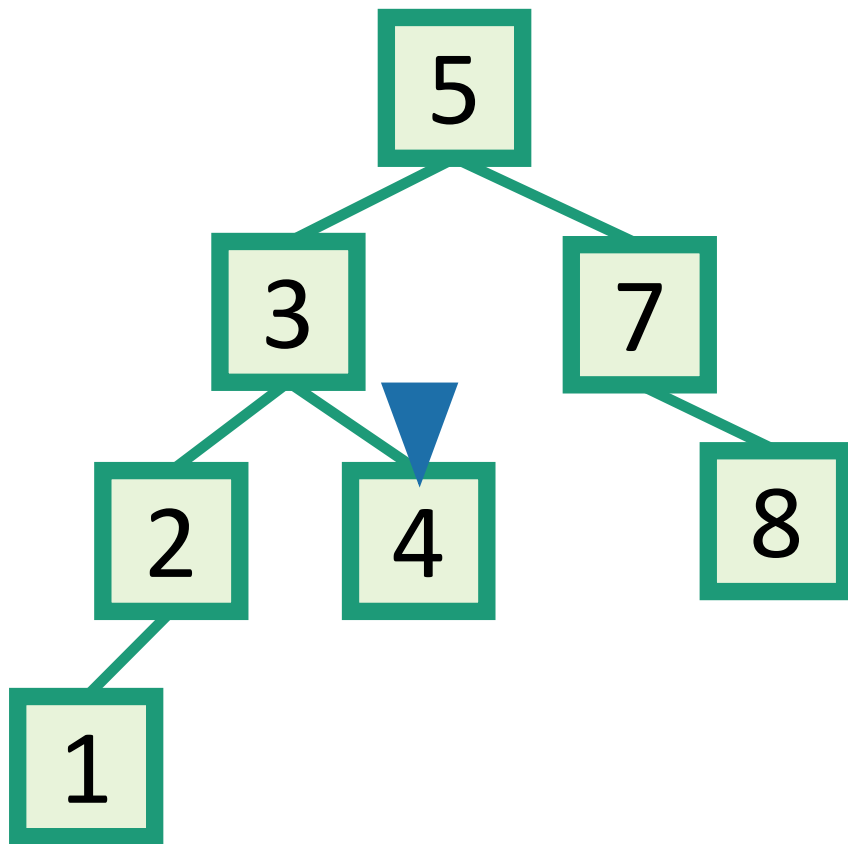
EXAMPLE: Search for 4.



SEARCH in a Binary Search Tree

definition by example

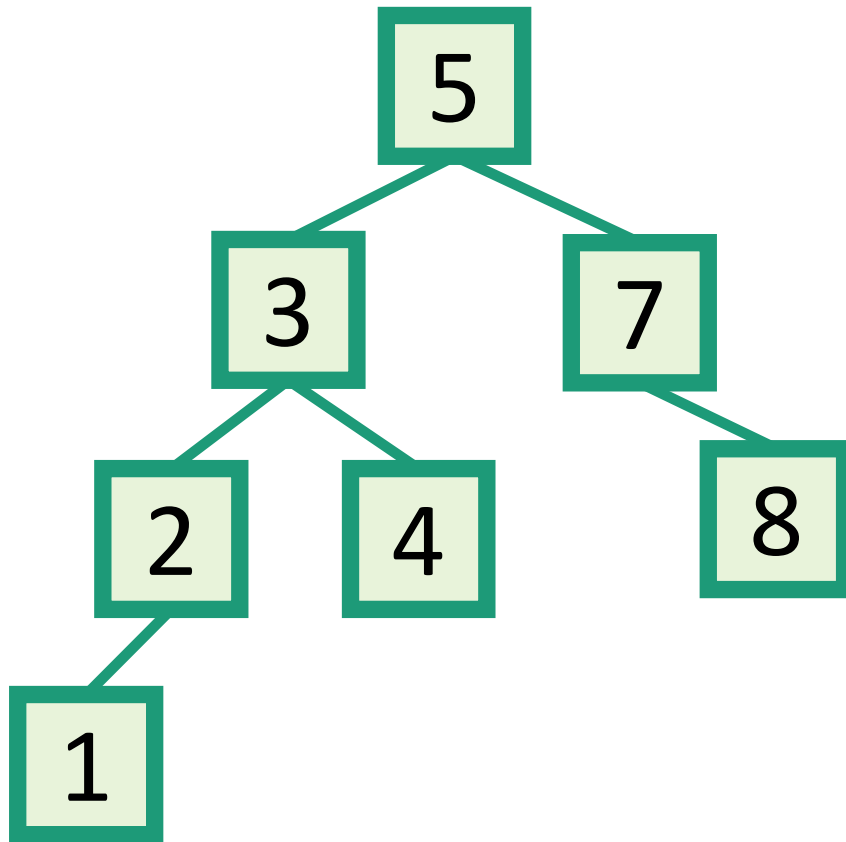
EXAMPLE: Search for 4.



SEARCH in a Binary Search Tree

definition by example

EXAMPLE: Search for 4.

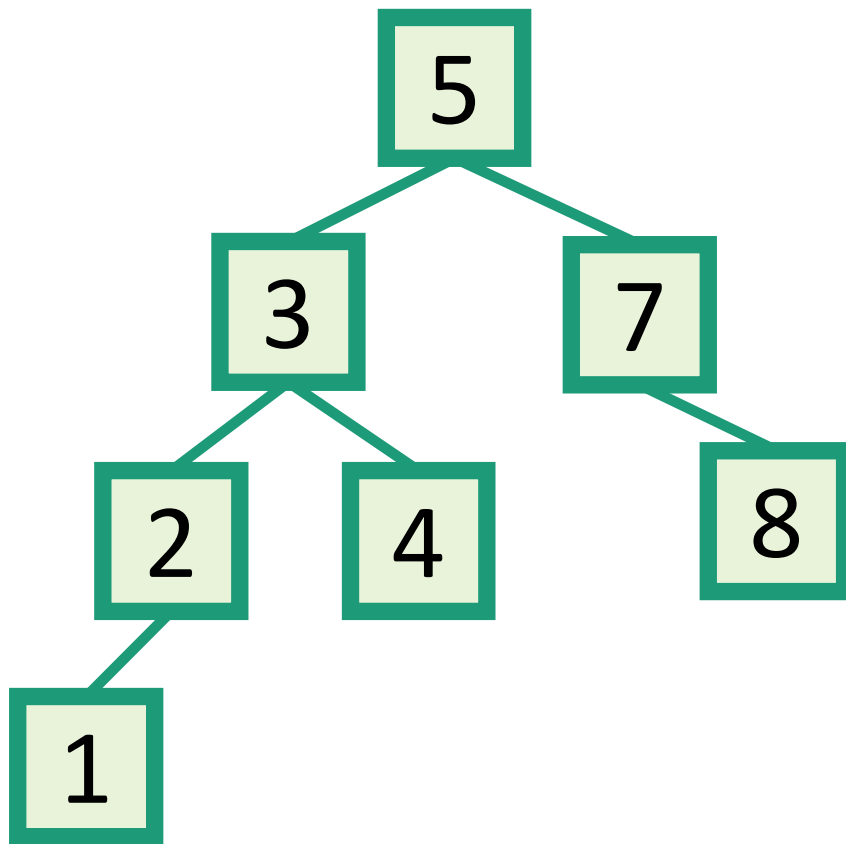


SEARCH in a Binary Search Tree

definition by example

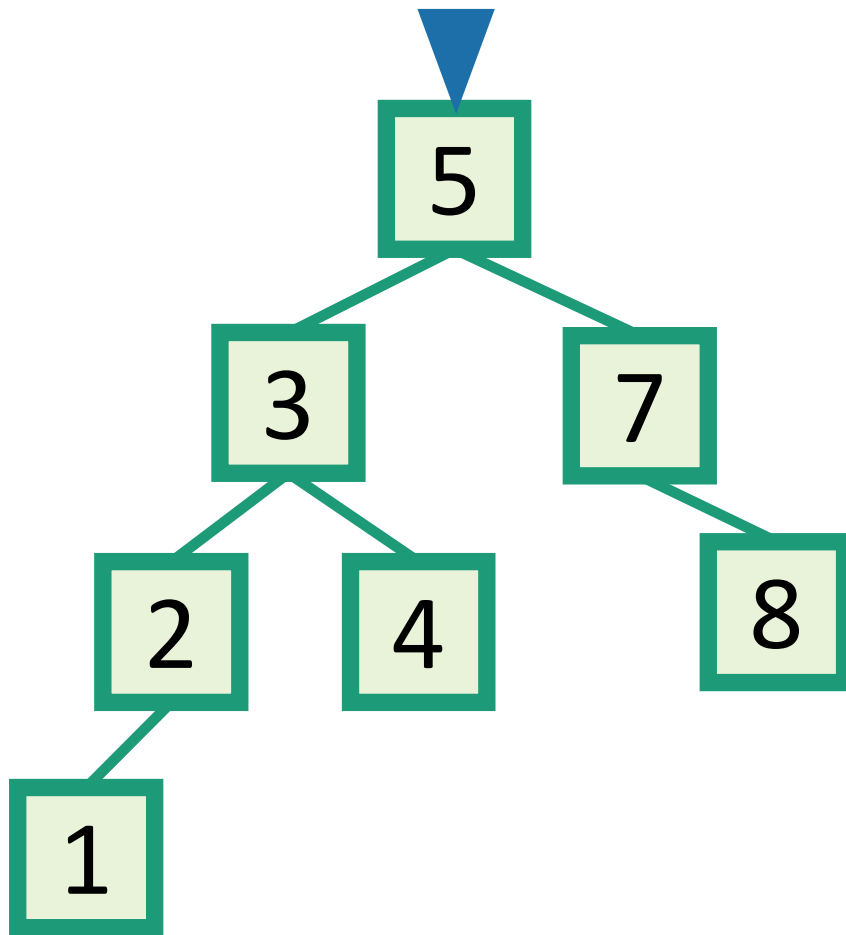
EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5



SEARCH in a Binary Search Tree

definition by example



EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

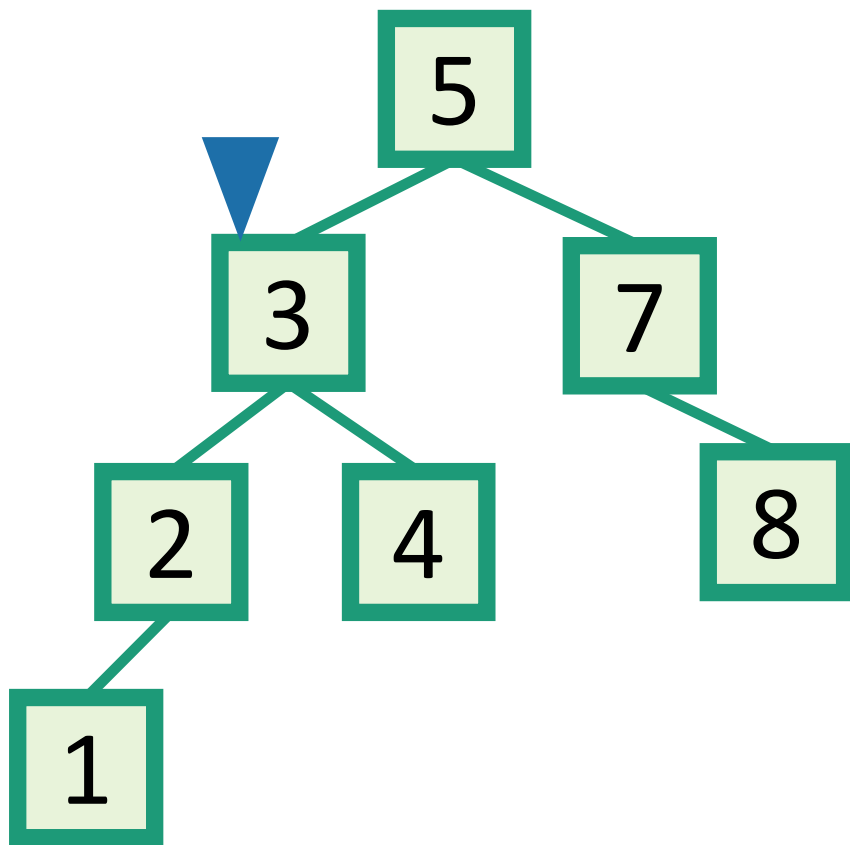


SEARCH in a Binary Search Tree

definition by example

EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

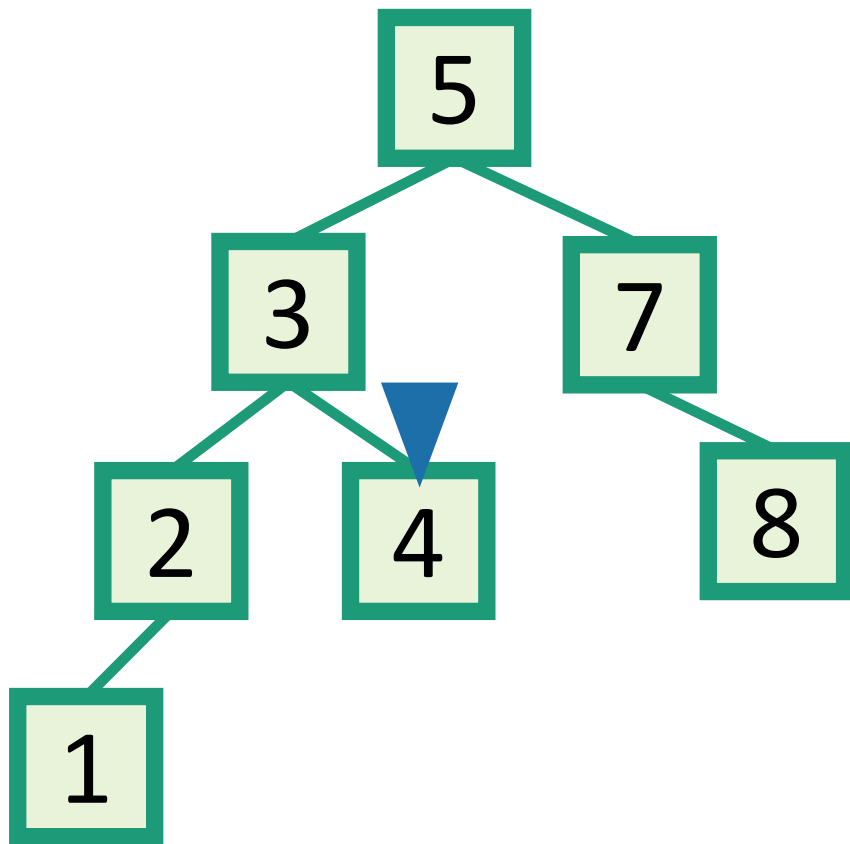


SEARCH in a Binary Search Tree

definition by example

EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

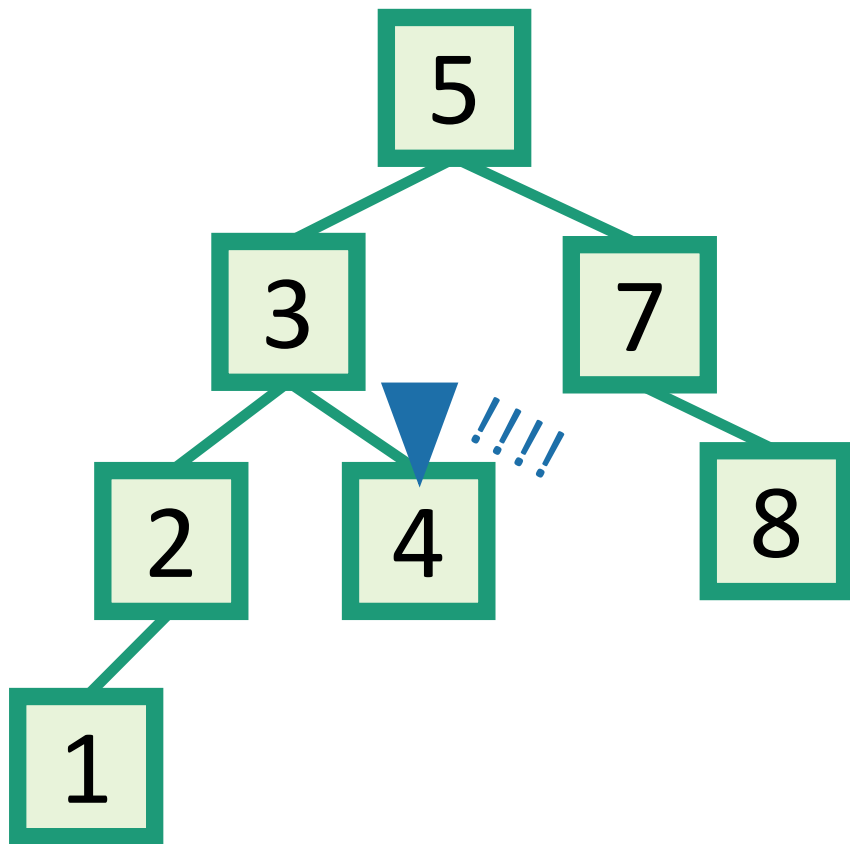


SEARCH in a Binary Search Tree

definition by example

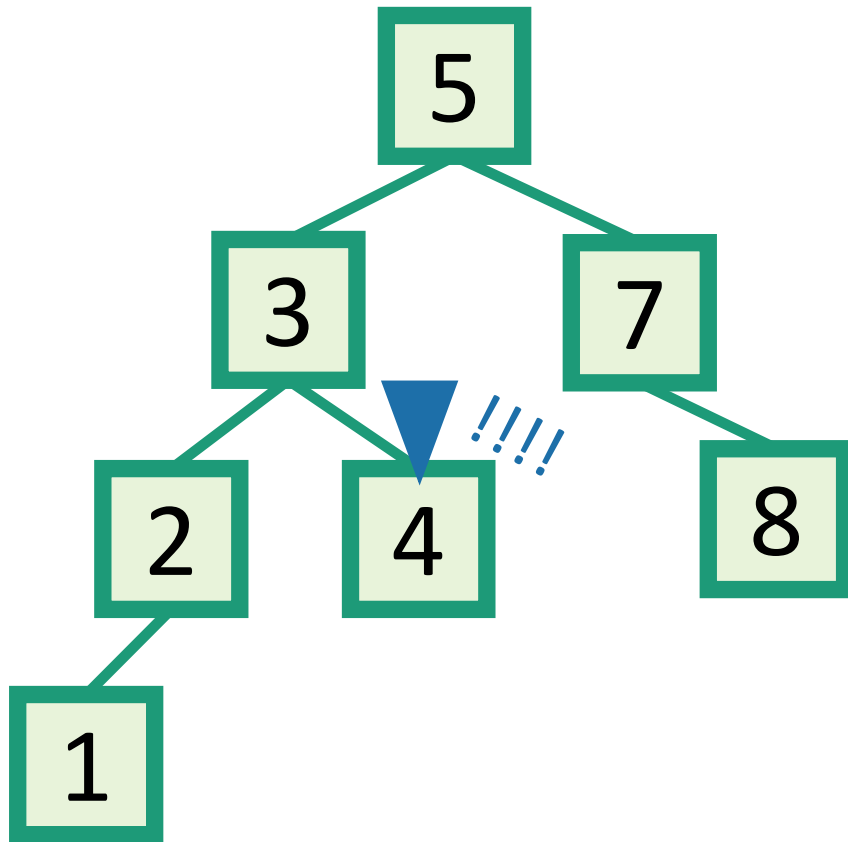
EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5



SEARCH in a Binary Search Tree

definition by example



EXAMPLE: Search for 4.

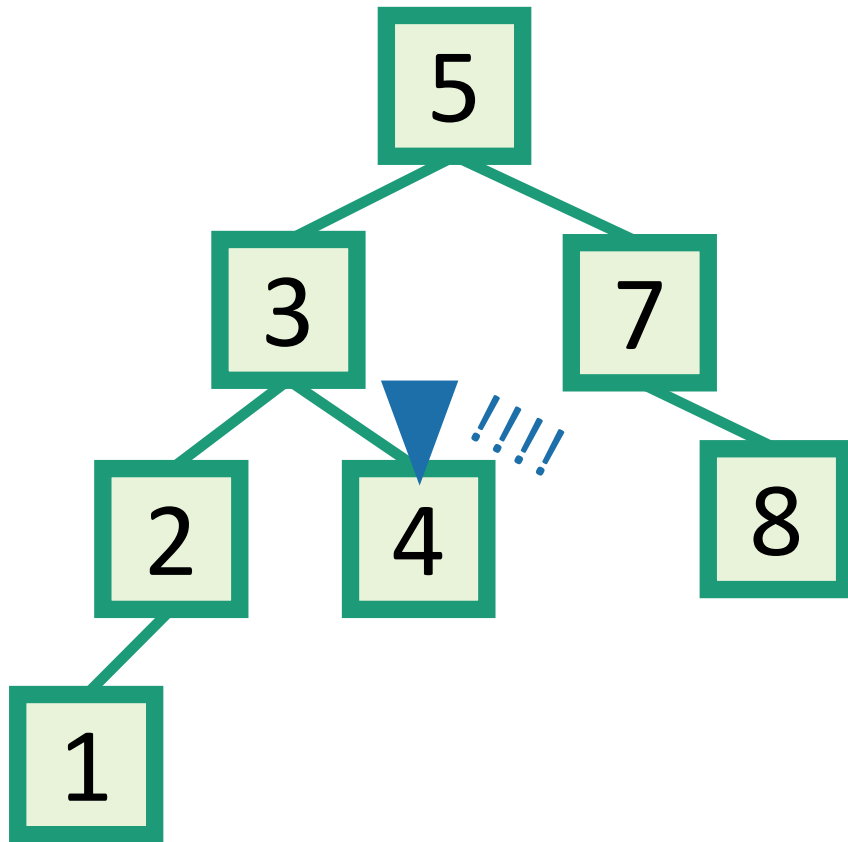
EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case



SEARCH in a Binary Search Tree

definition by example



EXAMPLE: Search for 4.

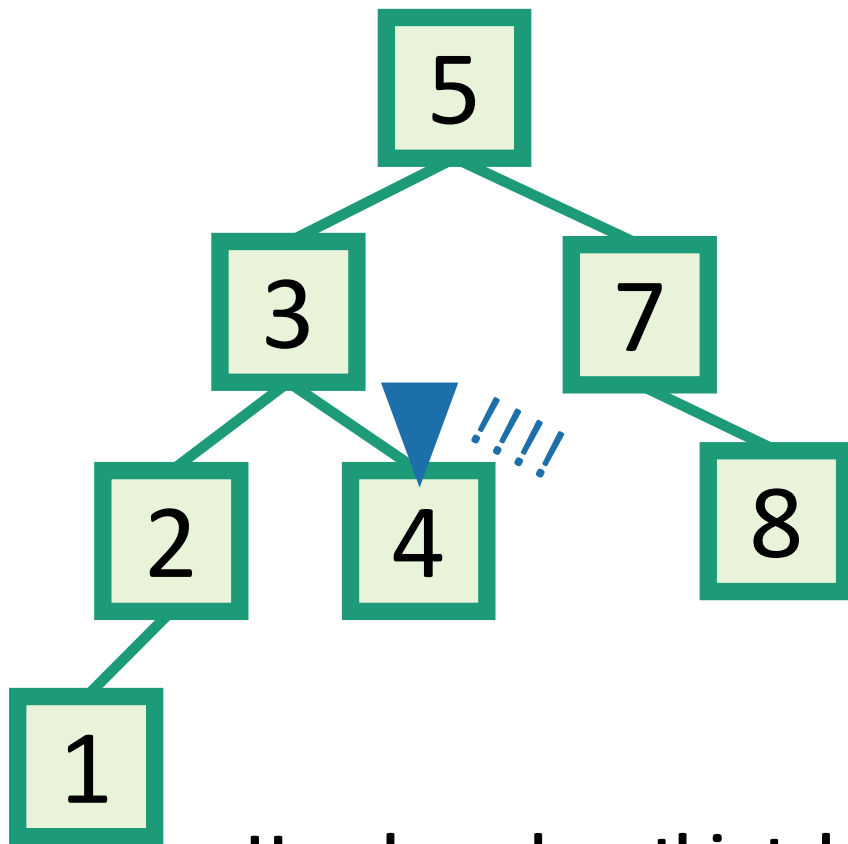
EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)



SEARCH in a Binary Search Tree

definition by example



How long does this take?

EXAMPLE: Search for 4.

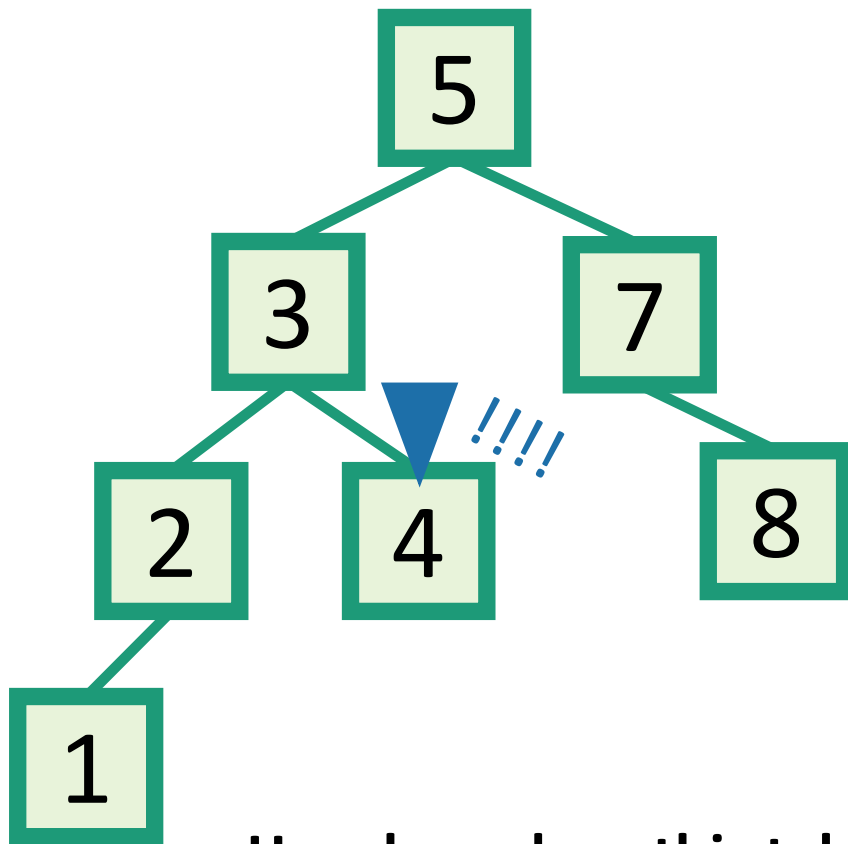
EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)



SEARCH in a Binary Search Tree

definition by example



EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

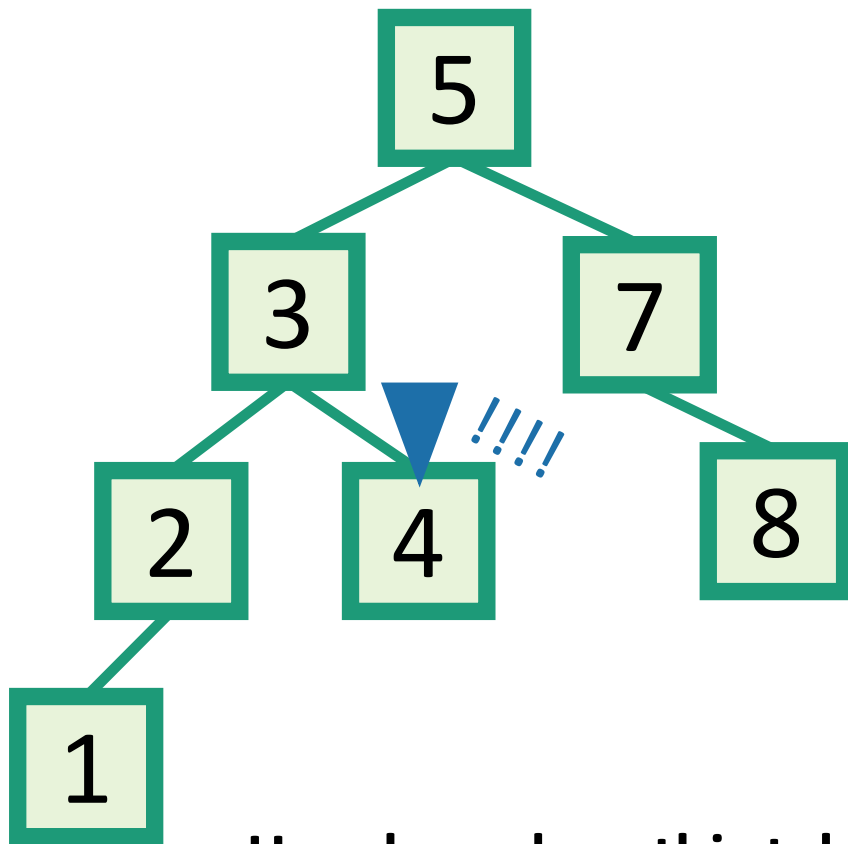
How long does this take?

$O(\text{length of longest path}) = O(\text{height})$



SEARCH in a Binary Search Tree

definition by example



How long does this take?

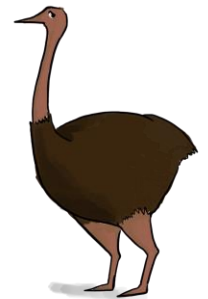
$O(\text{length of longest path}) = O(\text{height})$

EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

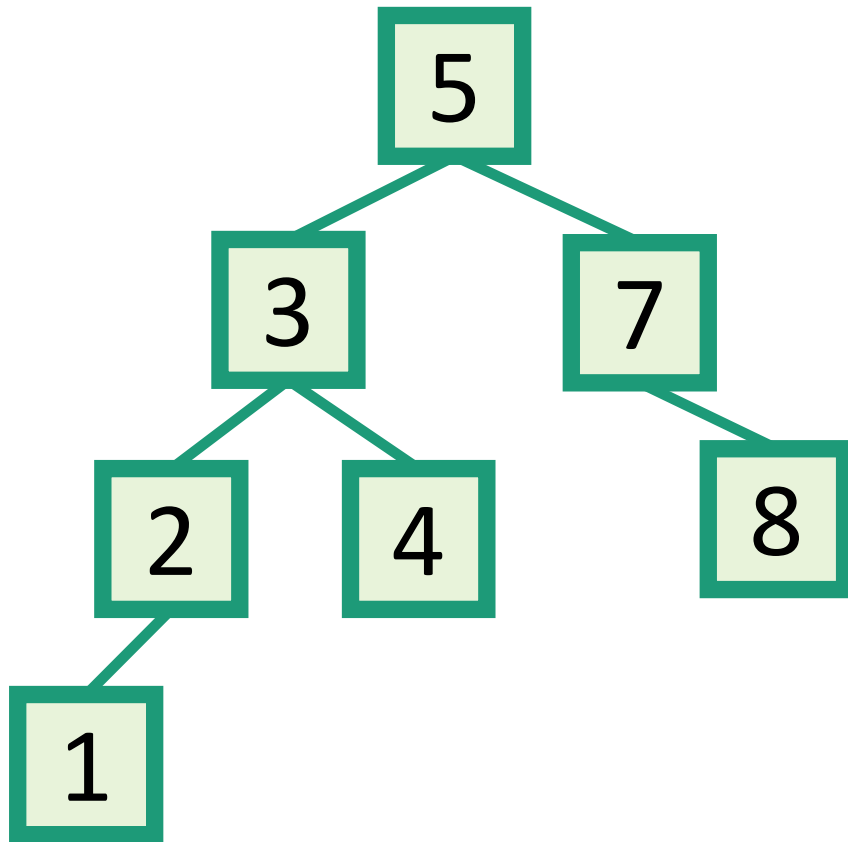
Write pseudocode
(or actual code) to
implement this!



Ollie the over-achieving ostrich

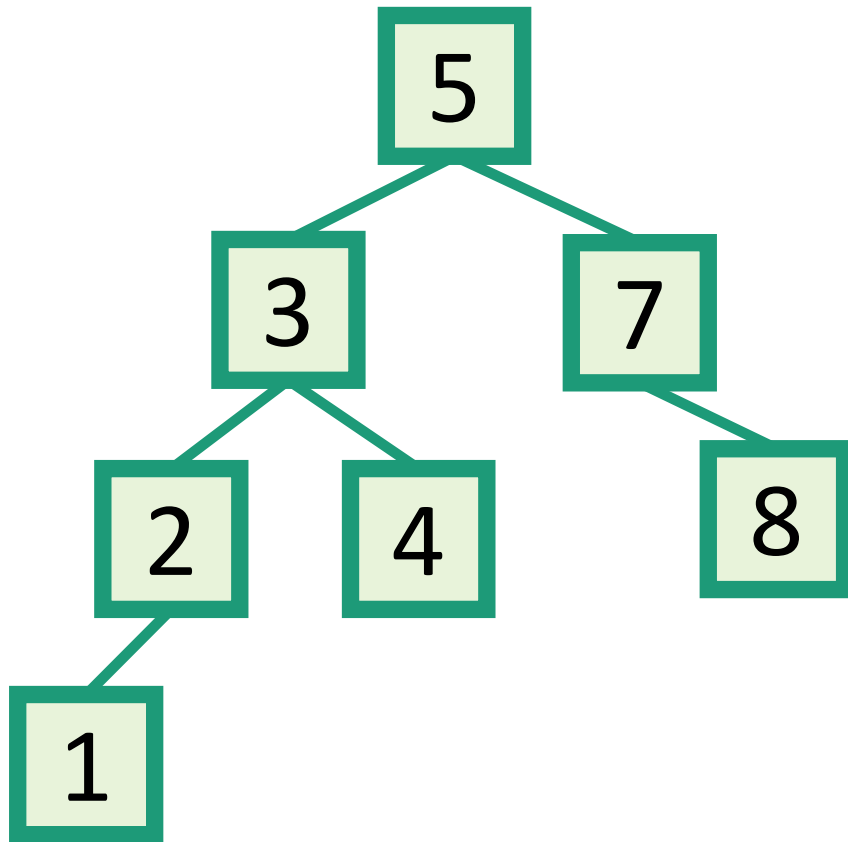


INSERT in a Binary Search Tree



INSERT in a Binary Search Tree

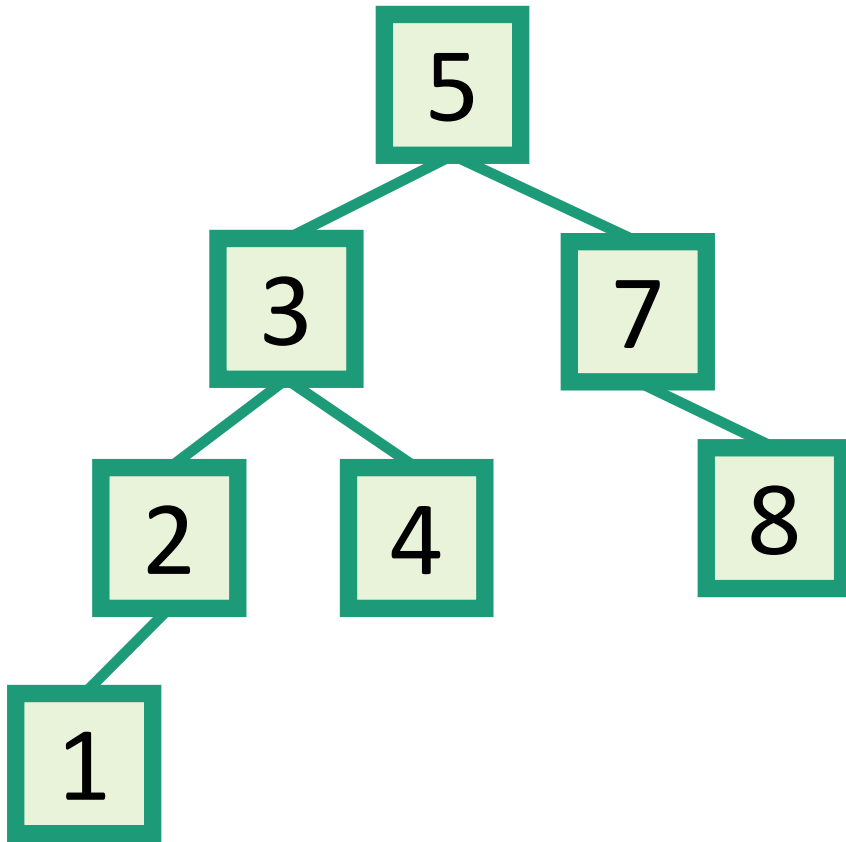
EXAMPLE: Insert 4.5



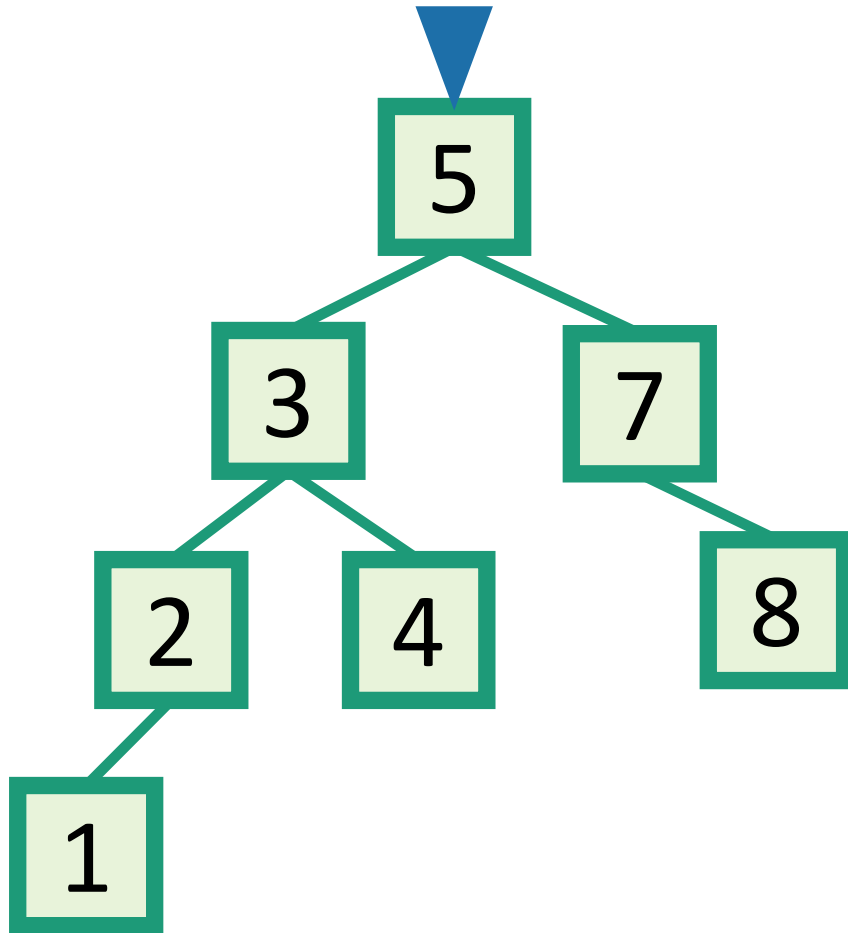
INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5

- **INSERT**(key):



INSERT in a Binary Search Tree



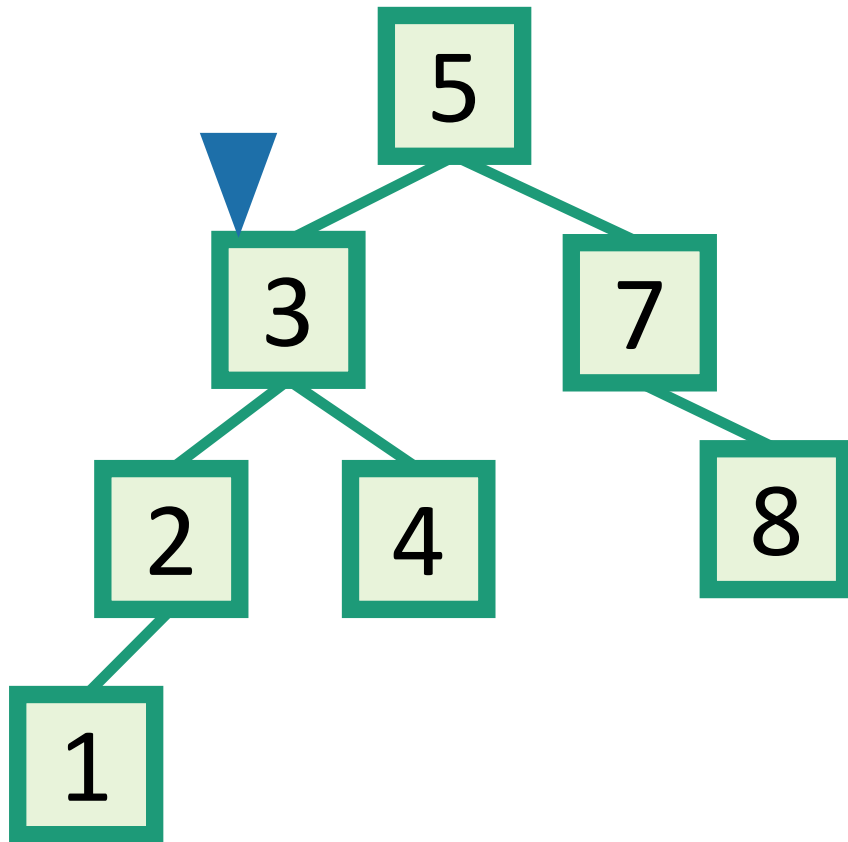
EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$



INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5



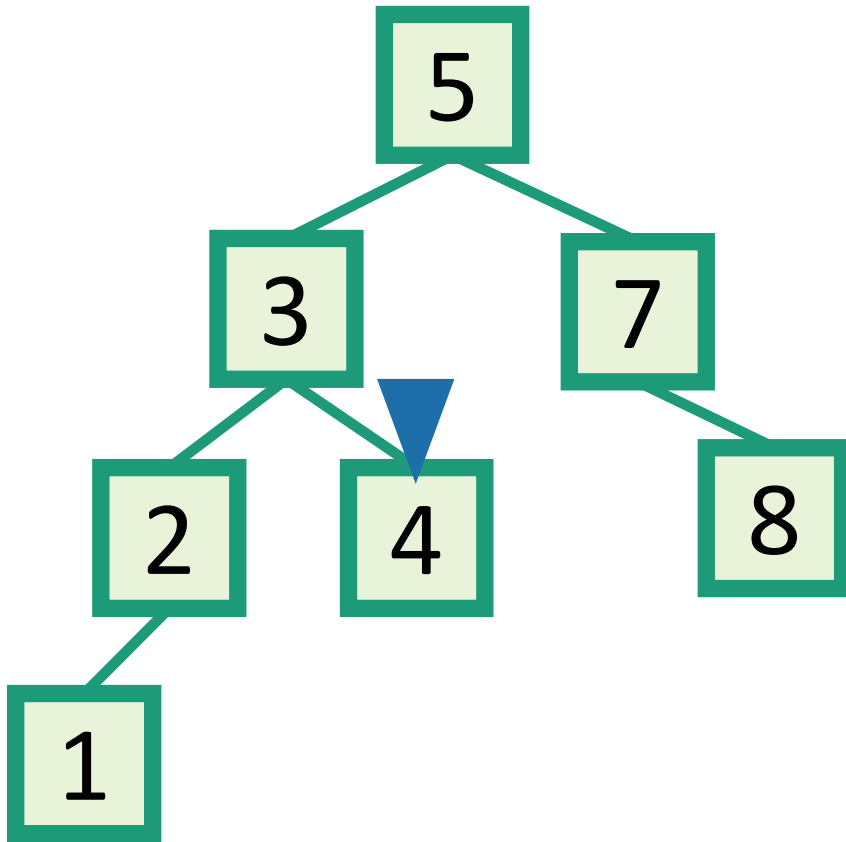
- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$



INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5

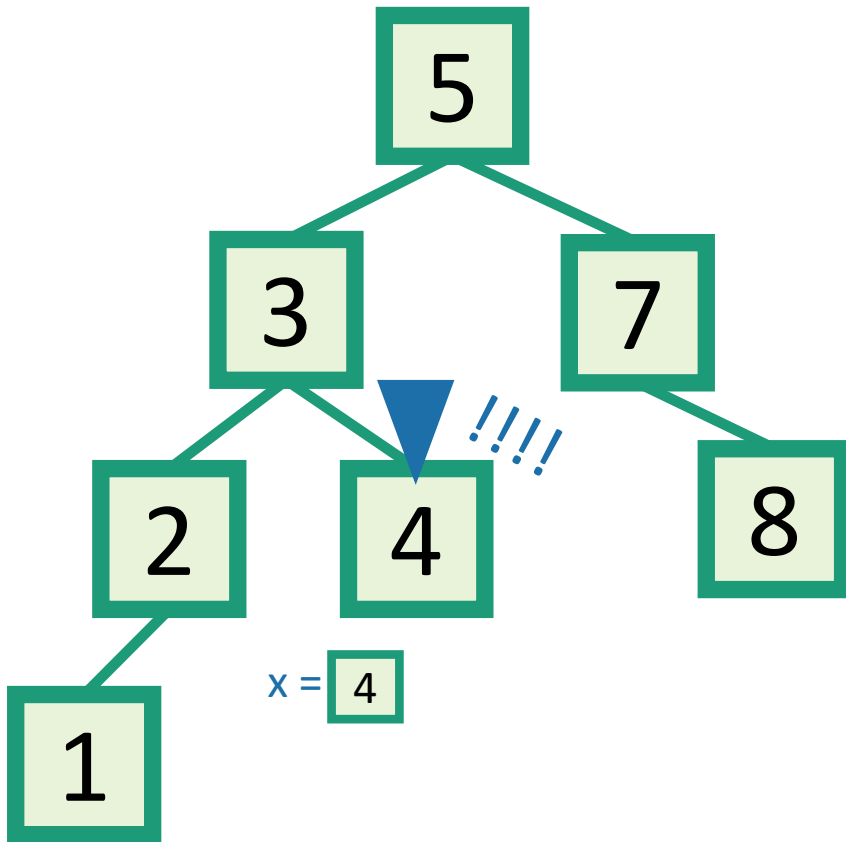
- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$



INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5

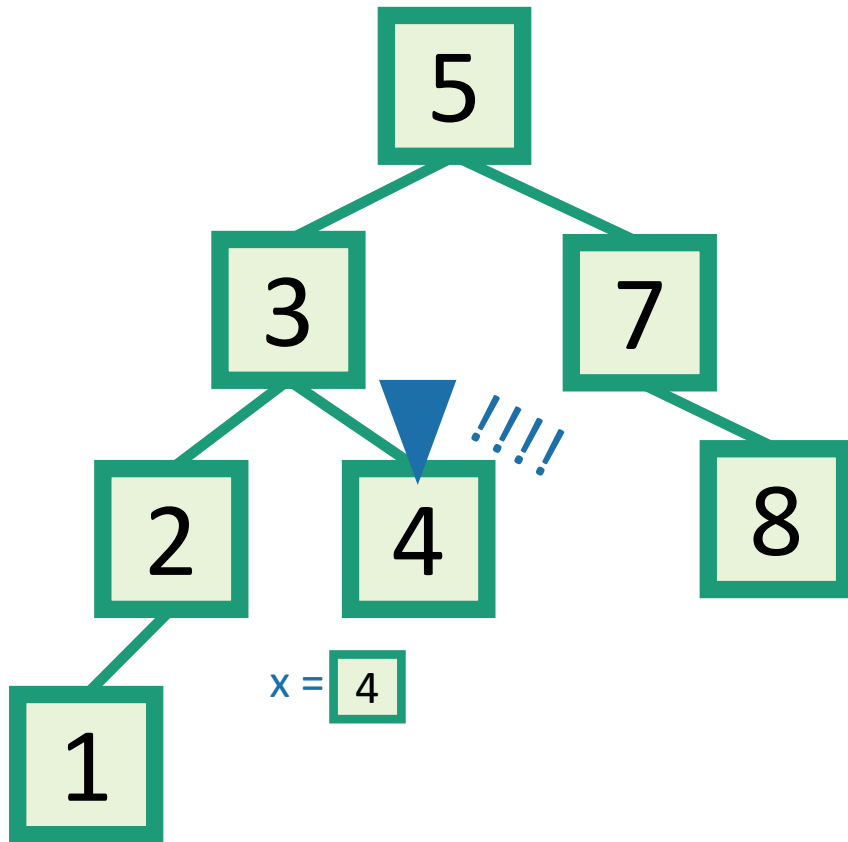
- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$



INSERT in a Binary Search Tree

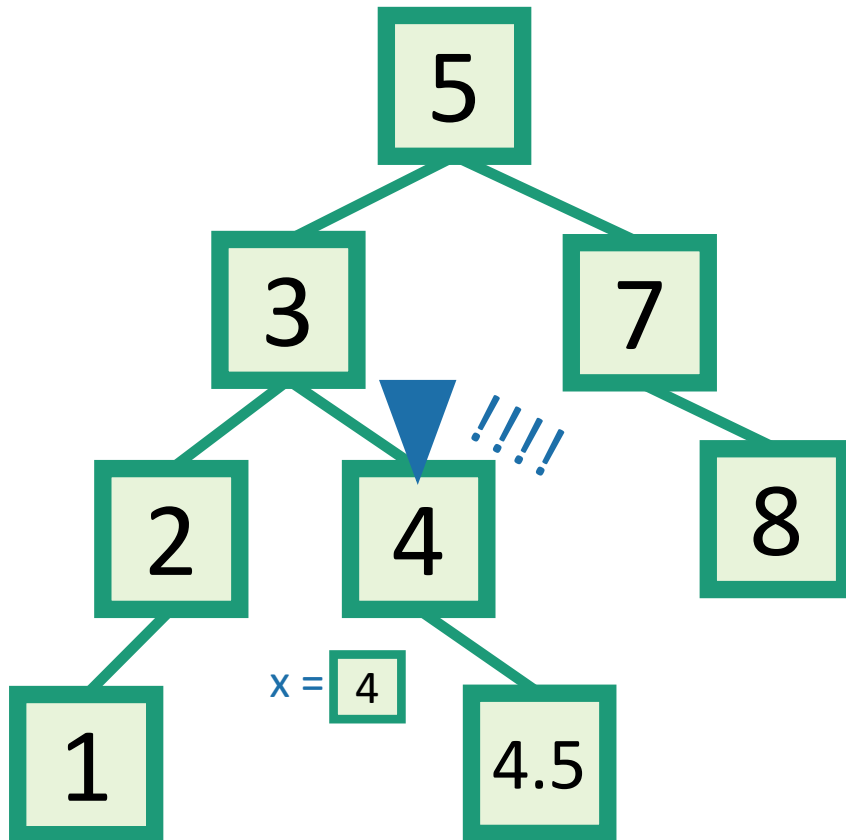
EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **Insert** a new node with desired key at x ...



INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5

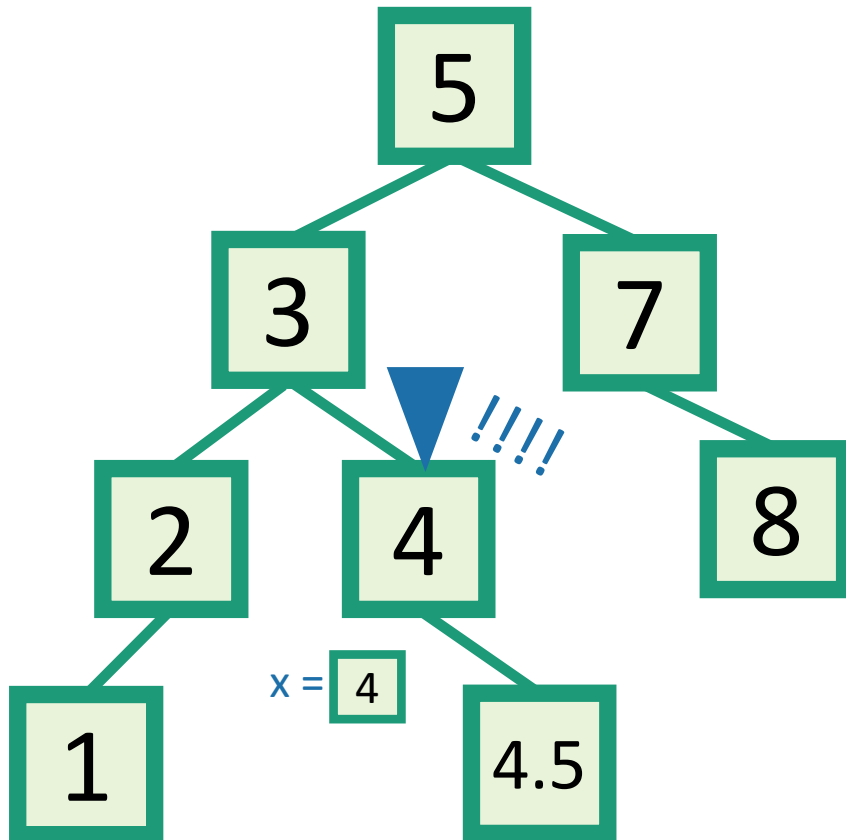


- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **Insert** a new node with desired key at x ...



INSERT in a Binary Search Tree

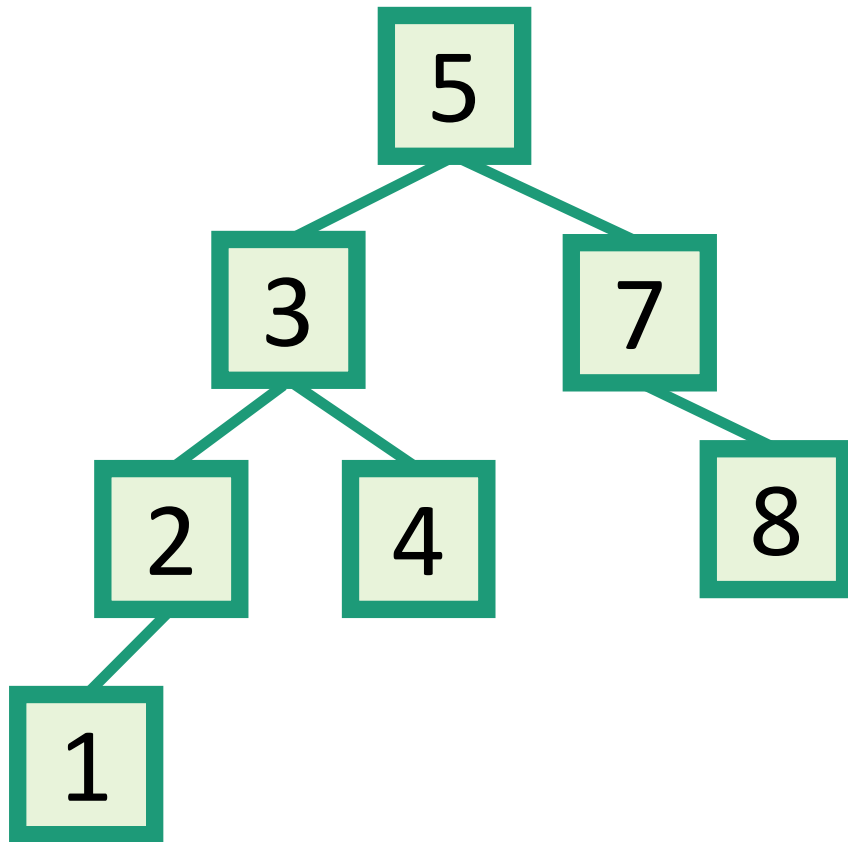
EXAMPLE: Insert 4.5



- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $\text{key} > x.\text{key}$:
 - Make a new node with the correct key, and put it as the right child of x .
 - **if** $\text{key} < x.\text{key}$:
 - Make a new node with the correct key, and put it as the left child of x .
 - **if** $x.\text{key} == \text{key}$:
 - **return**

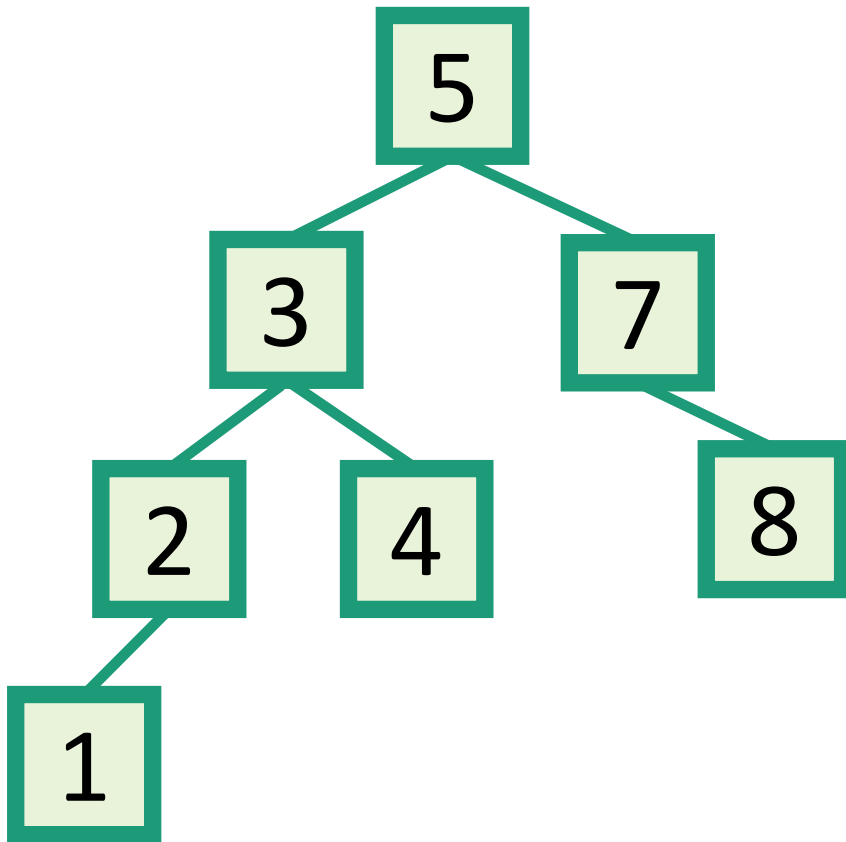


DELETE in a Binary Search Tree



DELETE in a Binary Search Tree

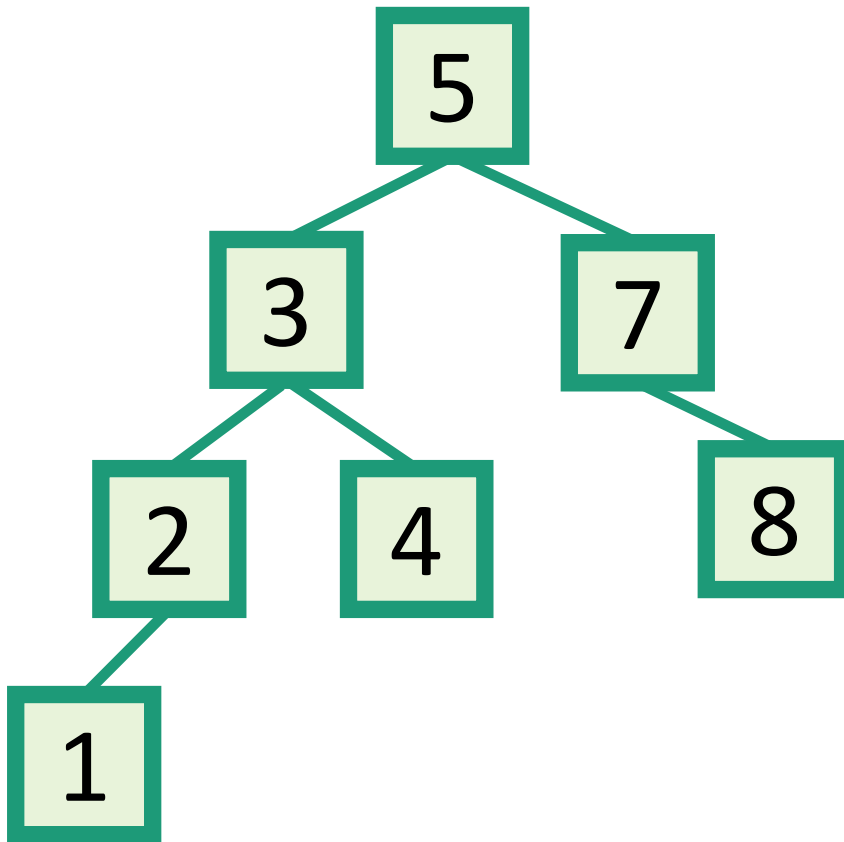
EXAMPLE: Delete 2



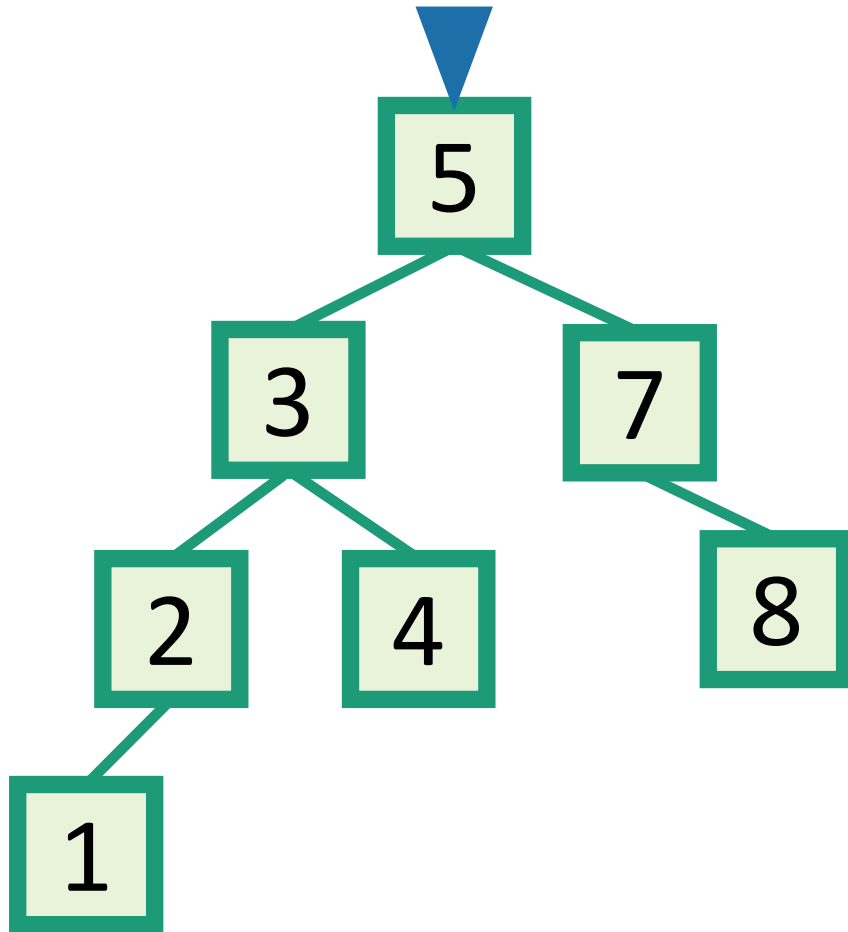
DELETE in a Binary Search Tree

EXAMPLE: Delete 2

- **DELETE**(key):



DELETE in a Binary Search Tree



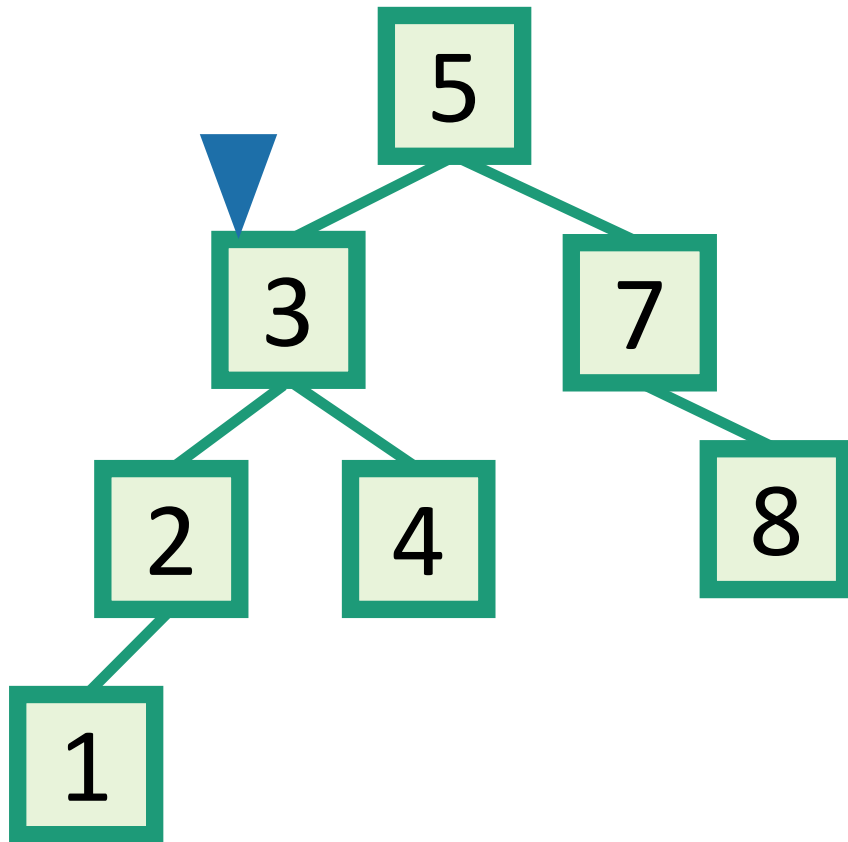
EXAMPLE: Delete 2

- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$



DELETE in a Binary Search Tree

EXAMPLE: Delete 2



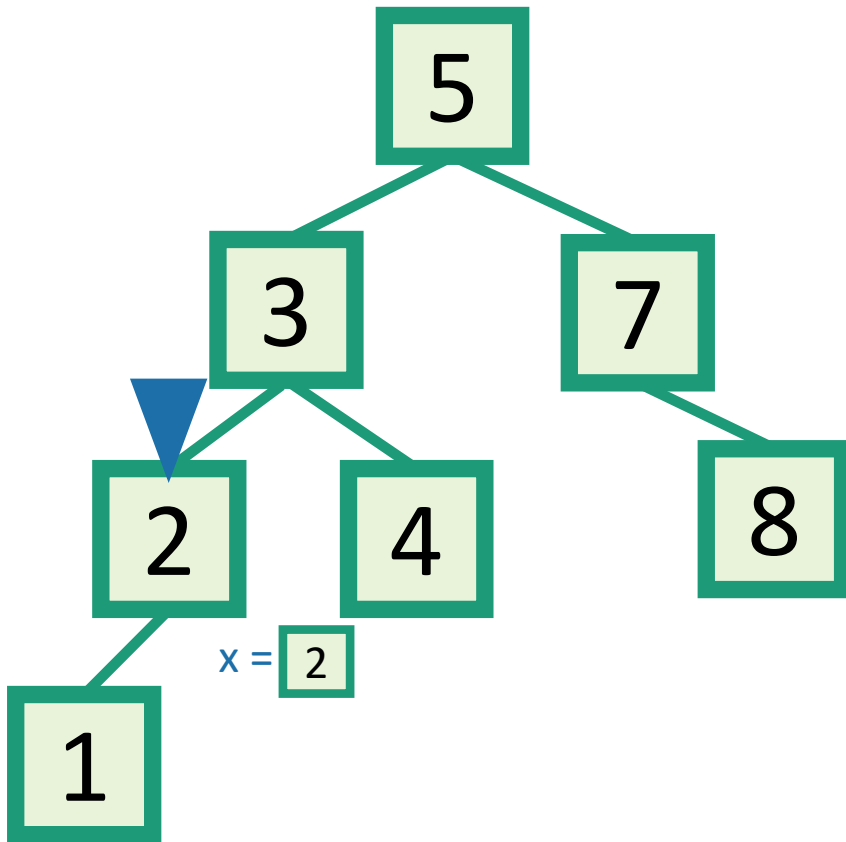
- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$



DELETE in a Binary Search Tree

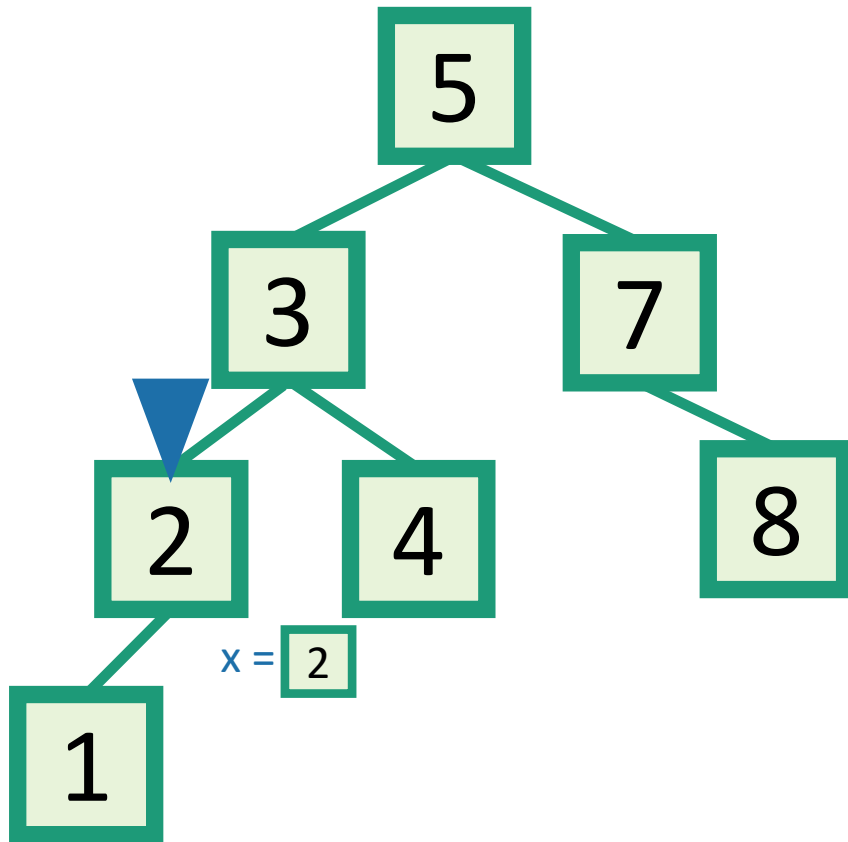
EXAMPLE: Delete 2

- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$



DELETE in a Binary Search Tree

EXAMPLE: Delete 2

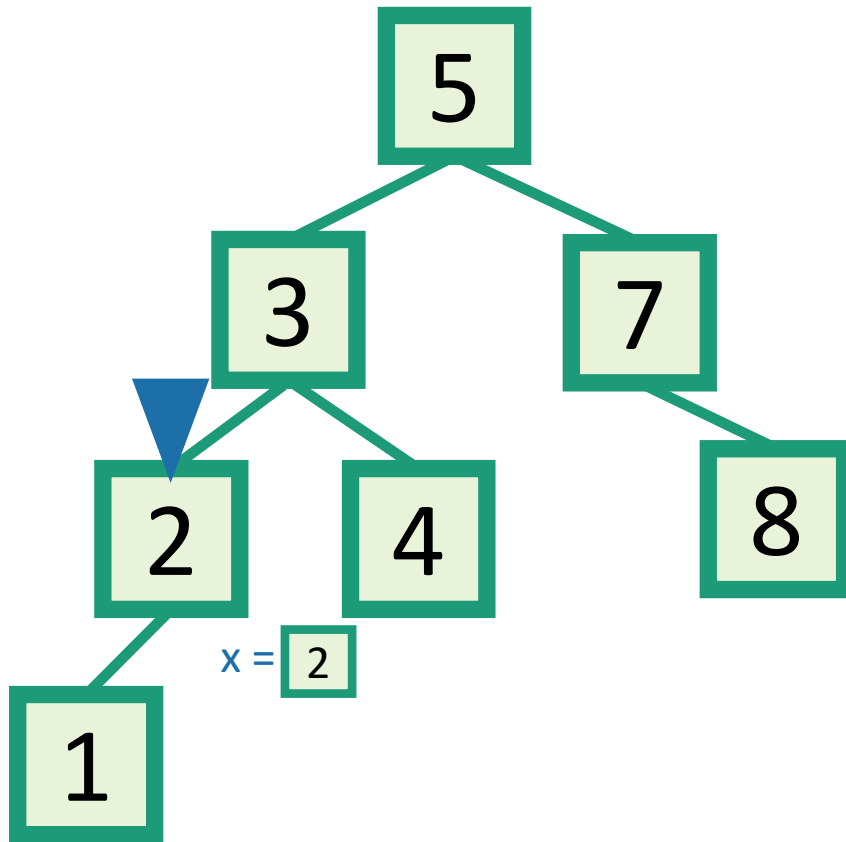


- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:



DELETE in a Binary Search Tree

EXAMPLE: Delete 2

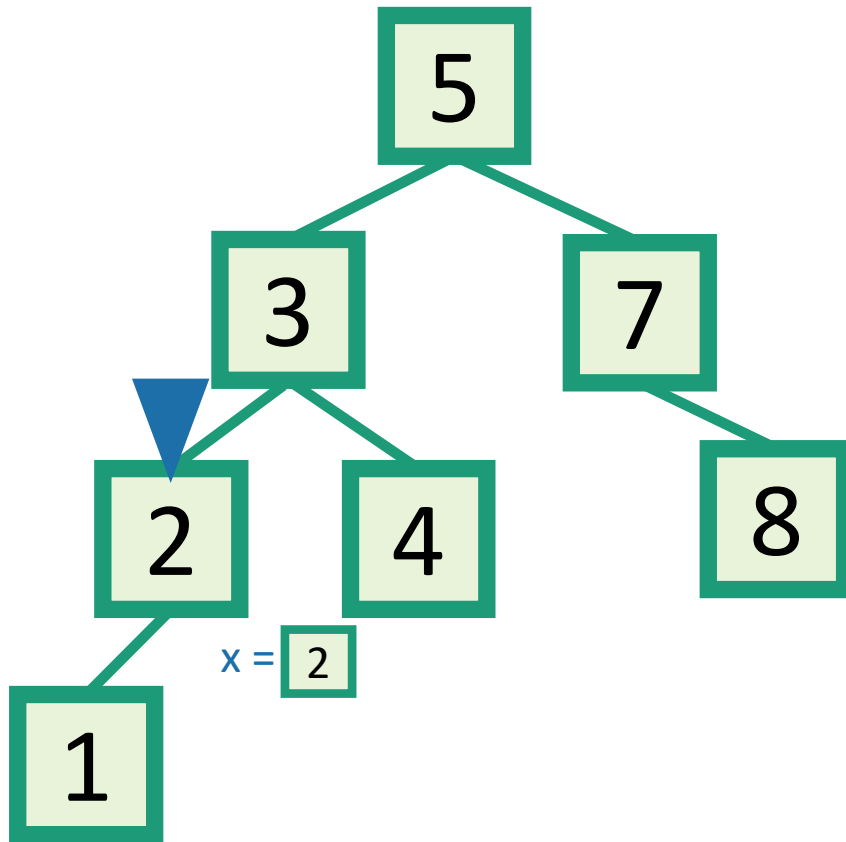


- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x



DELETE in a Binary Search Tree

EXAMPLE: Delete 2



- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x

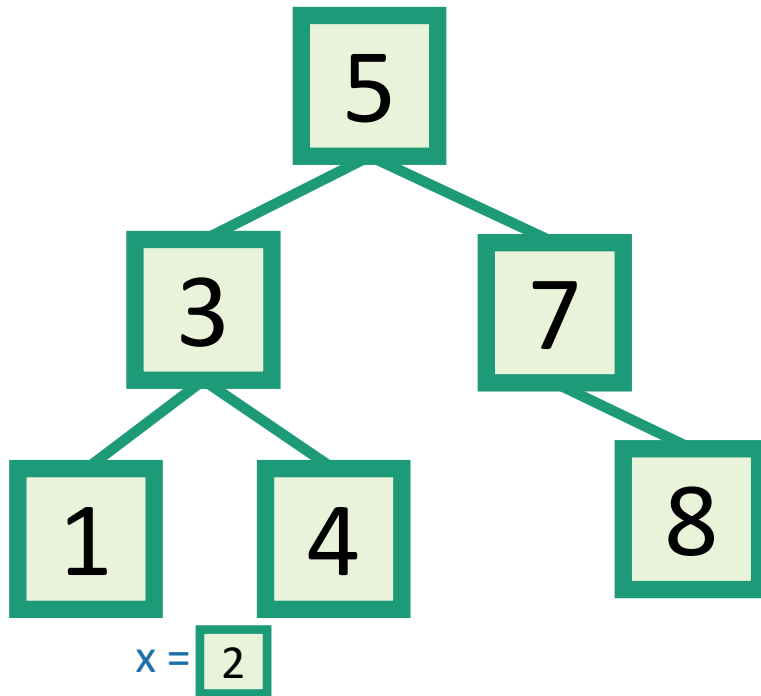
↖

This is a bit more complicated...see the slides for some pictures of the different cases.



DELETE in a Binary Search Tree

EXAMPLE: Delete 2



- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x....

↖

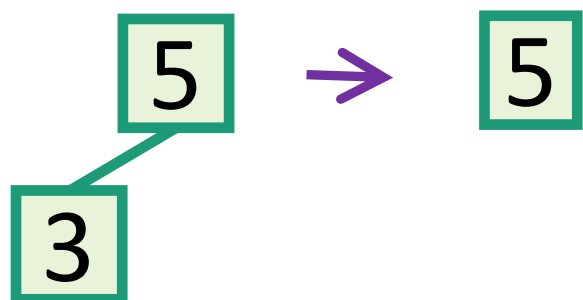
This is a bit more complicated...see the slides for some pictures of the different cases.



DELETE in a Binary Search Tree

several cases (by example)

say we want to delete 3

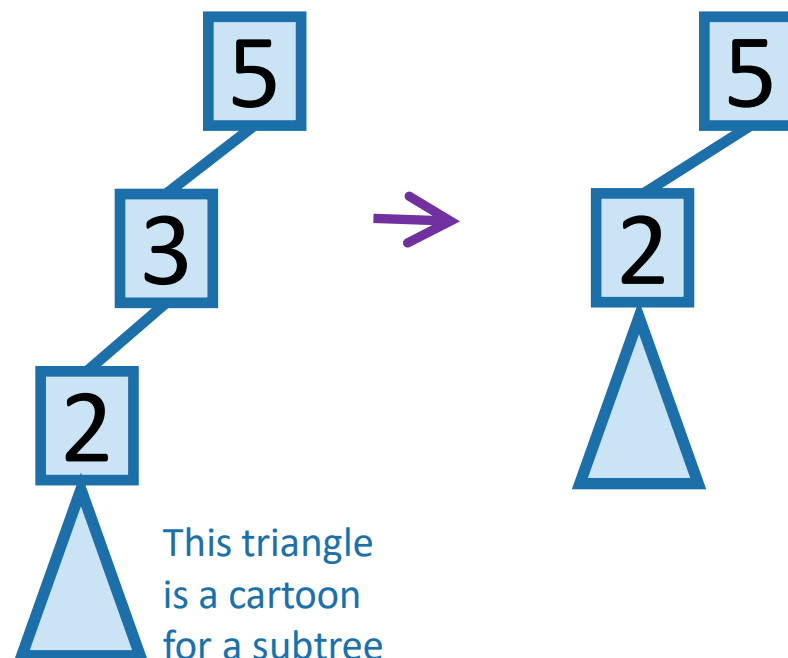


Case 1: if 3 is a leaf,
just delete it.

Write pseudocode for all of these!



Siggie the Studious Stork



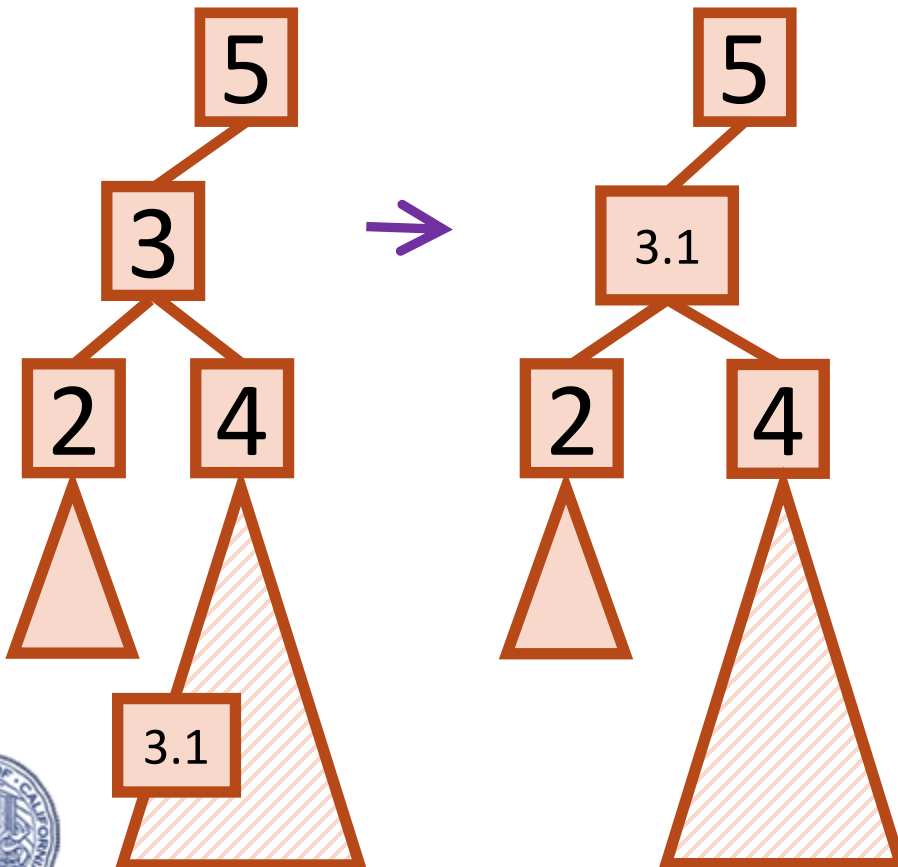
Case 2: if 3 has just one child,
move that up.



DELETE in a Binary Search Tree

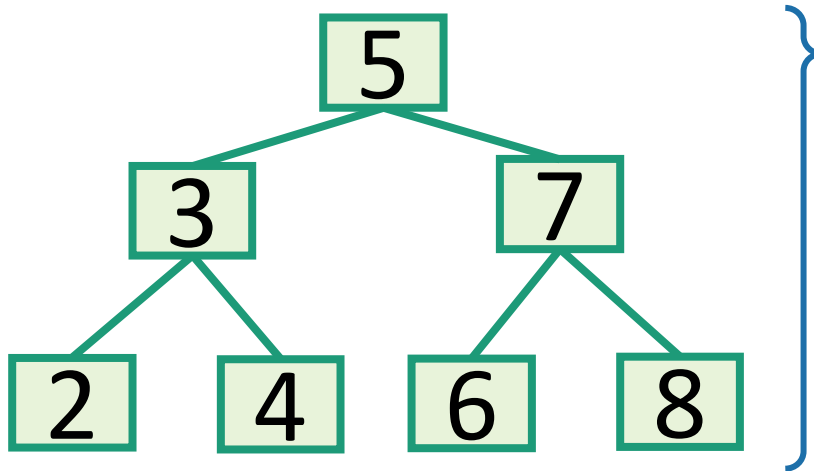
Case 3: if 3 has two children,
replace 3 with its **immediate successor**.
(aka, next biggest thing after 3)

- Does this maintain the BST property?
 - Yes.
- How do we find the immediate successor?
 - SEARCH for 3 in the subtree under 3.right
- How do we remove it when we find it?
 - If [3.1] has 0 or 1 children, do one of the previous cases.
- What if [3.1] has two children?
 - It doesn't.



How long do these operations take?

- **SEARCH** is the big one.
 - Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.



Time = $O(\text{height of tree})$

Trees have depth $O(\log(n))$. **Done!**



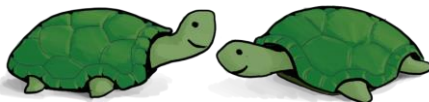
Lucky the
Lackadaisical Lemur

Wait a
second...

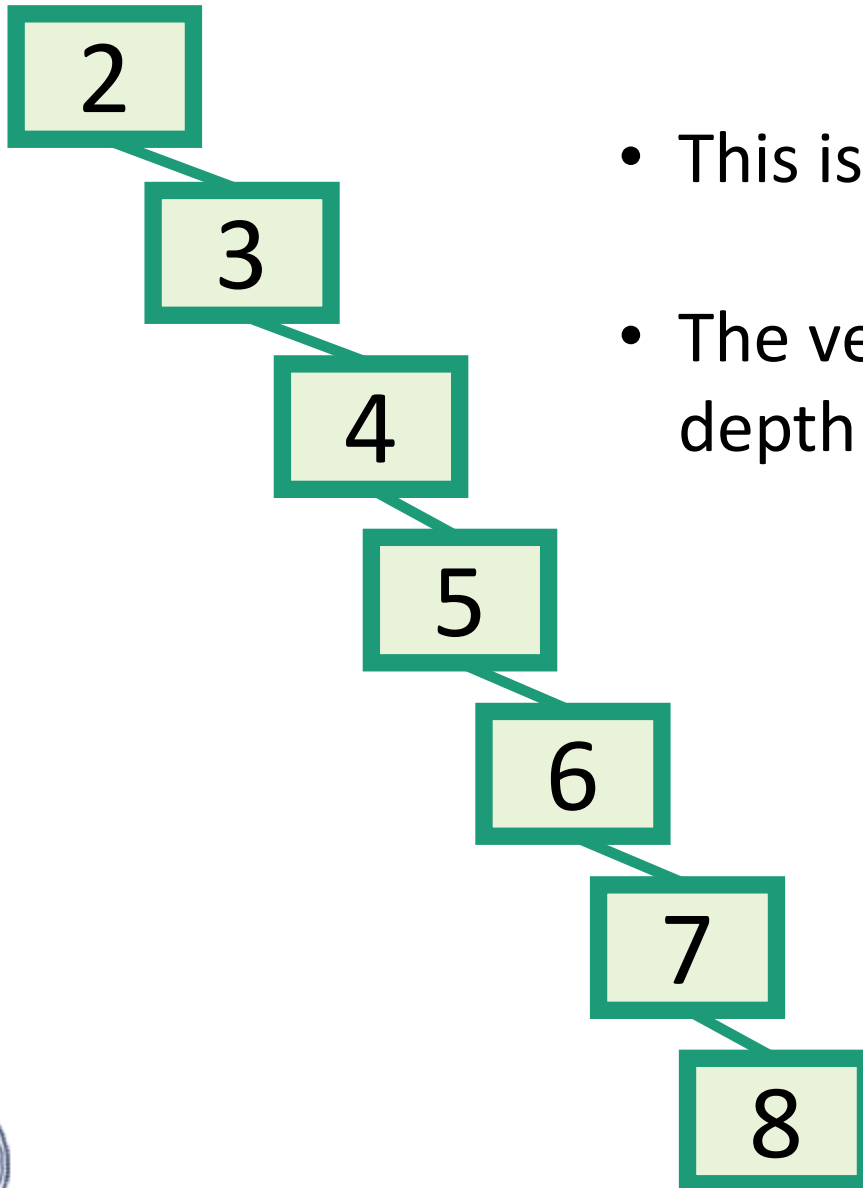


Plucky the
Pedantic Penguin

How long does search take?



Search might take time $O(n)$.



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.



What to do?

How often is “every so often” in the worst case?
It’s actually pretty often!



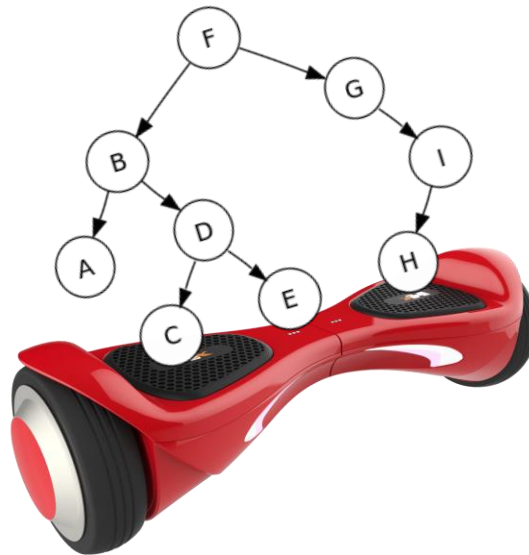
Ollie the over-achieving ostrich

- Goal: Fast **SEARCH/INSERT/DELETE**
- All these things take time $O(\text{height})$
- And the height might be big!!! ☹️
- Idea 0:
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
 - At least $\Omega(n)$ every so often....

• Turns out that’s not a great idea. Instead we turn to...

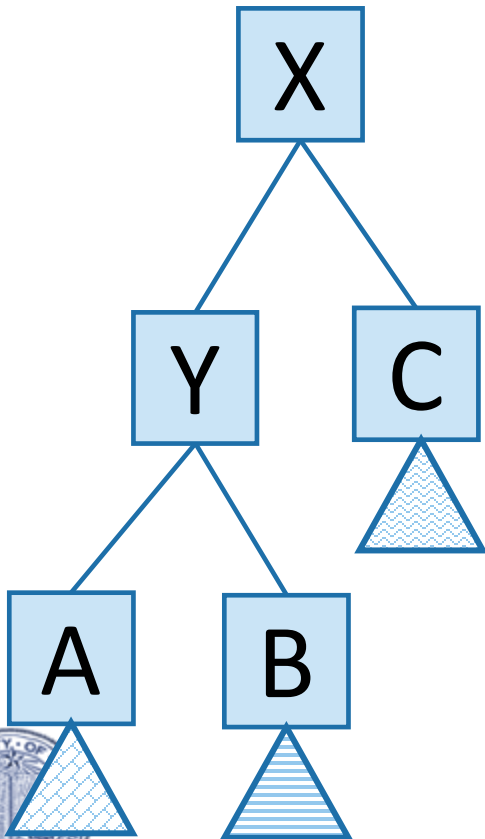


Self-Balancing Binary Search Trees



Idea 1: Rotations

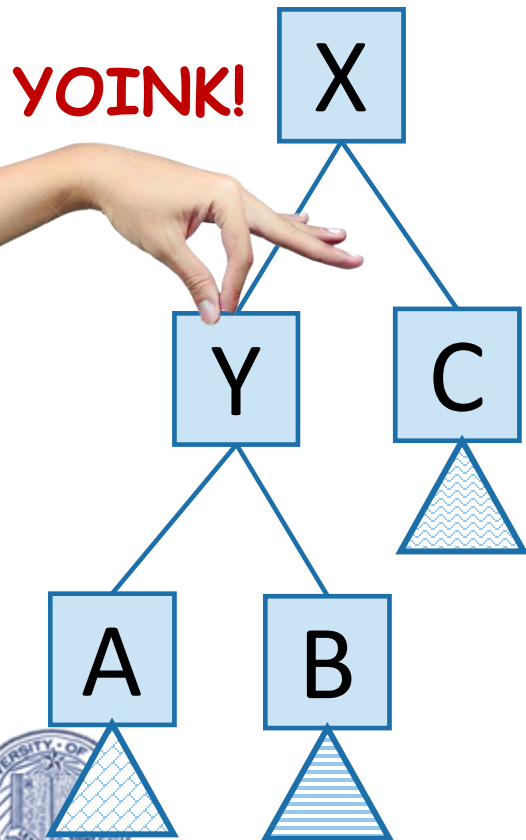
- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

Idea 1: Rotations

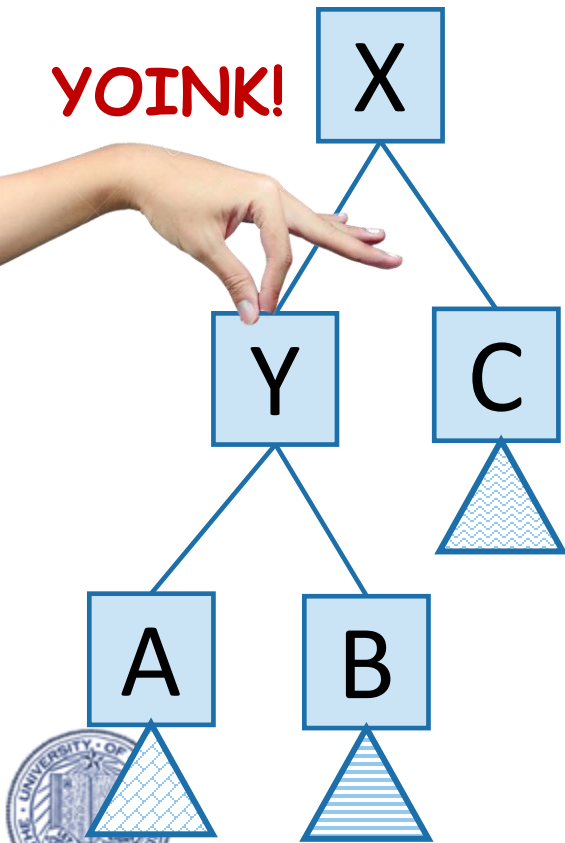
- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

Idea 1: Rotations

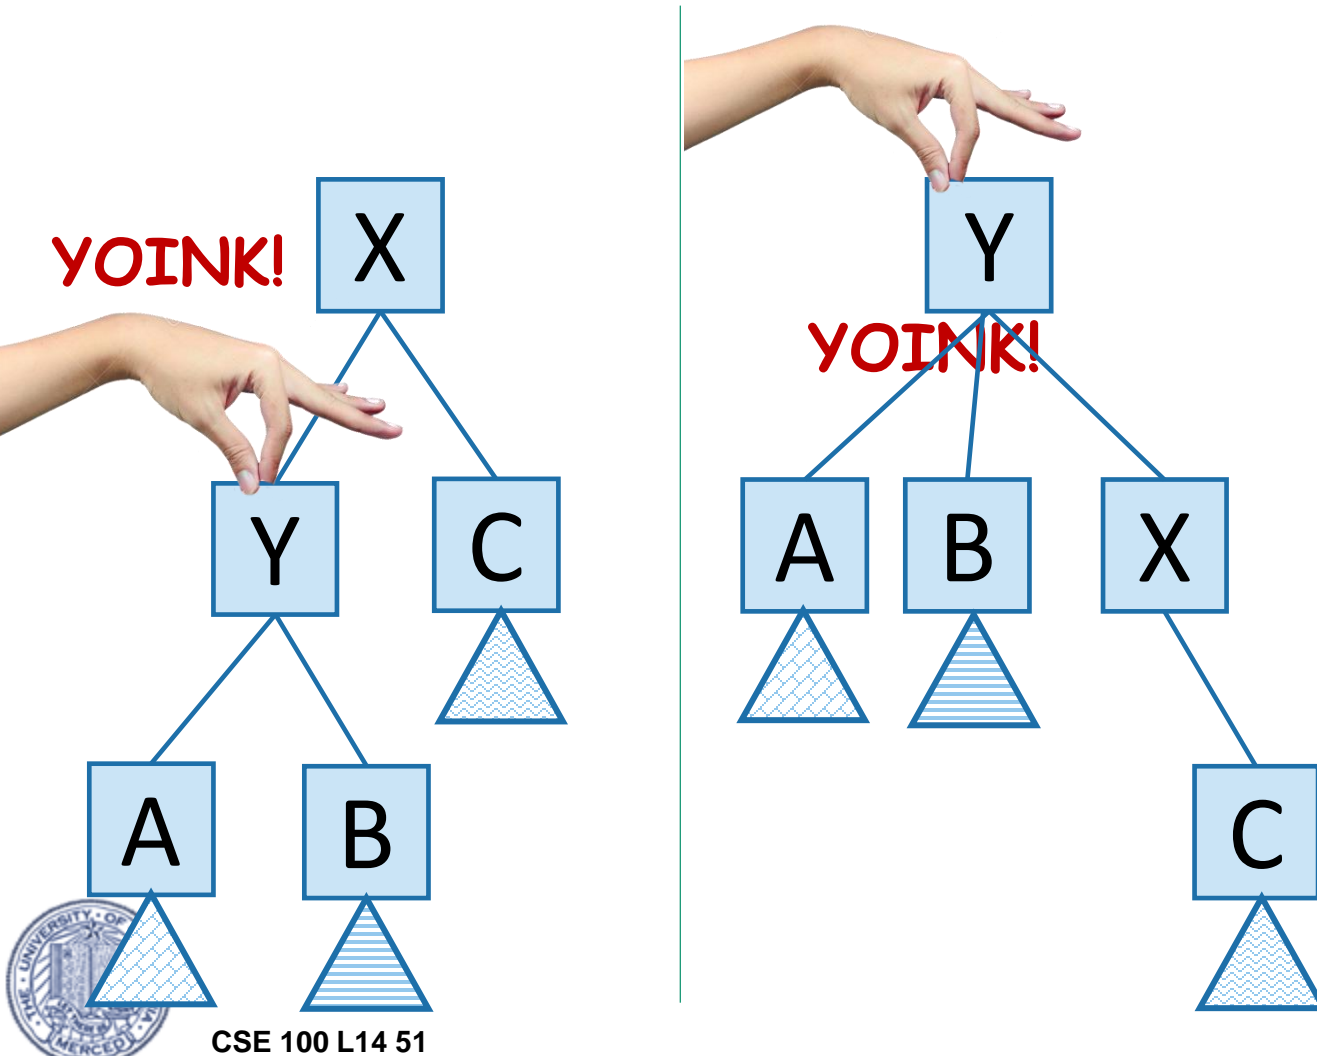
- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

Idea 1: Rotations

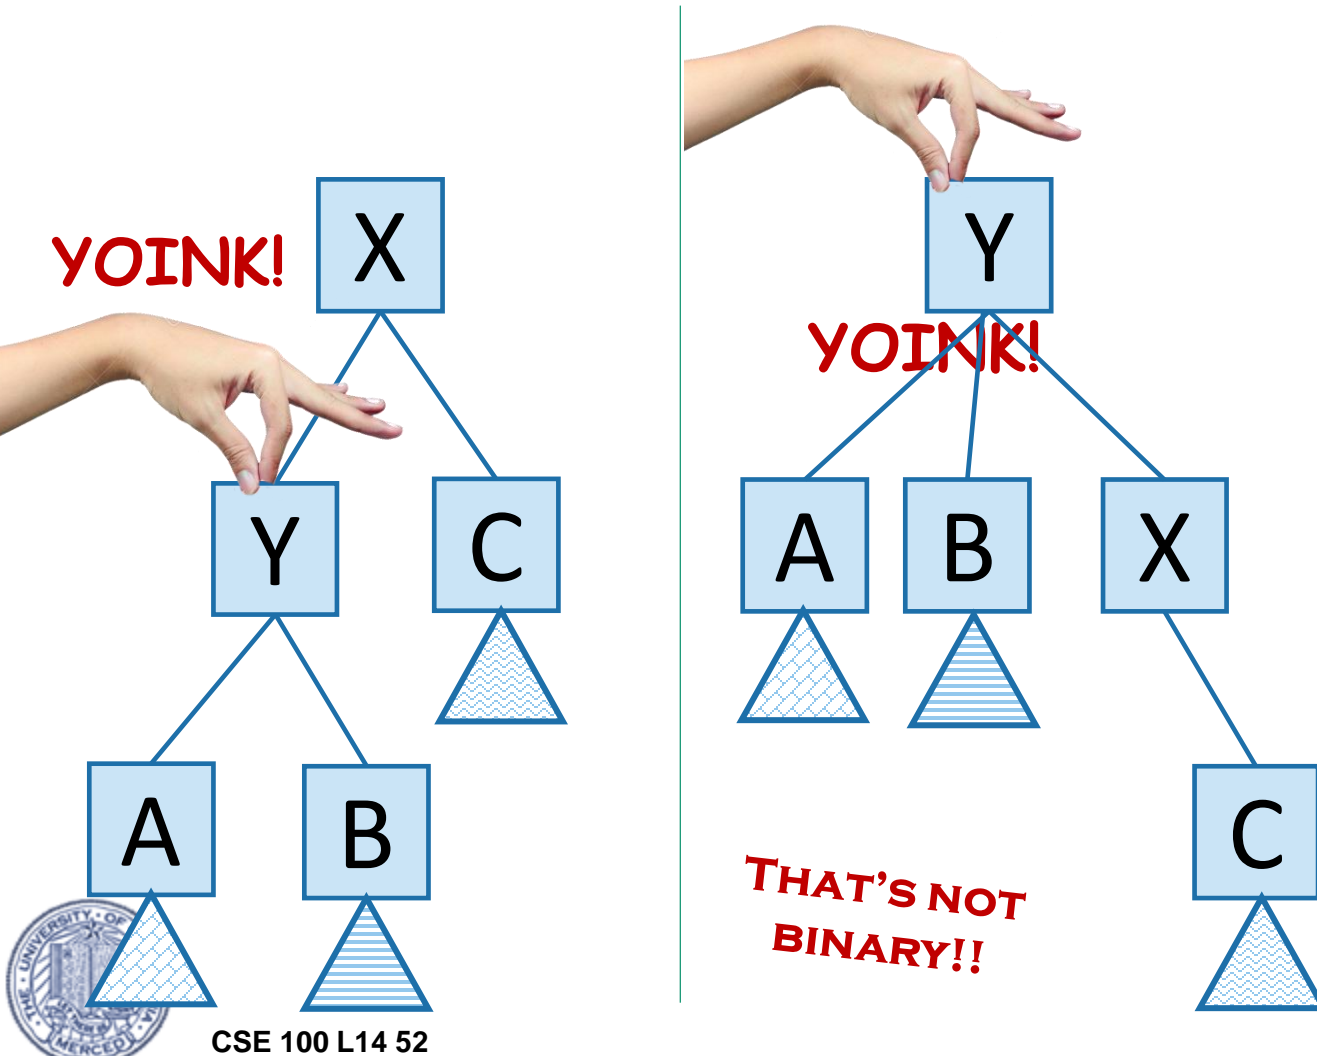
- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

Idea 1: Rotations

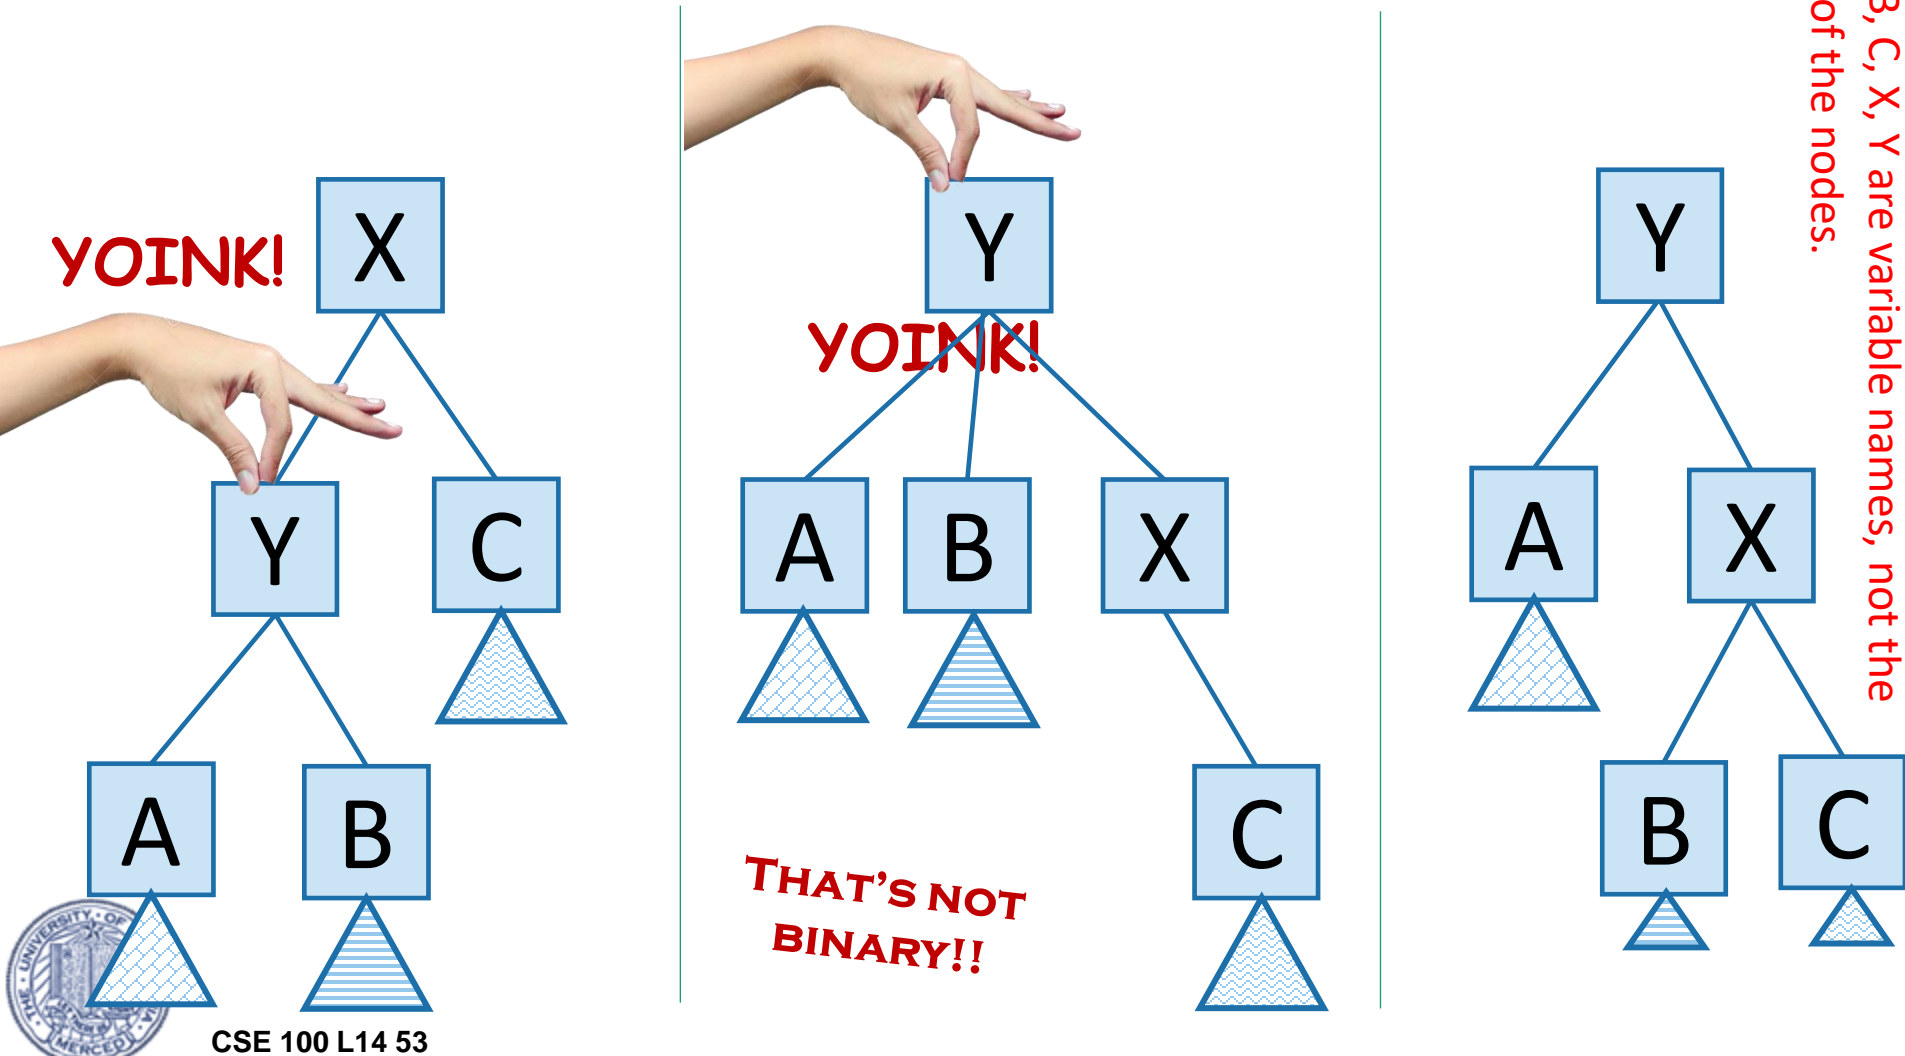
- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

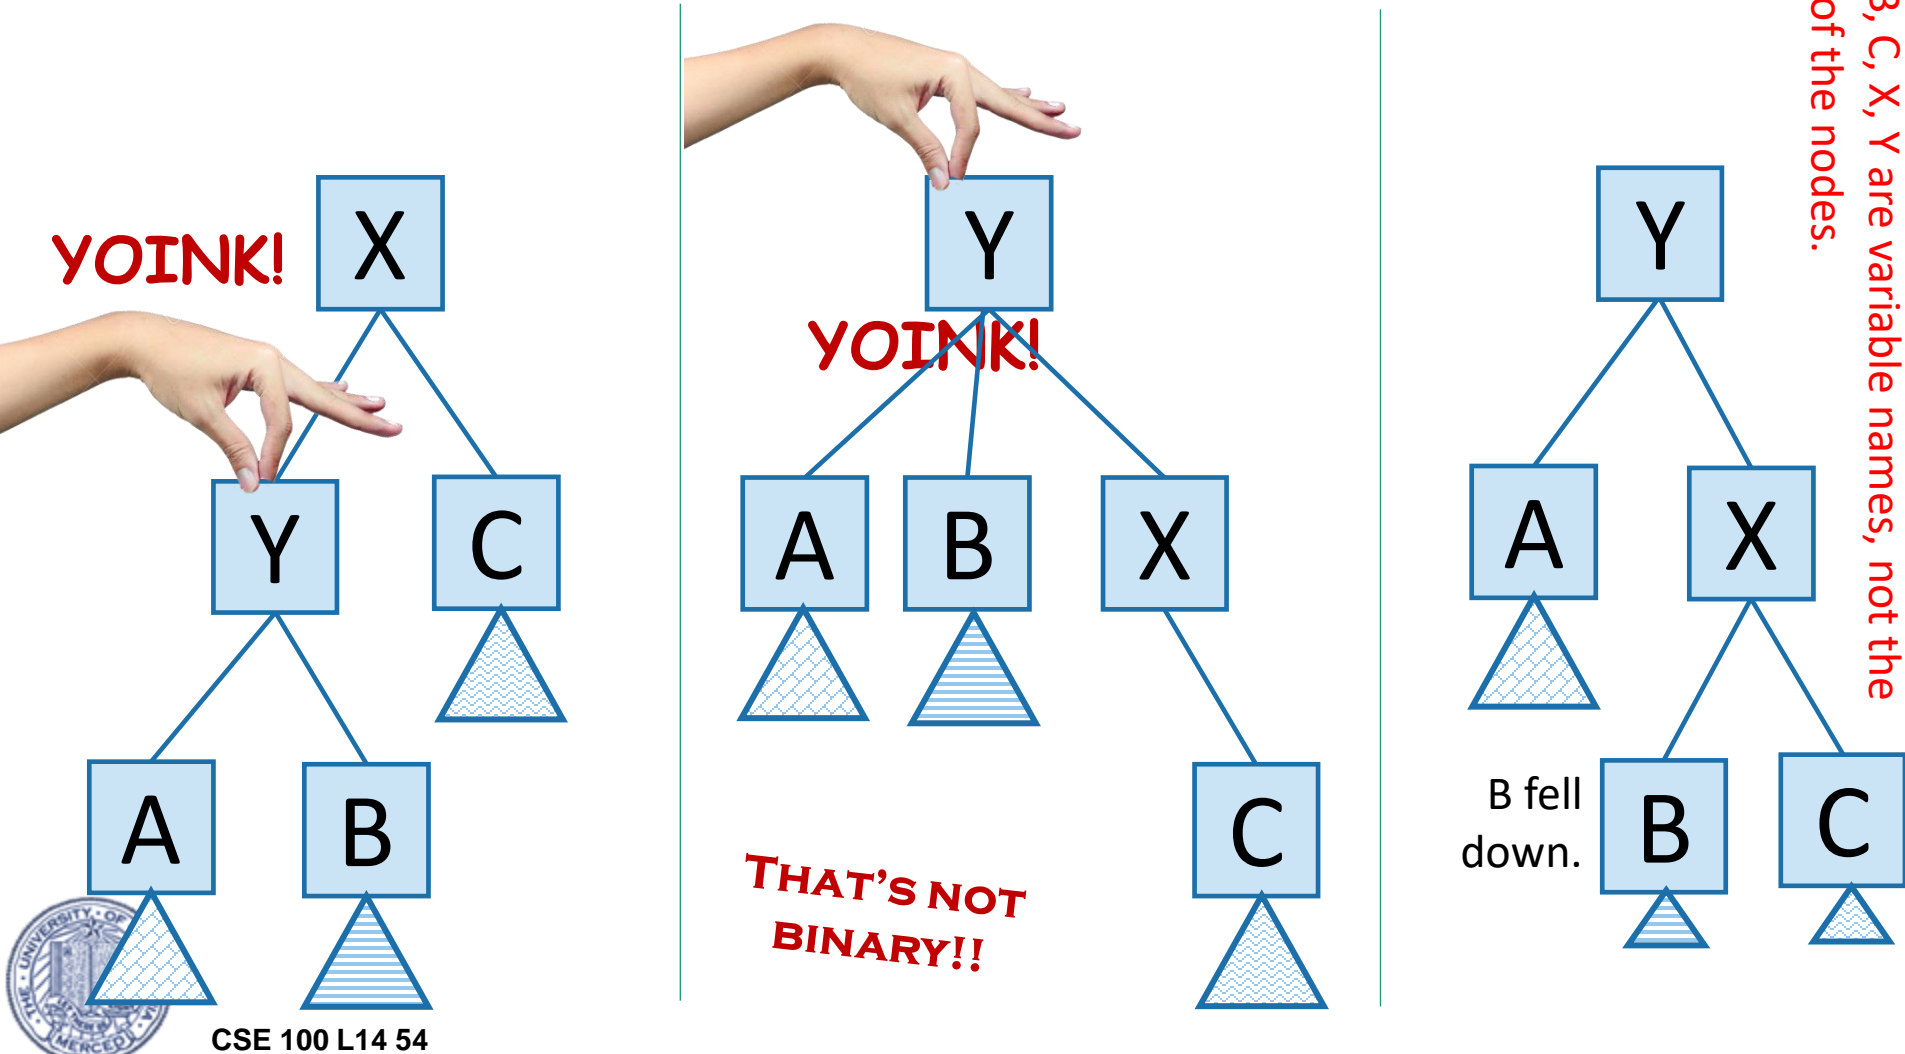
Idea 1: Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.



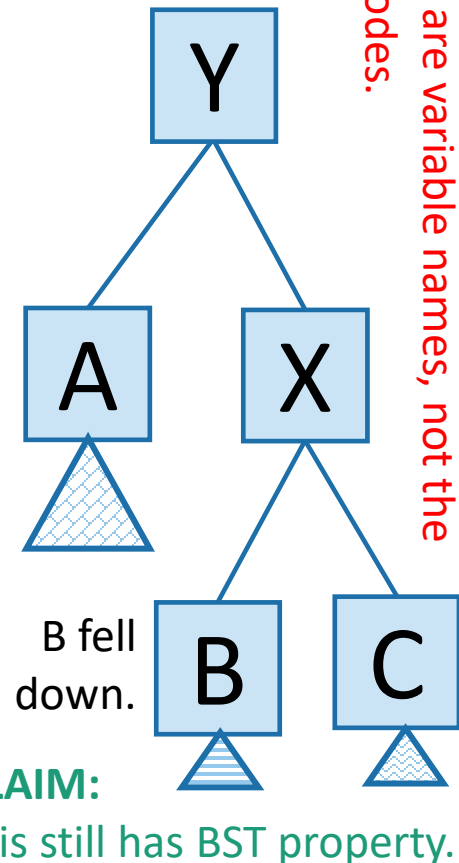
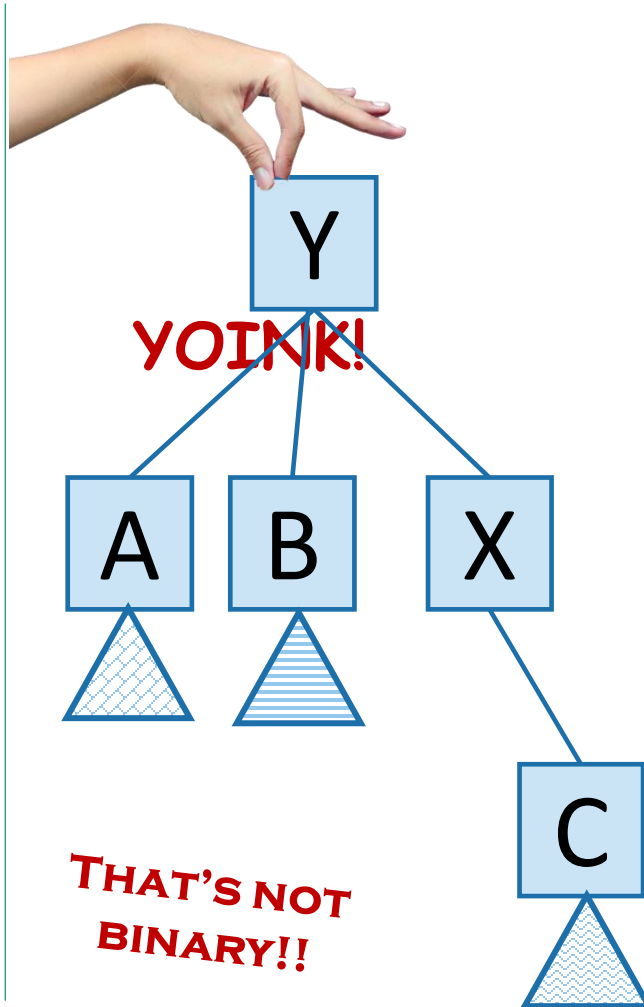
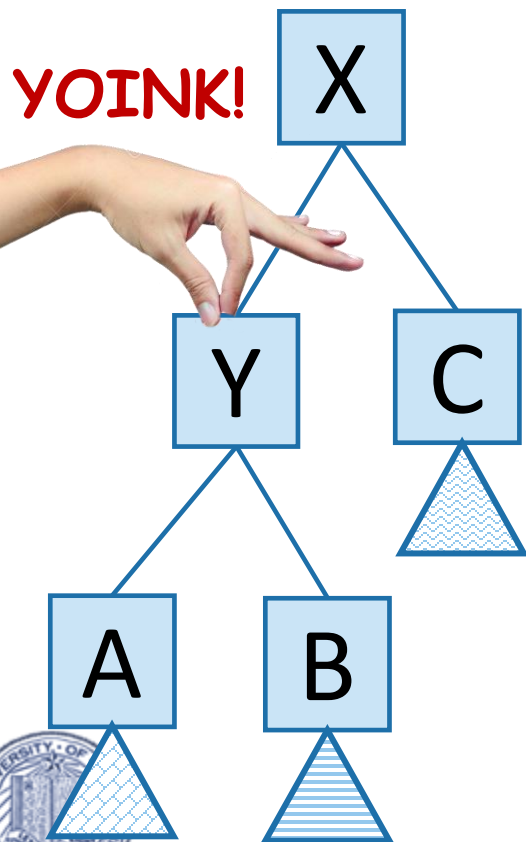
Idea 1: Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.



Idea 1: Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.



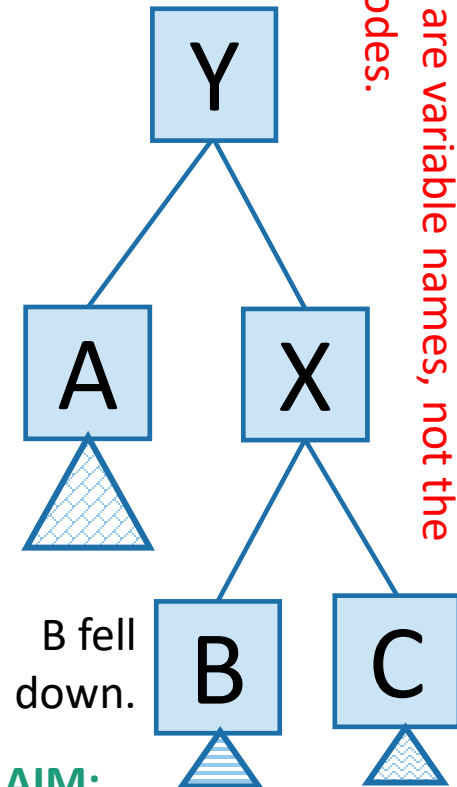
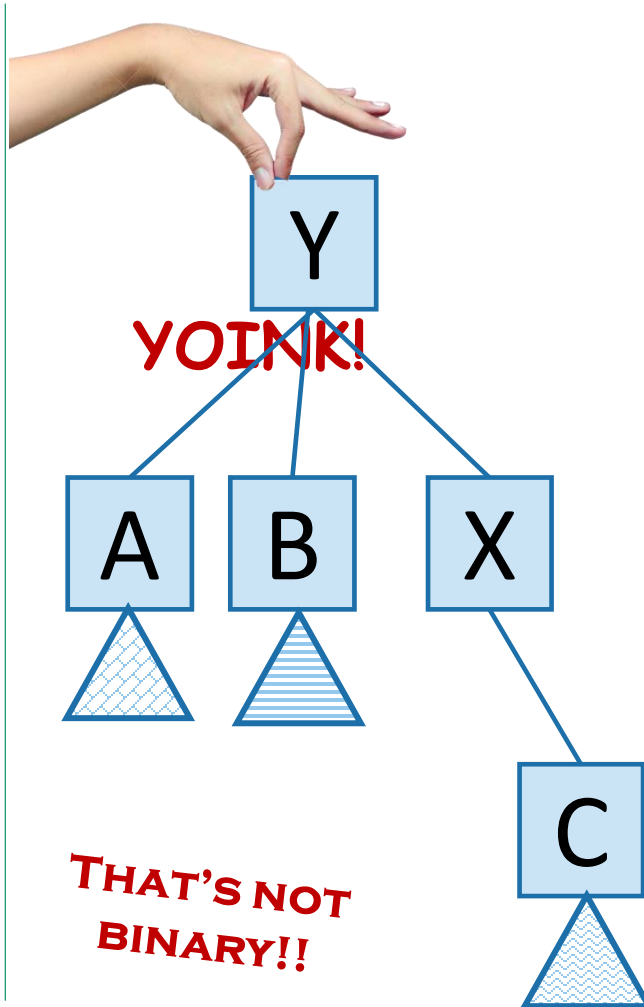
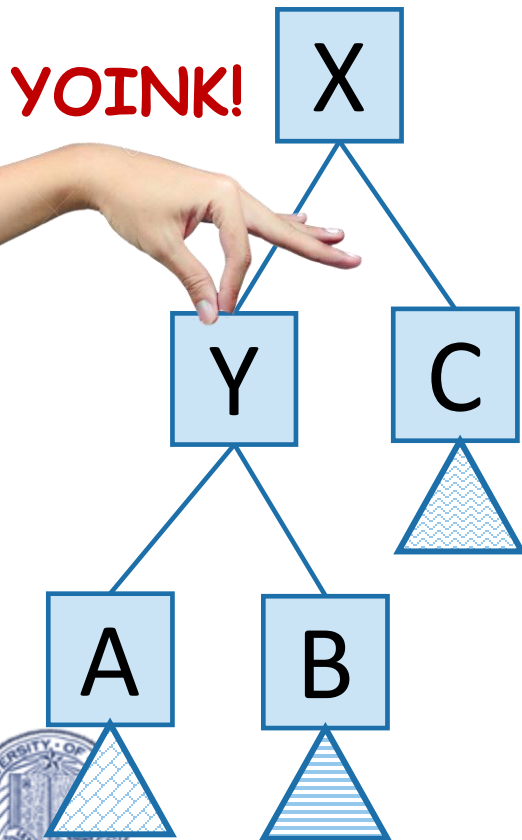
Note: A, B, C, X, Y are variable names, not the contents of the nodes.

Idea 1: Rotations

No matter what lives underneath A,B,C,
this takes time $O(1)$. (Why?)

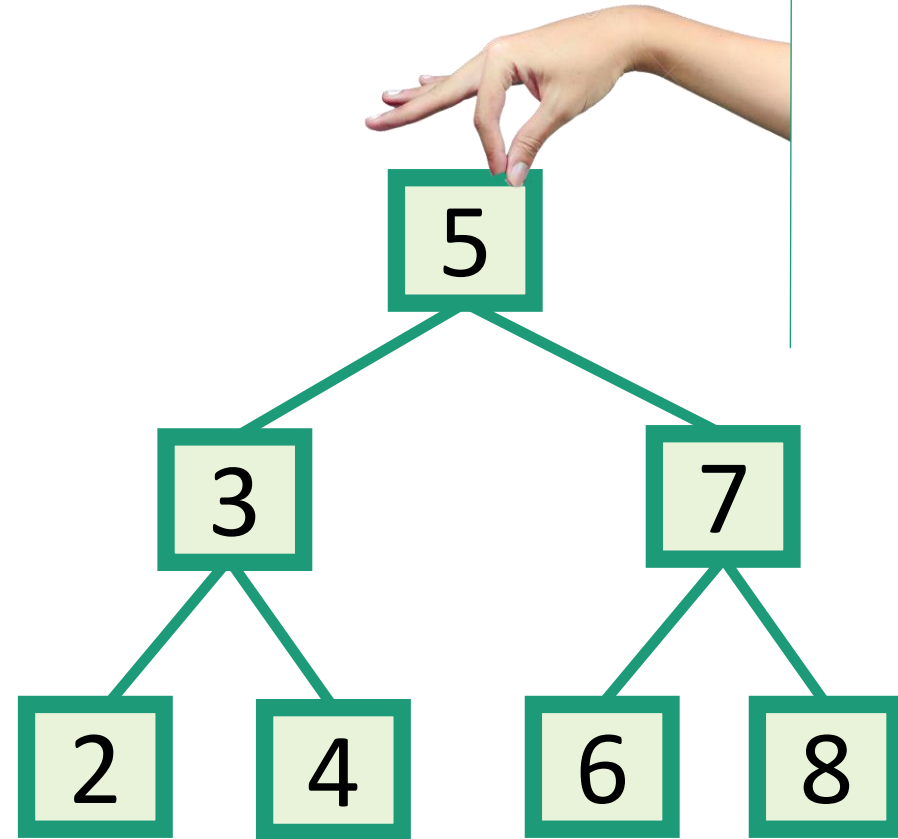
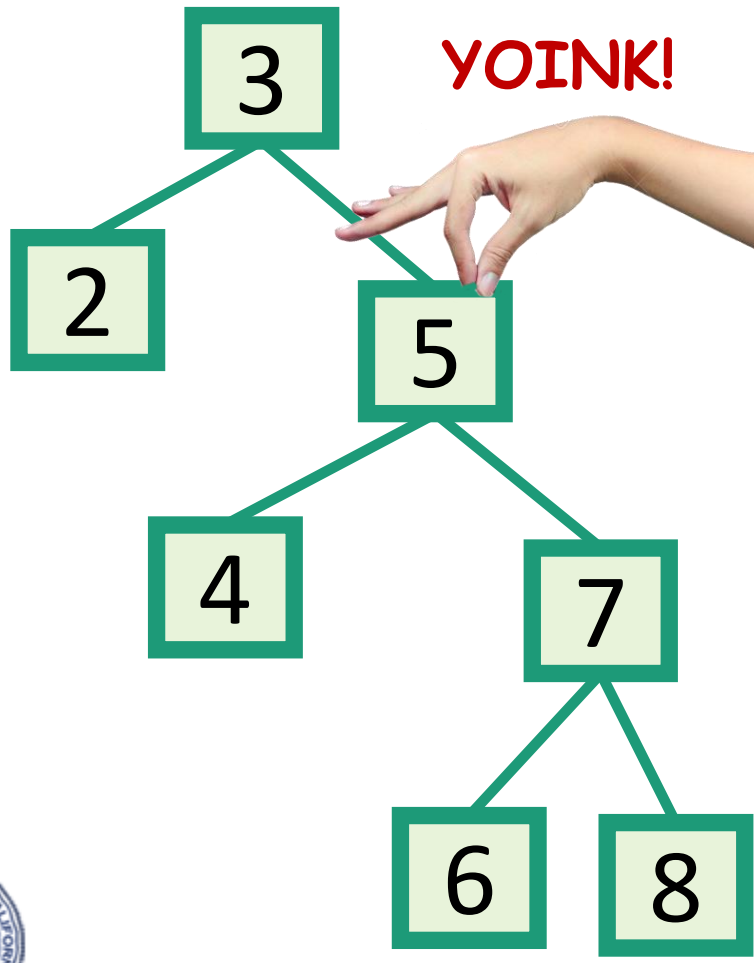
- Maintain Binary Search Tree (BST) property, while moving stuff around.

Note: A, B, C, X, Y are variable names, not the contents of the nodes.



CLAIM:
this still has BST property.

This seems helpful



Strategy?

- Whenever something seems unbalanced, do rotations until it's okay again.



Even for Lucky this is pretty vague.
What do we mean by “seems unbalanced”? What’s “okay”?

Lucky the Lackadaisical Lemur



Idea 2: have some proxy for balance

- Maintaining **perfect balance** is too hard.
- Instead, come up with some **proxy for balance**:
 - If the tree satisfies **[SOME PROPERTY]**, then it's pretty balanced.
 - We can maintain **[SOME PROPERTY]** using rotations.



There are actually several ways to do this, but today we'll see...



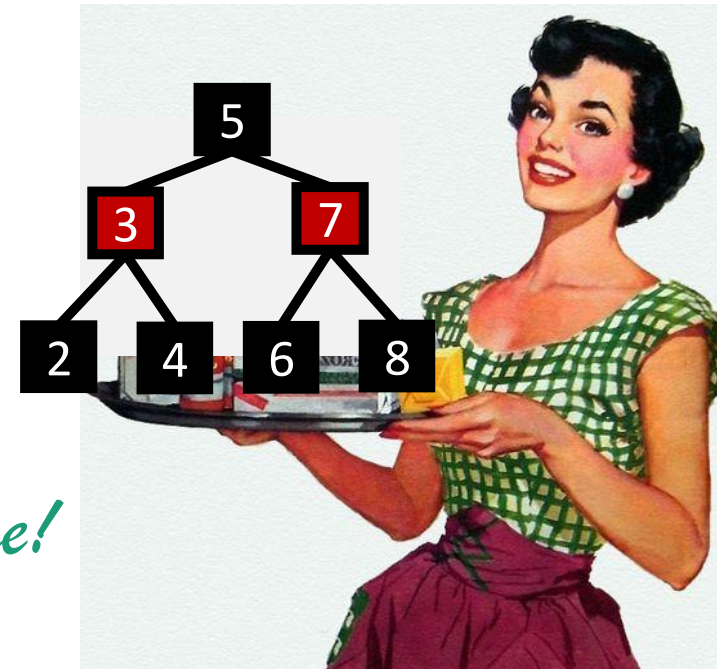
Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!
- Be the envy of your friends and neighbors with the time-saving...

Red-Black tree!

Maintain balance by stipulating that
black nodes are balanced, and
that there aren't too many **red**
nodes.

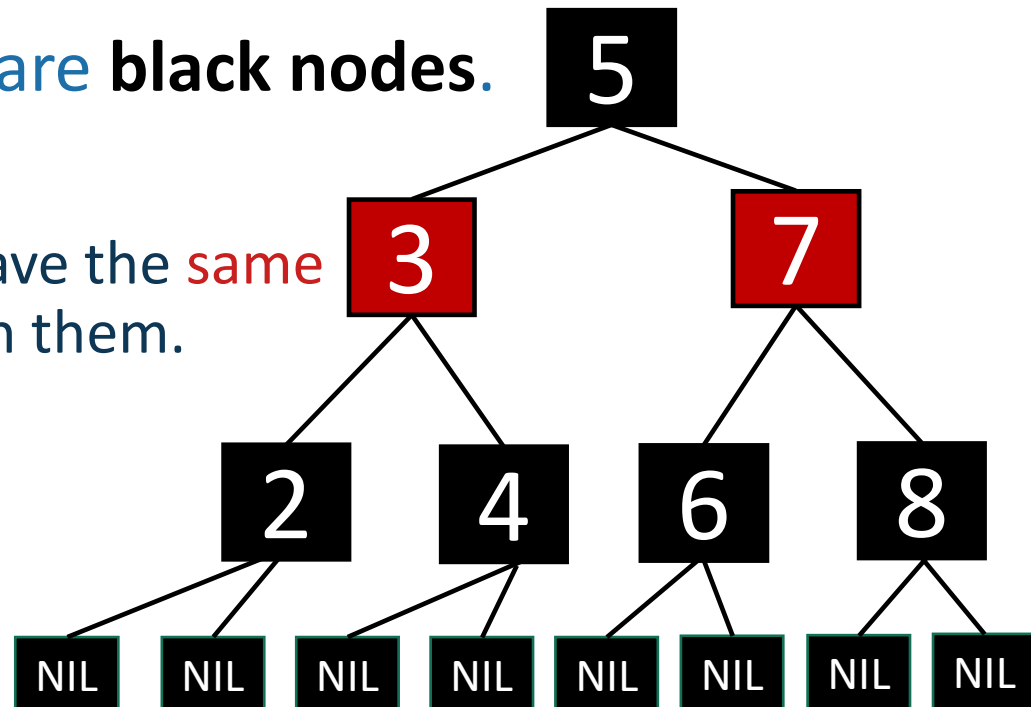
It's just good sense!



Red-Black Trees

obey the following rules (which are a proxy for balance)

1. Every node is colored **red** or **black**.
2. The root node is a **black node**.
3. NIL children count as **black nodes**.
4. Children of a **red node** are **black nodes**.
5. For all nodes x:
 - all paths from x to NIL's have the same number of **black nodes** on them.

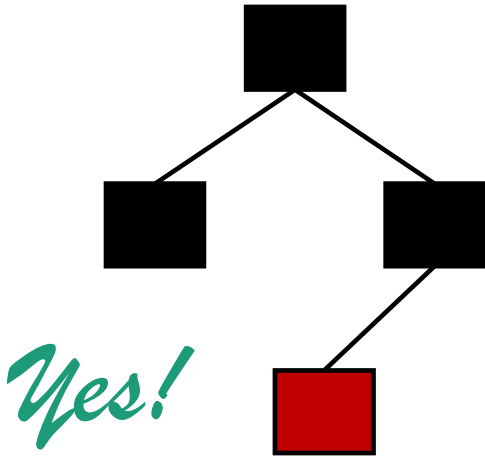


I'm not going to draw the NIL children in the future, but they are treated as black nodes.

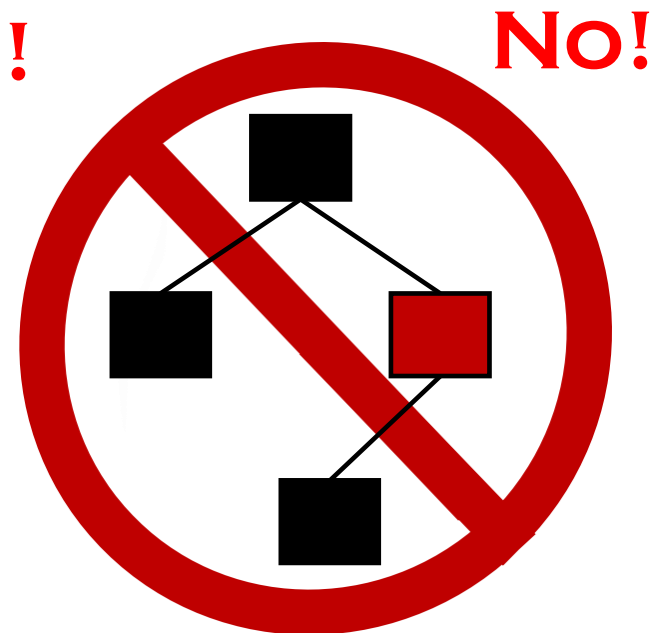
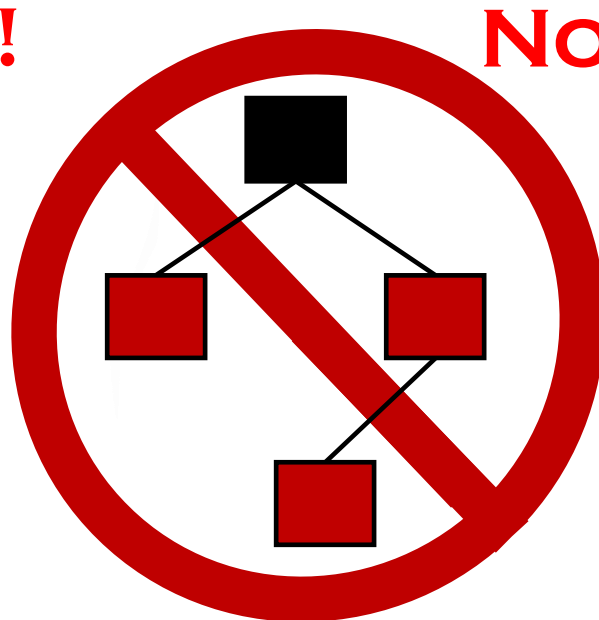
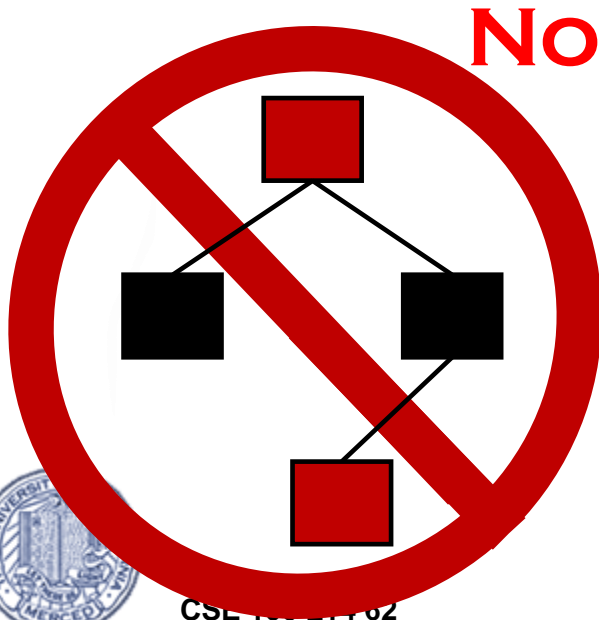
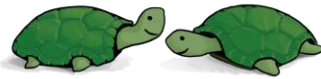


Examples(?)

1. Every node is colored **red** or **black**.
2. The root node is a **black** node.
3. NIL children count as **black** nodes.
4. Children of a **red** node are **black** nodes.
5. For all nodes x :
 - all paths from x to NIL's have the **same** number of **black** nodes on them.

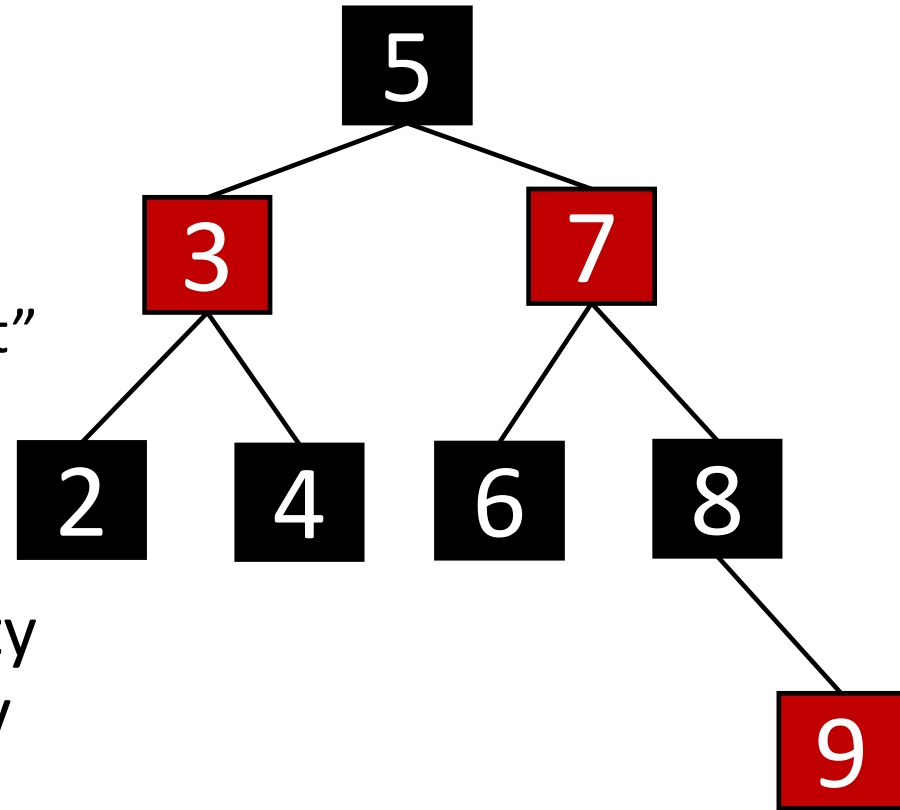


Which of these
are red-black trees?
(NIL nodes not drawn)



Why these rules??????

- This is pretty balanced.
 - The **black nodes** are balanced
 - The **red nodes** are “spread out” so they don’t mess things up too much.
- We can maintain this property as we insert/delete nodes, by using rotations.



This is the really clever idea!

This **Red-Black** structure is a **proxy for balance**.

It’s just a smidge weaker than perfect balance, but we can actually maintain it!

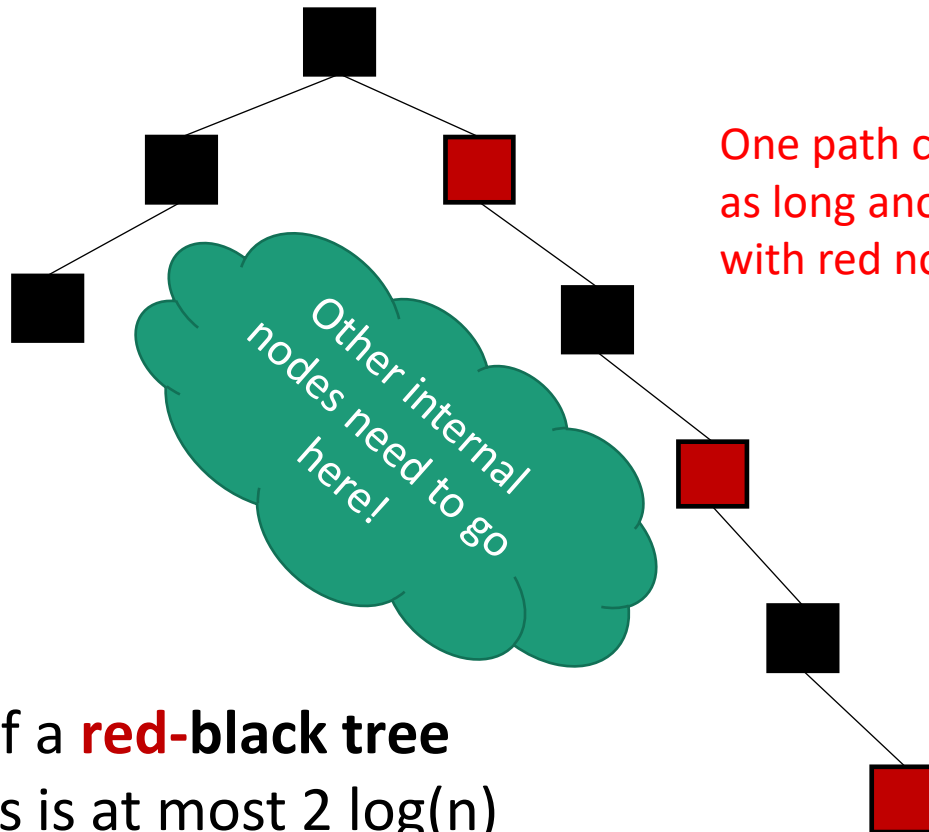




Lucky the
lackadaisical
lemur

This is “pretty balanced”

- To see why, intuitively, let's try to build a Red-Black Tree that's unbalanced.

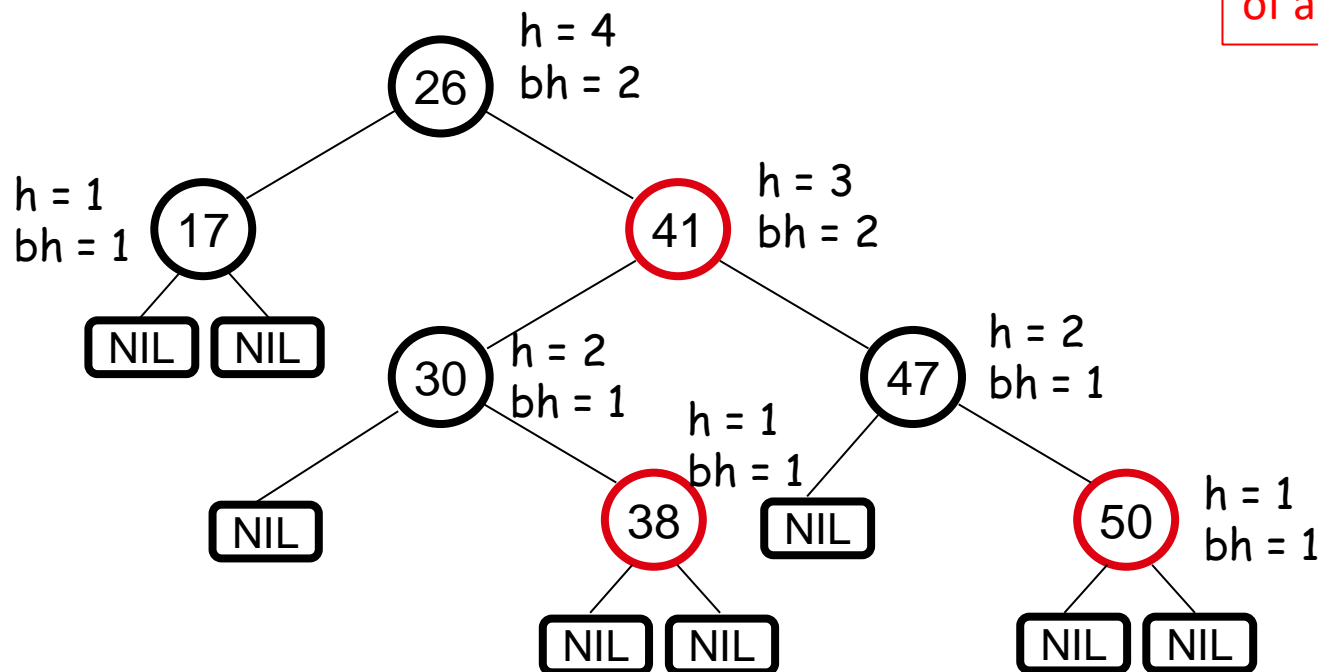


Conjecture:
the height of a **red-black tree**
with n nodes is at most $2 \log(n)$



Definitions

Aside: Formal
Proof on the
maximum height
of an RB Tree



- **Height of a node $h(x)$** : the number of edges in the **longest** path to a leaf (including the edge from leaf to NIL).
- **Black-height $bh(x)$** of a node x : the number of black nodes (including NIL) on the path from x to a leaf, not counting x .



Height of Red-Black-Trees

Aside: Formal
Proof on the
maximum height
of an RB Tree

A red-black tree with n internal nodes
has height at most $2\log(n + 1)$

- Need to prove two claims first ...



Claim 1

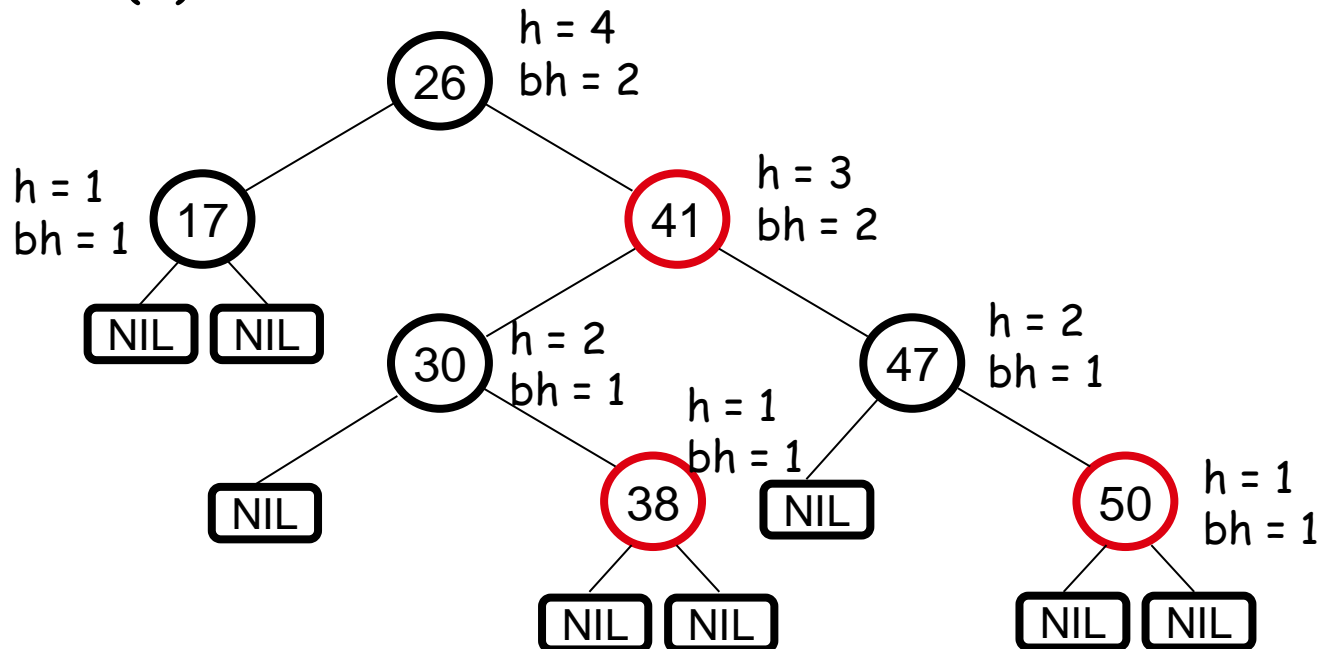
- Any node x with height $h(x)$ has

$$bh(x) \geq h(x)/2$$

- Proof**

- By property 4, at most $h(x)/2$ **red** nodes on the path from the node x to a leaf
- Hence at least $h(x)/2$ are **black**

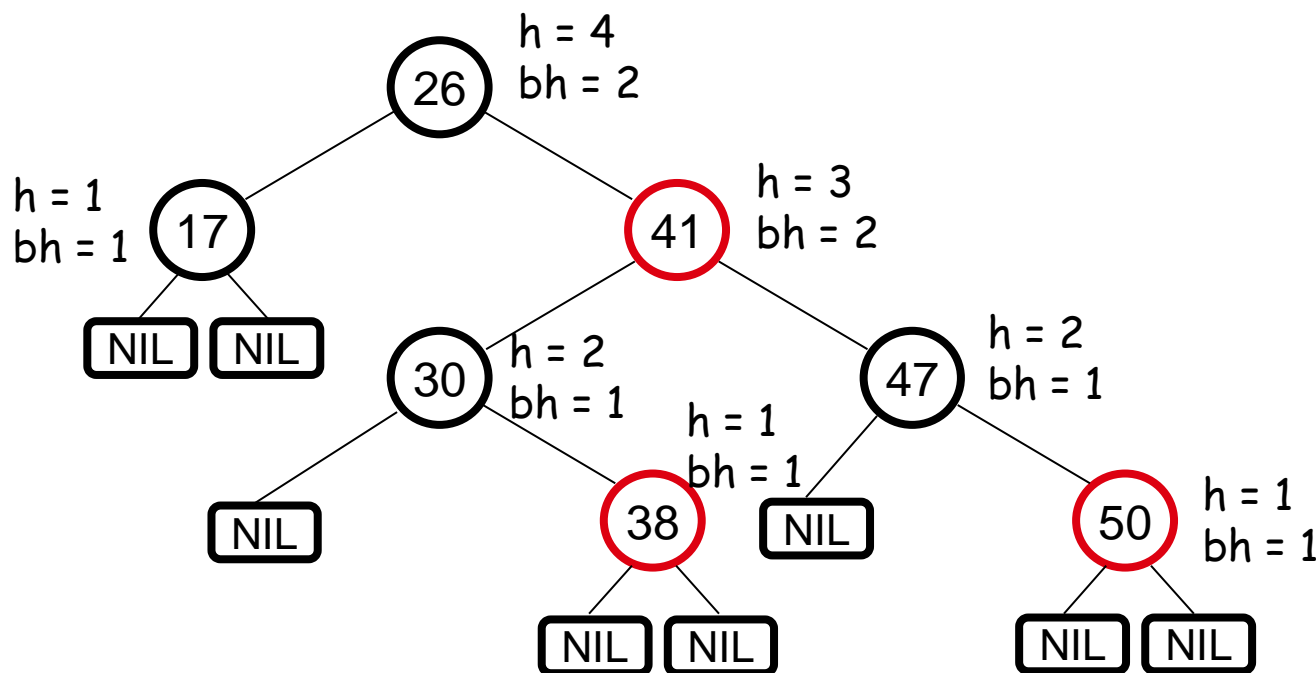
- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x :
 - all paths from x to NIL's have the **same** number of **black nodes** on them.



Claim 2

Aside: Formal
Proof on the
maximum height
of an RB Tree

- The subtree rooted at any node x contains **at least** $2^{bh(x)} - 1$ internal nodes



Claim 2 (cont'd)

Proof: By induction on $h(x)$

Basis: $h(x) = 0 \Rightarrow$

x is a leaf (**NIL**) \Rightarrow



$bh(x) = 0 \Rightarrow$

of internal nodes: $2^0 - 1 = 0$

Inductive Hypothesis: assume it is true for $h(x) = h-1$



Claim 2 (cont'd)

Aside: Formal
Proof on the
maximum height
of an RB Tree

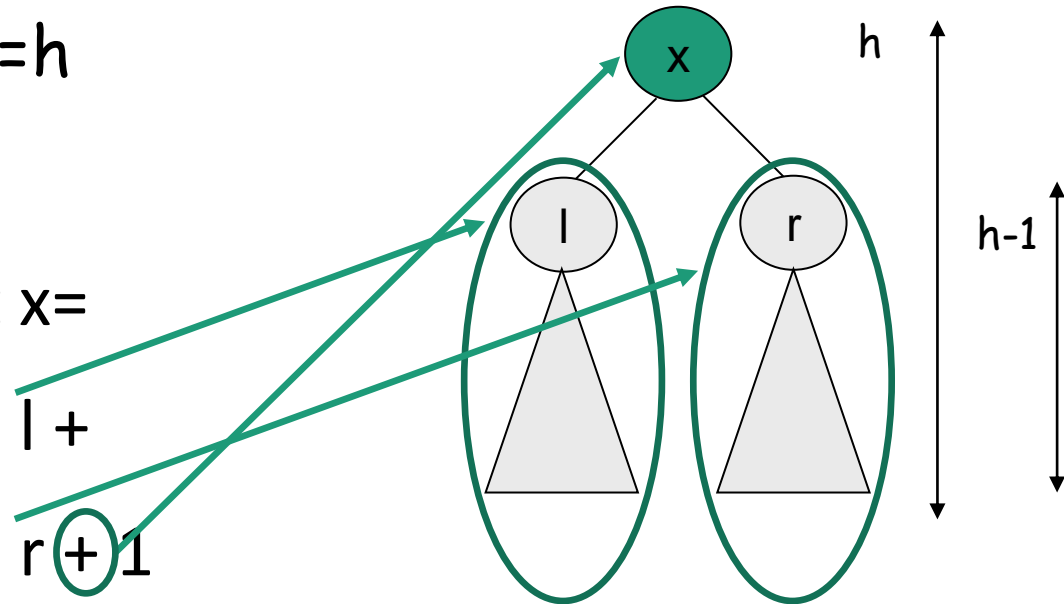
Inductive step:

- Prove it for $h(x)=h$

internal nodes at $x =$

internal nodes at $l +$

internal nodes at $r + 1$



Using inductive hypothesis:

$$\text{internal nodes at } x \geq (2^{bh(l)} - 1) + (2^{bh(r)} - 1) + 1$$

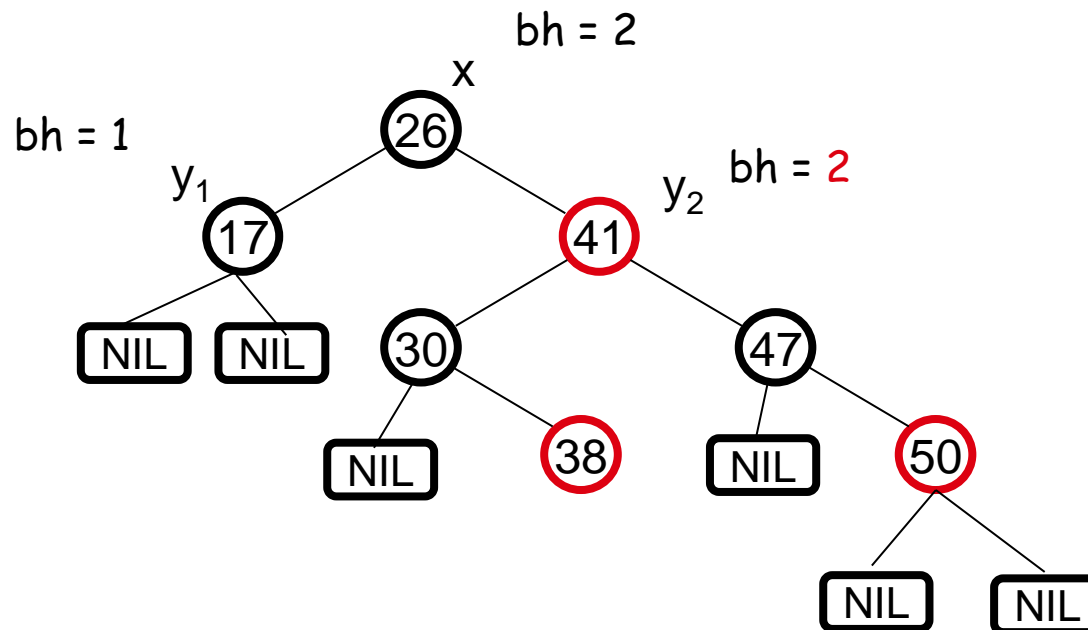


Claim 2 (cont'd)

- Let $bh(x) = b$, then any child y of x has:

- $bh(y) = b$ (if the child is **red**), or
- $bh(y) = b - 1$ (if the child is **black**)

$$\left. \begin{array}{l} \bullet \\ \bullet \end{array} \right\} bh(y) \geq bh(x) - 1$$



Claim 2 (cont'd)

- So, back to our proof:

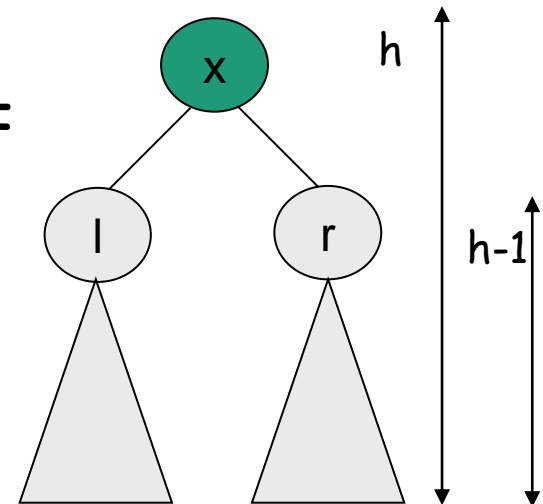
Aside: Formal
Proof on the
maximum height
of an RB Tree

internal nodes at $x \geq (2^{bh(l)} - 1) + (2^{bh(r)} - 1) + 1$

$$\geq (2^{bh(x) - 1} - 1) + (2^{bh(x) - 1} - 1) + 1 =$$

$$2 \cdot (2^{bh(x) - 1} - 1) + 1 =$$

$$2^{bh(x)} - 1 \text{ internal nodes}$$



Height of Red-Black-Trees

A red-black tree with n internal nodes has height at most $2\log(n+1)$.

Proof:

n
number
of internal
nodes

$$bh(\text{root}) = b$$

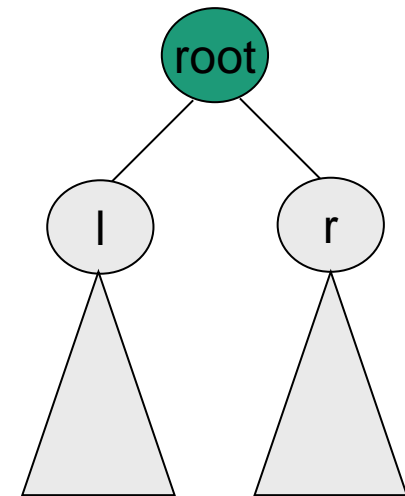
$$\geq 2^b - 1$$

Claim 2

$$\text{height}(\text{root}) = h$$

$$\geq 2^{h/2} - 1$$

Claim 1



• Solve for h :

$$n + 1 \geq 2^{h/2}$$

$$\log(n + 1) \geq h/2 \Rightarrow$$

$$h \leq 2 \log(n + 1)$$



-
- A diagram of a binary tree structure. The root node is a red square labeled 'x'. It has two children: a red square labeled 'y' on the left and a red square labeled 'z' on the right. Node 'y' has three children: a black square, a grey triangle, and a black square. The leftmost black square child of 'y' has two children: a black square and a black square. The bottommost black square child of 'y' is labeled 'NIL'. A thick purple line highlights a path from the root 'x' down to the bottommost 'NIL' node. A green arrow points upwards from the bottom towards the tree. A small penguin is in the top right corner.

Then:

$$\geq 2^{\text{height}/2} - 1 \quad \text{bh}(\text{root}) \geq \text{height}/2 \text{ because of RBTree rules.}$$

Rearranging:

$$n + 1 \geq 2^{\text{height}/2} \Rightarrow \text{height} \leq 2\log(n + 1)$$